

The Daemon, the GNU, and the Penguin

A History of Free and Open Source

by Peter H. Salus

试用水印

Dr. Salus is the author of "[A Quarter Century of UNIX](#)" and [several other books](#), including "HPL: Little Languages and Tools", "Big Book of Ipv6 Addressing Rfcs", "Handbook of Programming Languages (HPL): Imperative Programming Languages", "Casting the Net: From ARPANET to INTERNET and Beyond", and "The Handbook of Programming Languages (HPL): Functional, Concurrent and Logic Programming Languages". There is an interview with him, audio and video, "[codebytes: A History of UNIX and UNIX Licences](#)" which was done in 2001 at a USENIX conference. Dr. Salus has served as Executive Director of the USENIX Association.

This work is licensed under the Creative Commons Attribution-NonCommercial-NoDerivs License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-nc-nd/2.0/> or send a letter to Creative Commons, 559 Nathan Abbott Way, Stanford, California 94305, USA.

This easily printable version was produced by [Scott Dowdle](#).

Table of Contents

Preface	
Chapter 0. 1968 and 1969	
Chapter 1. Ancient History	
Chapter 2. UNIX	
Chapter 3. The Users	
Chapter 4. A Tale of Two Editors	
Chapter 5. UUCP and USENET	
Chapter 6. 1979	
Chapter 7. BSD and the CSRG	
Chapter 8. "Free as in Freedom"	
Excursus: Hardware	
Chapter 9. MINIX	
Chapter 10. SUN and gcc	
Excursus: UUNET	
Chapter 11. OSF and UNIX International	
Chapter 12. GNU, the GPL and Cygnus	
Chapter 13. USL v The Regents of the University of California	
Chapter 14. BTL after UNIX: Plan 9 and Inferno	
Chapter 15. Commercial UNIXes to BSDI	
Excursus: the GPL and other Licenses	
Chapter 16. The Hurd and BSDI	
Chapter 17. The Web	
Chapter 18. "Just for Fun"	
Chapter 19. Tanenbaum and Torvalds	
Chapter 20. Proliferating Penguins - Part I	
Chapter 21. Daemonology	
Chapter 22. Yet More Penguins	
Chapter 23. Oceans of FUD	
Chapter 24. The Documents of Freedom	
Chapter 25. The URL on Your Cereal Box	
Some Errata	

Preface

The activities of a distributed and unorganized band of scholars led to the conceptual revolution that produced the modern world. For example, Copernicus (1473-1543) observed the heavens and recorded his measurements. In 1563, Tycho Brahe (1546-1601) noted that Copernicus' figures weren't quite right, so, from 1577 to 1597, Tycho recorded extraordinarily accurate astronomical measurements. In 1599 Tycho moved from Denmark to Prague, where Johannes Kepler (1571-1630) was his assistant, until he succeeded him in 1601, when Tycho died.

Copernicus established heliocentricity. Tycho found that circular orbits just didn't work, and devoted decades to better measurements, which Kepler later used to determine that the orbits were ellipses, not circles. (In 1610, Galileo [1564-1642] pointed out that one could observe phases on Venus, and that therefore Venus must be nearer the Sun than the Earth was.) And, Newton (1643-1727) showed us the force (gravity) that held everything in place.

Poland. Denmark. Austria. Italy. Germany. England. Despite the Papacy, the 30 Years' War, turmoil in the Netherlands, in France, and in England, thought moved in print and in correspondence. Though countries were at war and religions were in conflict, scientific exchange of ideas and sharing of data persisted.

During the Renaissance it could take months for findings to reach those interested in other countries. In the seventeenth and eighteenth centuries lengthy epistles between scholars were distributed to others beyond the addressees. Scientific journals followed. Thanks to the progress of communications media, it now takes seconds where it once took decades for an idea or a discovery to proliferate. The fact is undeniable: Invention and scholarship have been the motor driving the development of civilization and culture.

The revolution of knowledge has led us to exploration and discovery. The computer, the Internet, and the Web have led to a similar revolution. While certainly no computer user, Thomas Jefferson, in a letter to Isaac McPherson (13 August 1813), wrote:

If nature has made any one thing less susceptible than all others of exclusive property, it is the action of the thinking power called an idea, which an individual may exclusively possess as long as he keeps it to himself; but the moment it is divulged, it forces itself into the possession of every one, and the receiver cannot dispossess himself of it.

My aim is to show how the advent of the computer and the Internet have given rise to the expansion of the academic/scholarly notions of sharing, and how this in turn has brought us free and open software, which will bring about a major change in the way we do business.

This effort is more than a history of Linux, of the Free Software Foundation (FSF), the Internet, software licensing, and myriad other topics. It will contain a number of histories within it, which (I hope) will serve as an antidote to the cloud of FUD stirred up by those who fear that change will mean that their businesses will fail (certainly more a sign of lack of imagination and flexibility than of anything else).

On the contrary: change yields opportunity. But change also requires adaptability. We are embarking on

a new business model, which will change the way we do business as much as mass production and global electronic communication did over the 19th and 20th centuries.

Since 1990, there has been an insistent drumbeat of anti-FSF FUD. Since 2000, this has focused on Linux. Some examples of this are:

- On June 1, 2001, Steve Ballmer, CEO of Microsoft, told the Chicago Sun-Times: "Linux is cancer."
- On October 15, 2002, Darl McBride, CEO of The SCO Group, said: "We are more committed to Linux than ever before."
- On March 4, 2003, Blake Stowell, SCO director of Public Relations, said: "C++ is one of the properties SCO owns."
- On May 14, 2004, the Alexis de Tocqueville Institution issued a press release in which it revealed that its Director, Ken Brown, had discovered that Linus Torvalds had not "invented" Linux.
- On August 26, 2004, Kieran O'Shaughnessy, director of SCO Australia and New Zealand, told LinuxWorld: "Linux doesn't exist. Everyone knows Linux is an unlicensed version of Unix."

The remarks are noise. But though ludicrous, statements like these make businessfolk fearful. They then hug Windows the way a different Linus clutches his blanket. My goal here is to show a wider audience just what went into the creation of open source and its worldwide network of contributors and users over the past 50 years.

Over four centuries have passed since our static heliocentric universe was replaced by a dynamic one. Today, the business model that has persisted since the late eighteenth century is being replaced. Here's how it's happening.

Chapter 0. 1968 and 1969

1. In June 1968, the Federal Communications Commission's "Carterphone" decision compelled AT&T to allow its customers to connect non-Western Electric equipment to the telephone network. [FCC Docket Number 16942; 13 FCC 2nd, 420].
2. In July 1968, Andrew Grove and Gordon Moore founded Intel.
3. In August 1968, William G. McGowan established Microwave Communications of America [MCI] and the FCC ruled that MCI could compete with AT&T, using microwave transport between Chicago and St. Louis.
4. In December 1968, the Defense Advanced Research Projects Agency let a contract to Bolt, Beranek and Newman of Cambridge, MA, for just over \$1 million. The contract was for a packet-switching network of four nodes.

Four more events of importance followed the next year.

1. In August, humans landed on the moon.
2. Summer saw the invention of UNIX.
3. In the autumn, those first four nodes of the ARPAnet went up.
4. And, in December, Linus Torvalds was born.

Had anyone asked, I would have thought the first of these events was the most important. Outside of his immediate family, I seriously doubt whether anyone even knew about the last of these.

As of the outset of the Twenty-First Century, the moon landing has taken us nowhere. The other items in this list though are the stuff of revolution.

Chapter 1. Ancient History

While mechanical calculation goes back to the seventeenth century, computation is far more recent. Though first conceived by Charles Babbage in 1823, the computer as we know it needed more than a century to come into being. The first true electro-mechanical computer was Harold Aiken's Mark I (conceived in 1937 and put into operation in 1944) and the first fully electronic machine was Maurice Wilkes' EDSAC (1949).

IBM and SHARE

The first commercial computer, the IBM 701, wasn't completed until late in 1952. The first production machine was shipped from Poughkeepsie to the IBM headquarters building in Manhattan that December. The second machine was destined for Los Alamos, and production continued in IBM's Poughkeepsie facility through June 1954, when machine 18 was shipped to Lockheed in Burbank. That's rather slow production by our standards, but literally everything was new in the early 1950s.

Prior to the 701, all computers had been one-offs. Aiken's, Wilkes', ENIAC, etc.; each was sui generis. The 701 was a genuine breakthrough. On 7 May 1954, the redesigned 701 was announced as the IBM 704. It was more than merely a redesign. The 704 was incompatible with the 701. It had 4096 words of magnetic core memory. It had three index registers. It employed the full, 36-bit word (as opposed to the 701's 18-bit words). It had floating-point arithmetic. It could perform 40,000 instructions per second. While deliveries began in late 1955, the operators (today we would think of them as system administrators) of the eighteen 701s were already fretful months earlier.¹

IBM itself had no solution to the problem. Though IBM had hosted a "training class" for customers of the 701 in August 1952, there were no courses, no textbooks. But several of the participants in the training class decided to continue to meet informally and discuss mutual problems. (According to Pugh², their first meeting was "in February 1953 during an AIEE-IRE Computer Conference in Los Angeles.") The participants agreed to hold a second meeting after their own 701s had been installed. The second meeting was hosted by Douglas Aircraft in Santa Monica in August 1953. There were other informal meetings and then, following an IBM Symposium, The RAND Corporation hosted a meeting in Los Angeles in August 1955 of representatives from all seventeen organizations that had ordered 704s. It was at this meeting that the world's first computer user group was formed. It was called SHARE.

IBM encouraged the operators to meet, to discuss their problems, and to share their solutions to those problems. IBM funded the meetings as well as making a library of 300 computer programs available to members. SHARE, 50 years later, is still the place where IBM customers gain information. (A number of the earliest contributed programs are still available.)

The importance of SHARE can be seen in the fact that in December 1955, early purchasers of Remington Rand's ERA1103A formed an organization called USE [= Univac Scientific Exchange]. In 1956, user groups for Burroughs and Bendix computers were formed, as well as IBM's GUIDE, for users of their business computers. Though SHARE was vendor-sponsored at the outset, today it is an independent organization.

User groups are one thread in the complex fabric which we employ today. Another is communication.

DARPA and IPTO

In response to the USSR's launching of Sputnik in October 1957, the US Department of Defense set up the Defense Advanced Research Projects Agency (DARPA³). That charge was "to think independently of the rest of the military and to respond quickly and innovatively to national defense challenges."

In 1962, Jack Ruina, the Director of DARPA, hired J. C. R. Licklider to be the first Director of DARPA's new Information Processing Techniques Office (IPTO).⁴

Originally, the IPTO was to extend research into the computerization of the air defense system. The IPTO funded research into advanced computer (and networking) technologies and funded fifteen groups to do research in human-computer interaction and distributed systems. (Among the research sites were: Carnegie-Mellon University, MIT, the RAND Corporation, the Stanford Research Institute, the System Development Corporation, UC Berkeley, UC Santa Barbara, UCLA, the University of Southern California, and the University of Utah.) In 1963, Lick (as many called him) funded Project MAC at MIT, headed by Robert Fano.⁵ Project MAC explored the potential for communities on time-sharing machines. That is, relationships among the uses and the users of shared mainframes.

And this leads directly to the next strand in our narrative: time-sharing.

Time-Sharing

John McCarthy had begun thinking about time-sharing in the mid-1950s. But it was only at MIT in 1961-62 that he, Jack Dennis and Fernando Corbato talked seriously about permitting "each user of a computer to behave as though he were in sole control of a computer."⁶

When McCarthy went to MIT from Dartmouth in 1957, it was clear that time-sharing the IBM 704 would require an interrupt system which didn't exist yet. So McCarthy proposed a hardware solution involving a relay whereby the 704 could be set to "trapping mode" by an external signal. But, like many other brilliant insights, McCarthy's notion went undeveloped for several years.

Four years later, MIT had a transistorized computer, the IBM 7090, and so Corbato wrote CTSS (Compatible Time-Sharing System). While it had bugs, it was a wild success, influencing systems at Dartmouth (DTSS) and the Incompatible Time-Sharing System (ITS) for the PDP-10s at MIT (more about this later).

At the same time, Lick's imagination led him to note how many different multi-million dollar computers he was funding, each of which was a solitude, unable to communicate with others. In early 1963 he sent a memo to "Members and Affiliates of the Intergalactic Computer Network." He consistently asserted that the computer was a communications, not a computation device. Then he returned to MIT.

Lick's successor at the IPTO was Robert Taylor. He was interested in networking and, in 1966, was funding 17 sites with a variety of incompatibilities. He needed help; and he found it in Larry Roberts.

Roberts had been working at the Lincoln Laboratory in Massachusetts since 1963. While there, he and Thomas Marill had conducted a networking experiment connecting the Systems Development Corporation's AN/FSQ-32 in Santa Monica, CA, with the TX-2 at Lincoln via a 1200 bps dedicated phone link. This permitted any program on one machine to dial the other computer, log in and run a

program from a server (somewhat like a subroutine call).⁷

While this was quite an achievement, it really did not further the aim of ARPA, except to demonstrate that long-distance data transfer via telephone wires was indeed feasible.⁸

In April 1967, Roberts and Taylor took advantage of the meeting of the IPTO Principal Investigators in Ann Arbor, MI, to talk up their ideas of a network. Some of the PIs were interested in "resource sharing," but the contractors in attendance set up a sub-group, "Communication Group," to work on problems. Among the problems were the conventions to be used in communications and the kinds of communications lines.

It was agreed that work should be begun on the conventions and that the connections would be via dial-up lines. The plan as developed was for the computer sites to be connected via commercial phone lines and data sets, so that each computer could be connected with every other computer via circuit-switching. During the discussion, Wesley Clark (who had moved to Washington University in St. Louis from Lincoln) had an idea. He thought about it and described it to Roberts after the meeting during a shared cab ride between Ann Arbor and the Detroit airport.

Clark's idea was that the problems of working out the many possible connections could be solved by placing a mini-computer on each site. These mini-computers would communicate with each other and each site would only have to concern itself with the task of communicating with its mini. Roberts incorporated the idea into his summary of the meeting, "Message Switching Network Proposal," which he sent on April 27, 1967. He called the mini an "Interface Message Processor." The IMP was born.⁹

Nearly a year later, on March 1, 1968, the IPTO reported to the Director of ARPA that the specifications were "essentially complete." Larry Roberts submitted a "program plan" to the Director on June 3rd and it was approved on June 21st. The ARPA budget for 1968 earmarked \$500,000 for the ARPANET.

ARPA sent out a Request for Quotation to 140 potential bidders. The Defense Supply Service - Washington received twelve proposals. Four of the bidders were deemed to be in contention and, finally, the week before Christmas 1968, the contract was awarded to BBN in Cambridge, Massachusetts. Work began on January 2, 1969. At the end of December, there were four nodes on the ARPAnet: UCLA, SRI, UCSB, and the University of Utah.

Excursus: Law I

In 1949, the Truman Department of Justice filed suit against AT&T and Western Electric, claiming the companies were acting "in restraint of trade." On 24 January 1956, Judge Thomas F. Meaney entered a "consent decree," in which the companies were enjoined "from commencing ... manufacture for sale or lease any equipment" other than that used in providing telephone or telegraph services; from "engaging ... in any business not of a character or type engaged in by Western or its subsidiaries ..."; and AT&T was enjoined "from engaging ... in any business other than the furnishing of common carrier communications services."

There were a few exceptions. Exception (b) was "experiments for the purpose of testing or developing new common carrier communications services."

AT&T was further required to reveal the patents it held and to license these when asked. No one could

have foreseen the problems that this consent decree would entail.

¹ Actually, there were nineteen 701s, the first having gone to IBM World Headquarters. Eventually (by 1960), 123 IBM 704 systems were sold.

² E.W. Pugh, *Building IBM*; MIT Press, 1995; p.186

³ DoD directive 5105.15 (7 February 1958) set up "The Advanced Research Projects Agency" (ARPA). On 23 March 1972, by DoD directive, the name was changed to DARPA. On 22 February 1993, DARPA was "redesignated" ARPA, and on 22 February 1996, Public Law 104-106 (Title IX of the FY 1996 Defense Authorization Act) directed an "organizational name change" to DARPA. The basic "charge" of the Agency was not changed significantly.

⁴ Licklider (1915-1990) has been called the Father of Artificial Intelligence, the Father of Cybernetics, the Father of the ARPAnet, and of many other things. See M.M. Waldrop, *The Dream Machine* (Viking Press, 2001).

⁵ Fano (1917-) was born in Turin, Italy, and studied there until he emigrated to the US in 1939. He received his Sc.D. from MIT and joined its faculty in 1947. He has done important work in information theory (with Shannon), microwave transmission and networking.

⁶ McCarthy received his Ph.D. from Princeton in 1951 and coined the term "Artificial Intelligence" at the 1955 Dartmouth Conference. He was the creator of Lisp and received the 1971 Turing Award.

⁷ The TX-2 had been installed at Lincoln Lab in 1958, the successor to the TX-0 (1955), the first transistorized computer. The original team included Wesley Clark as the designer and Ken Olsen -- who would go off to found DEC -- as the engineer-in-charge. The TX-2 was the computer on which Ivar Sutherland in 1959 designed and ran Sketchpad, the first graphics program. SDC's AN/FSQ-32 ran TSS (Time-Sharing System) in 1963, which had been designed in response to a challenge from Licklider.

⁸ For a fuller narrative of how Roberts got to Washington, see my *Casting the Net* (1995), chapter 3.

⁹ Detailed pre-history of the Internet can be found in *Casting the Net*.

Chapter 2. UNIX

In spring 1969, AT&T decided to terminate its involvement in a project called Multics -- Multiplexed Information and Computing Service -- which had been started in 1964 by MIT, GE and Bell Labs. This left those at AT&T Bell Labs who had been working on the project -- notably Doug McIlroy, Dennis Ritchie and Ken Thompson -- at loose ends. Doug immediately got involved with other things in Murray Hill, NJ, but Dennis and Ken had been interested in the project per se and wanted to explore several of its ideas.

Ken has said:

Dennis and [Rudd] Canaday and I were discussing these ideas of the general nature of keeping the files out of each other's hair and the nitty-gritty of expanding, of the real implementation where you put block addresses... We did it in Canaday's office, and, at the end of the discussion, Canaday picked up the phone; there was a new service at Bell Laboratories that took dictation. You call up essentially a tape recorder and you give notes, and then the next morning the notes are typed and sent to you. The next day these notes came back, and all the acronyms were butchered, like 'inode' was 'eyen...'. So we got back these descriptions and they were copied, and we each had copies of them and they became the working document for the file system -- which was just built in a day or two on the PDP-7.

At first ... we used it for other things, you know, the famous Space Travel game, and it was the natural candidate as the place to put the file system. When we hacked out this design, this rough design of the file system on the dictation [machine] that day in Canaday's office, I went off and implemented it on the PDP-7.

I won't go into full detail on the evolution of that file system on the PDP-7 to Unics [Uniplexed Information and Computing Service, a pun on cut-down (emasculated) Multics devised by Peter Neumann, an inveterate punster.¹ Several people told me that Brian Kernighan had changed the spelling to UNIX, but Brian told me that he had not, and that no one recalled who had done it.] For now, it is important to realize that it was the cooperative product of several brilliant minds: Ritchie, Thompson and Canaday, of whom Robert Morris (who joined Bell Labs in 1960 and is now Chief Scientist at the National Computer Security Center in Maryland) said he was the "most underrated" of the original participants.²

In August 1969, Ken Thompson's wife Bonnie took their year-old son on a trip to California to show off to their families. As a temporary bachelor, Ken had time to work.

I allocated a week each to the operating system, the shell, the editor and the assembler [he told me]... and during the month she was gone, it was totally rewritten in a form that looked like an operating system, with tools that were sort of known, you know, assembler, editor, and shell -- if not maintaining itself, right on the verge of maintaining itself, to totally sever the GECOS [= General Electric Comprehensive Operating System, a clone of System/360 DOS] connection. ... Yeh, essentially one person for a month.

It didn't exist by itself for very long ... maybe a day or two before we started developing the things we needed.

While Multics certainly influenced UNIX, there were also profound differences.

Dennis Ritchie explained:

We were a bit oppressed by the big system mentality. Ken wanted to do something simple. Presumably, as important as anything was the simple fact that our means were much smaller -- we could get only small machines with none of the fancy Multics hardware.

So UNIX wasn't quite a reaction against Multics, it was more a combination of these things. Multics wasn't there for us any more, but we liked the feel of interactive computing that it offered; Ken had some ideas about how to do a system that he had to work out; and the hardware available as well as our inclinations tended to trying to build neat small things, instead of grandiose ones.

Thompson "scarfed up" a PDP-7 and "did this neat stuff with it," Ritchie told me, modestly. Thompson created a new toy that would initiate work on a new system all over the world.

Soon a PDP-11 was acquired and UNIX was rewritten and expanded and rewritten. With McIlroy prodding, Dennis and Ken produced a UNIX Programmer's Manual (dated "November 3, 1971"). A "Second Edition" was issued June 12, 1972: "the number of Unix installations has grown to 10, with more expected," the Preface told us. Third Edition of the manual appeared "February, 1973," and noted that there were "now 16 installations." That was soon to wax quite rapidly.

All of the first 10 installations were at AT&T in New Jersey. In the late summer of 1972, UNIX leaped across the Hudson River to an office on the 14th floor of 330 Madison Avenue in Manhattan. Neil Groundwater had joined New York Telephone upon graduating from Penn State. He commuted from his apartment in Manhattan to Whippany, NJ., where he worked on programming for the Electronic Switching System. But being in Whippany placed him in proximity to Bell Labs and he began learning about UNIX. It was no easy task. "There was documentation on some parts," he told me. "But as we would come to say years later, 'Use the source, Luke' was the sole answer to many questions."³

In October 1973, Dennis Ritchie and Ken Thompson drove up the Hudson Valley to the new IBM Research Center at Yorktown Heights to deliver the first UNIX paper at the Symposium on Operating System Principles.

"It was a beautiful fall day," Dennis remarked. Ken, who delivered the paper, told me: "The audience was several hundred. I was pretty nervous. The response was the normal, polite applause. I don't recall any questions."

Ken was over-modest. The audience was quite enthusiastic. Ken and Dennis were immediately asked for copies of the new system.

This put the AT&T lawyers in a bind: was a computer operating system part of "common carrier communications services"? Was AT&T required to distribute UNIX?

The decision of the corporate lawyers was that Bell Labs should distribute UNIX to academic and research institutions at the cost of the media involved plus a shipping charge. Within a few months, several dozen institutions requested UNIX.

¹ Peter G. Neumann holds doctorates from Harvard and Darmstadt. After a decade at Bell Labs, he moved to SRI in 1971 and has remained there. Among other things, Neumann is the co-founder of People for Internet Responsibility and chairs the National Committee for Voting Integrity.

² Canaday graduated from Harvard in 1959 and subsequently received MS and PhD degrees from MIT. He spent 25 years at Bell Labs, taking early retirement in 1989. Long interested in business applications, Canaday was the manager of the Programmer's Workbench "gang." He is the founder of SumTime.

³ A fuller version of Groundwater's narrative is in chapter 7 of A Quarter Century of UNIX.

试用水印

Chapter 3. The Users

Let's go back to the mid-1950s. At the time that Judge Meaney was considering the action against AT&T, IBM was coming out with the 704, an upgrade of the 701. As mentioned earlier, the transitioning from the 701 to the 704 wasn't easy, so some of IBM "operators" formed the organization still known as SHARE.

Soon, many computer manufacturers were sponsoring user organizations. DECUS -- the DEC Users' Society -- first met in 1961. It soon had a British branch (DECUS UK), and rapidly became yet more international. Remington Rand, Bendix and Burroughs had formed user groups. And in but a few years, Prime and Apollo had user organizations as well -- PRIMUS and ADUS.

So, by the beginning of 1974 there were a number of user groups exchanging information and a new operating system that was beginning to get folks excited. No one had thought seriously about licensing. And there were 40 nodes on the ARPAnet.

Early in 1974, Mel Ferentz (then at Brooklyn College)¹ and Lou Katz (then at Columbia's College of Physicians and Surgeons)² called a meeting of UNIX users in New York in May. Ken Thompson supplied them with a list of those who had requested a copy of UNIX after the SOSF meeting. Nearly three dozen in under six months. The meeting took place on May 15, 1974. The agenda was a simple one: descriptions of several installations and uses; lunch; "Ken Thompson speaks!"; interchange of UNIX hints; interchange of DEC hints; free-for-all discussion. Lou told me that he thought there were about 20 people in attendance; Mel thought it might have been a few more than that. That's the organization that's now the USENIX Association.

The Ritchie-Thompson paper appeared in the July 1974 issue of Communications of the ACM. The editor described it as "elegant." Soon, Ken was awash in requests for UNIX.

Mike O'Dell's reaction to the article is typical. In 1974, Mike was an undergraduate at the University of Oklahoma. He told me:

When the famous 1974 CACM issue appeared, I was working at the OU Computer Center. We had this thing called ITF, the Intermittent Terminal Facility, which had the world's worst implementation of BASIC, and one of the guys had written some routines which let you do I/O on terminals -- and this was a non-trivial feat. So a group of us sat down and tried to figure out whether we could do something interesting. ...

The UNIX issue came. I remember going down the hall and getting it out of my mailbox and saying to myself, Oh, ACM's got something on operating systems, maybe it's worth reading. And I started reading through it. I remember reading this paper on the UNIX time-sharing system. It was sort of like being hit in the head with a rock. And I reread it. And I got up and went out of my office, around the corner to George Maybry who was one of the other guys involved with this. And I threw the issue down on his desk and said: "How could this many people have been so wrong for so long?"

And he said: "What are you talking about?"

And I said: "Read this and then try to tell me that what we've been doing is not just nuts. We've been crazy. This is what we want."

The CACM article most definitely had a dramatic impact.

Today, things would be quite different. Lou Katz wouldn't have relied on written notices; Ferentz might not have produced a purple-Dittoed newsletter. O'Dell wouldn't have gleaned the news from CACM, but from email and the Internet and the Web.

By 1975, the ARPAnet (with 60 nodes and soon to turn into the Internet) was becoming a way of distributing information. In late 1969, what we would think of as telnet and ftp were all there was. Then, in 1971, Ray Tomlinson invented email (which soon became the principal use of the ARPAnet), and in May 1975, RFC 681, "Network UNIX," appeared. Written by Steve Holmgren, Steve Bunch and Gary Grossman, the RFC began:

INTRODUCTION

THE UNIX TIME-SHARING SYSTEM PRESENTS SEVERAL INTERESTING CAPABILITIES AS AN ARPA NETWORK MINI-HOST. IT OFFERS POWERFUL LOCAL PROCESSING FACILITIES IN TERMS OF USER PROGRAMS, SEVERAL COMPILERS, AN EDITOR BASED ON QED, A VERSATILE DOCUMENT PREPARATION SYSTEM, AND AN EFFICIENT FILE SYSTEM FEATURING SOPHISTICATED ACCESS CONTROL, MOUNTABLE AND DE-MOUNTABLE VOLUMES, AND A UNIFIED TREATMENT OF PERIPHERALS AS SPECIAL FILES.

The secret, such as it was, was out. Several people have expressed their strong feelings as to just how this "put UNIX on the Net." I feel that the effect was more powerful: over the next few years, the result was that the Internet was run on UNIX. The protocols all were in tune with the "UNIX Philosophy." What we would now call "source" was widely available. Anyone actually running UNIX had accessible source. This meant that there could be true communication and we were approaching interoperability. The direct result was that UNIX was soon in use throughout the world: Japan and Australia; most of Europe; North America.

Just how widespread UNIX was can be seen from Ferentz' first mailing list (July 30, 1975) published in UNIX NEWS:

The First Mailing List

Bell Telephone Labs
Brooklyn College
Carleton College
Case Western Reserve University
The Children's Museum
City University of New York
Columbia University
Duke Medical Center
East Brunswick High School
Harvard University

Hebrew University of Jerusalem
Heriot-Watt University
Johns Hopkins University
Knox College
Naval Postgraduate School
Oregon Museum of Science
Polytechnic University of NY
Princeton University
The Rand Corporation
St. Olaf College
Stanford University
The Spence School
Univ. Catholique de Louvain
University of Alberta
U. of California, Berkeley
U. of Manitoba
U. of North Carolina
U. of Saskatchewan
U. of Texas at Dallas
U. of Toronto
U. of Utah
U. of Waterloo
U. of Wisconsin

The US, Scotland, Belgium, and Canada; universities and museums; a public high school and a private girls' school. In one year from publication. But in mid-1975, few of these establishments had electronic connectivity. In a few years, that would change and many (if not all) of the user sites would have some sort of network connection.

Another problem was hardware. In 1975, if you wanted to run UNIX, you needed a PDP-11 from DEC. That, too, was to change.

That change came about first at Princeton and then, simultaneously, at two sites about half the world apart from one another: Bell Labs in Murray Hill, NJ, and the Wollongong satellite campus of the University of New South Wales in Australia.³

First, at Princeton, in 1976 and 1977, Tom Lyon enabled some parts of UNIX to run under VM/360 on an IBM 360. It was only the first step.

At the Labs, in 1977-78, Dennis Ritchie and Steve Johnson ported UNIX to the Interdata 8/32; in Australia, Richard Miller and his colleagues were porting UNIX to the Interdata 7/32.⁴ Dennis Ritchie has said that porting to the Interdata was both a challenge and an achievement he was most proud of, for it demonstrated that UNIX could be ported to non-DEC hardware. Steve Johnson told me that once one had ported something to an alien architecture, one knew better than to try it again. He referred to the Interdata as the "Intersnail."

Australia? Yes.

John Lions read the CACM article in the summer of 1974, when the University of New South Wales was about to get a PDP-11/40, and the University negotiated a license with Western Electric.⁵ In 1975-76, UNIX was a real hit on the UNSW campus. But Lions had a problem. He wanted to use UNIX in teaching operating systems. But there was no textbook and there was no explicated version of the code -- v6. So Lions decided to do something about the lack: he wrote a commentary on the code (9073 lines at that time) and received permission from Western Electric to print out the code and commentary for instructional purposes. UNSW duplicated the code in red cardboard covers and the commentary in orange. They were as big a hit as the system.

The March 1977 issue of UNIX NEWS (vol. 2, no. 3) announced the availability of the books (to licensees) together with a note by Mel Ferentz: "Ken Thompson has seen the first version of the book and reports that it is a good job" (quite a review). The price, including airmail, was \$A17.70 (under \$20 US, at that time). The UKUUG Newsletter announced the availability of the code and commentary, too, but the next issue said that future orders should be placed with Bell Laboratories and by 1978 the volumes were no longer available. (The Labs' reproductions were in a single volume bound in black.) Someone at AT&T/Western Electric had woken up.

Once again, the proverbial cat was out of the bag.

Over the years, over nearly two decades, John Lions' Code and Commentary became the most copied work in computing. They carry the appropriate copyright notices and the restriction to licensees, but there was no way that Western Electric could stem their circulation. They were just too valuable. (I admit that I possess an nth-generation photocopy as well as treasured copies, in orange and red covers, inscribed to me by John Lions.)⁶

Why care? Because here we are in the mid-1970s with the users taking control and determining what to distribute where information was concerned. Luckily, Western Electric was no more successful at controlling information than Popes Paul V and Urban VIII were when Galileo wrote of heliocentricity. But note again: In the 1970s, you received Lions' work in hard copy, via airmail from Sydney, Australia.

Similarly, the inability of the AT&T/Western Electric lawyers to decide just what was permissible led an announcement in UNIX NEWS (30 April 1976) that Lew Law of the Harvard Science Center was

willing to undertake the task of reproducing and distributing the manuals for UNIX. ... 'The UNIX PROGRAMMER'S MANUAL' Sixth Edition dated May 1975 will be reproduced in its entirety. Most installations will want to remove several pages...

The May-June 1976 issue announced "the first mailing from the Software Exchange." This first software tape contained Harvard software; the duplication and mailing was done by Mike O'Brien, then at the University of Illinois at Chicago Circle. The idea had come to him earlier.

It depends on what you mean by "began". Actually, I was one of the "forty people in a classroom" at the meeting called much earlier than the Urbana meeting by Mel Ferentz [1975]. It was at that meeting that the idea of hosting a "Unix Users' Group tape exchange" hit me. I came home from that meeting, cleared it with my management, and declared myself open for business. By the time Urbana came around [1977], the UNIX Users' Group Software Distribution Center had been a going concern for some time.

The second software tape was announced in November 1976, along with the following note from O'Brien:

I got the "diff" listing of all changes to Bell UNIX system proper from "standard" version 6 ... Anyway, I've itemized some 50 changes, and sent the list to Ken for verification and comments. The changes will be available through the center by special request.

The second distribution tape contained contributions from the RAND Corporation, the Naval Postgraduate School, the University of California at San Diego, Yale, and UIUC. The Third Software Distribution was announced in May 1977. The last USENIX distribution was in 1988 and consisted of two 10-inch reels. The 50-bugs tape has an interesting tale connected to it.

Ken Thompson told me:

The first thing to realize is that the outside world ran on releases of UNIX (V4, V5, V6, V7) but we did not. Our view was a continuum.

After V6, I was preparing to go to Berkeley to teach for a year. I was putting together a system to take. Since it was almost a release, I made a "diff" with V6. On the way to Berkeley, I stopped by Urbana-Champaign to keep an eye on Greg Chesson who was finishing up his Ph.D. (subtle recruiting).⁷ I left the "diff" tape there and told him that I wouldn't mind it if it got around. (I think I gave it to others too, perhaps Katz.)...

Lou Katz' version is a bit different:

A large number of bug fixes was collected, and rather than issue them one at a time, a collection tape was put together by Ken. Some of the fixes were quite important... I suspect that a significant number of the fixes were actually done by non-Bell people. Ken tried to send it out, but the lawyers kept stalling and stalling and stalling.

Finally, in complete disgust, someone "found" a tape on Mountain Avenue [The address of Bell Laboratories was 600 Mountain Avenue, Murray Hill, NJ] which had the fixes.

When the lawyers found out about it, they called every licensee and threatened them with dire consequences if they didn't destroy the tape ... after trying to find out how they got the tape. I would guess that no one would actually tell them how they came by the tape (I didn't). It was the first of many attempts by the AT&T lawyers to justify their existence and to kill UNIX.

At the 1994 USENIX technical meeting, there was a 25th birthday session after which Lou "confessed" that he had received a phone message at Columbia to the effect that if he drove down to Mountain Avenue "around 2pm," he'd "find" something of interest. So he and Reidar Bornholdt drove from Manhattan to Murray Hill and "found" the can with the tape in it. Ken told me he had "no idea" how the tape had gotten there. Dennis suggested that it might have "fallen from a truck." Everyone laughed.

At this time AT&T had a strict policy of

1. no advertising
2. no support

3. no bug fixes
4. payment in advance

This forced the users to band together and compelled them to share what they had learned and what they knew.

¹ Mel went on to Rockefeller University and later became one of the founders of NYSERNET.

² Lou moved to UC Berkeley in 1981. He was the founding President of USENIX. Lou was also the first recipient of a 9-track tape of UNIX, cut by Ken.

³ The University of Wollongong is now independent.

⁴ Interdata, later bought by Perkin-Elmer, brought out the 7/32 in 1974 and the 8/32 the following year.

⁵ See chapter 15 of A Quarter Century of UNIX.

⁶ In 1996, after a great deal of correspondence and with the active assistance of Dennis Ritchie, I succeeded in getting permission from both AT&T and the (original) Santa Cruz Operation to reprint Lions' work. ISBN 1-57398-013-1.

⁷ Chesson had brought UNIX to the University of Illinois, where he received his Ph.D. in 1977. He went on to become one of the founders of Silicon Graphics.

Chapter 4. A Tale of Two Editors

In Chapter 1, I mentioned CTSS and ITS. At that time, early in the 1960s, TECO (Tape Editor and COrrector; later, Text Editor ...) was the editor everyone used on the PDP-1 and, later, on the PDP-6. Not only was it widely used, but just about everyone modified it. (In RFC 681, quoted earlier, an editor "based on QED" is mentioned. QED was written by Butler Lampson, who wrote the QED text editor for the Berkeley Time-Sharing System on the SDS 940. It was character-oriented and based on TECO. Ken Thompson used this version of QED while a student at Berkeley, prior to going to Bell Labs in 1966.) Indirectly, TECO was the ancestor of vi; directly, it was the parent of Emacs (= Editing macros for TECO).

Interestingly, Bill Joy created vi in 1976 and Richard Stallman (together with Guy Steele and Dave Moon) created Emacs the same year. The original version was based on TECMAC and TMACS, two TECO editors. Stallman and Michael McMahon ported it to the Tenex [for the DEC-10] and TOPS-20 [for the DEC-20] operating systems. [James Gosling, the creator of Oak/Java, wrote the first Emacs for UNIX at Carnegie-Mellon in 1981. RMS began work on GNU EMACS in 1984.]

Joy's creation had a more complex origin.

The editor created by Ken Thompson in August 1969 was called ed. Ken had written a version of QED for CTSS on the IBM 7094 at MIT. He and Ritchie then wrote a version for the GE-635 at Bell Labs. The cut-down version of this for the PDP-7 was ed. While TECO was known for its complex syntax, ed must have been the most user-hostile editor ever created.

Across the Atlantic in London, George Coulouris at Queen Mary College (now Queen Mary and Wakefield College) had gotten UNIX v4 in late 1973. George explained to me how unhappy he had been with ed and how he created em (editor for mortals) so that QMC students could "exploit more effectively some vdu [visual display unit] that we had recently acquired..."

Then I spent the summer of 1976 as a visitor to the CS Department at Berkeley. I worked in a room full of teletype terminals using the departmental UNIX. I had brought em with me on DECtape and installed it there for my own use...

One day, sitting at the next terminal was this fairly frenzied hacker/Ph.D. Student [Bill Joy] who told me he was writing a Pascal compiler. I showed him em, and he said "that's nice, the systems support people might be interested in that." He took me and introduced me to them. They had a couple of PDP-11s ... supporting several rooms full of vdu terminals connected at 9600 baud, an environment in which em could really shine.

I explained that em was an extension of ed that gave key-stroke level interaction for editing within a single line, displaying the up-to-date line on the screen (a sort of single-line screen editor)...

The system support person [Jeff Schriebman] said something like: "That's very nice, but if we made it available to all of our users the overheads associated with running in raw mode would swamp the cpu."

I was rather depressed by this reaction, thinking "I guess I have been unrealistic in developing an editor that is so expensive to run..."

Nevertheless, Bill and the support people took a copy of my source to see if they would use it. I then went to the East Coast for a week or so. When I returned, I found that Bill had taken my code as a starting point and had got a long way towards what was to become ex and subsequently vi, and that the editor was installed on the service machines ...

1976! Created in 1969, ed had travelled east to Australia and west to Vienna. Coulouris had created em in London and brought it to Berkeley. Now the Berkeley editor, ex, would be available on the first UCB tape. But vi, which was available on 2BSD (1979), only made it into a BTL distribution with v8 (1985).

Even in 1976, international communication and access to source meant the distribution of new tools and new programs encouraged and enlivened the user community. Let's look at the landscape for a few minutes.

1. In 1974, Bob Kahn and Vint Cerf published the first paper describing what was to become TCP/IP.
2. In 1975, RFC 681 was published.
3. In January 1976, there were 63 hosts on the ARPAnet, which was on the verge of becoming the Internet.
4. And UNIX was available throughout the world -- but only on hardware that cost well over \$10,000.

Chapter 5. UUCP and USENET

Also in 1976, Mike Lesk at AT&T developed UUCP (UNIX-to-UNIX copy).¹ V2 was implemented in 1977.²

UUCP meant that information could be directed around the network (as it was). It also meant that one could establish a telephone connection and transmit information across that (relatively expensive) link. Two years later, three graduate students in North Carolina (Tom Truscott, Jim Ellis, and Steve Bellovin) took the next step.

Tom Truscott had an early interest in chess. While a student at Duke in 1974, he devised a chess program (Duchess) and played against Ken Thompson's Belle. Duchess lost on time. (In competitive chess, each side has a given time to make its next move; Duchess exceeded that time due to a core dump.) But Truscott competed in every ACM computer chess tournament from 1974 through 1980. He also attended the 1976 UNIX Users Group meeting at Harvard (1-2 April) and the 1978 meeting at Columbia (24-28 May), where he met Ken and others.³ In 1979, Truscott went to the Labs as a summer student and, on his return to Duke, arranged for a UUCP link. (He also attended the USENIX meeting in Toronto [20-23 June], to which we'll return in the next chapter.)

When he returned to Duke, he found that Jim Ellis had installed V7 on the Computer Science PDP 11/70. They employed the auto-dialer capacity to dial up two other Duke computers and one at the University of North Carolina. Ellis and Truscott then called a meeting to discuss their idea -- to have something like the mailing lists on the ARPAnet for computer sites that weren't on the ARPAnet. Steve Bellovin, then a graduate student at the University of North Carolina, attended and then wrote the first Netnews program -- three pages of shell script (later rewritten in C). The first implementation was between the Duke and UNC Computer Science departments; the Duke Medical Center Department of Physiology was added at the beginning of 1980. In January 1980, Ellis and Truscott went to Boulder, CO, and announced their Netnews design at the USENIX meeting (January 29 - February 1).

The first version of Netnews was fairly simple and efficient. It periodically checked the "last saved" time-stamp of each file in a specified directory, and then sent any file updated since the last check to another computer using UUCP across a modem connection.

Tom Truscott and Steve Daniel (also a graduate student at Duke) then rewrote the program to create what was called Netnews Version A. Since Netnews was designed for UNIX at a university, it was automatically categorized as public domain software under the conditions of the AT&T UNIX license, which greatly facilitated its subsequent use and adoption. This implementation appeared on the 1980 USENIX distribution tape, which was distributed at the Newark, DE, meeting (June 17-20). Duke University then invited other sites to join the network, which was made easier by the fact that the software was free, starting the first Usenet expansion -- to 15 sites. But one of those was Berkeley, which resulted in an explosive growth spurt.

That connection was the responsibility of Armando Stettner, then with DEC. Someone at the Delaware meeting complained about the inordinate cost of the long-distance telephone connections needed to get news to the West Coast. Armando spoke to Bill Shannon and they said that if they could get a news feed to decvax (in New Hampshire), they'd pick up the Berkeley phone bill. [Armando later supplied the first news feeds to Europe, Japan, and Australia, too.] The network soon spread to universities

across North America, quickly establishing a critical mass of useful information, which made it even more popular.

In under a year, Usenet grew to 100 sites and over 25 articles per day (the original protocol had been optimized for 1-2 messages per day). Mike Lesk had never contemplated such uses of uucp. Truscott, Ellis and Bellovin had never imagined such popularity. The system collapsed.

In 1982, Netnews Version B was developed by Mark Horton (a graduate student at Berkeley) and Matt Glickman (a high school student) to increase efficiency so that USENET could cope with the increasing loads on the growing network. Horton continued to maintain the system till 1984, when Rick Adams at the Center for Seismic Studies took over maintenance of Version B, releasing 2.10.2. There were 135 news groups at the end of March. Version 2.11 of Netnews was released in late 1986.

In 1989, Netnews Version C was developed by Henry Spencer and Geoff Collyer at the University of Toronto again increasing efficiency.

By the late 1980s there were 20,000 newsgroups. Looking at what was current became nearly impossible. The remedy was "searching." As Mike O'Dell put it: "The Internet is now a rich fabric of resources and capabilities, and it is no longer possible to simply know all the places where interesting stuff is available. We now need tools with which to discover and navigate this world-wide treasure trove" (Computing Systems 5.4 [Fall 1992] p. 373).

Searching in a Distributed Environment

By the beginning of 1992, there were a number of "resource discovery systems" (much of this section is drawn from the special issue of Computing Systems cited above).

The earliest of these were search engines which pointed at specific, dedicated servers: whois, which responded to queries about people, network numbers and domains across the Internet; and X.500 (ISO DIS 9594-1 [1988]), a distributed directory look-up service.

The next tool was archie, developed by Alan Emtage and Peter Deutsch when they were at McGill University in Montreal. archie maintained a directory of materials available on about 1100 UNIX archives accessible by anonymous FTP. In 1992 it contained about 2.6 million files with 150 gigabytes of information.

Gopher (from the University of Minnesota's "Golden Gophers") provided a simple menu-driven interface to data searching.

Prospero (Neuman, 1992) was an "enabling technology," allowing users to create their own views of information in a distributed file system.

WAIS, wide-area information service, developed by Brewster Kahle and colleagues at Thinking Machines, was a network of over 70 servers world-wide, offering access to over 300 databases by natural language keyword search.

There were also knowbots (developed at the CNRI) and Netfind and a host of others.

But the "winner" was the World Wide Web (WWW), invented by Tim Berners-Lee at CERN, and first

published in Electronic Networking in Spring 1992. We will return to the Web in about a dozen years.

¹ Lesk earned a Ph.D. in Chemical Physics and worked in the UNIX group, where he wrote tbl, refer, lex, and UUCP. He went on to Bellcore; was head of a Division at the National Science Foundation (1998-2002); and is currently Professor of Library and Information Science at Rutgers University.

² UUCP has a rich history; see Chapter 15 of Casting the Net.

³ This was the meeting at which the name was changed to USENIX, spurred by a letter from an AT&T lawyer stating that Western Electric had not granted permission to use "UNIX" in "UNIX User Group."

试用水印

Chapter 6. 1979

In Chapter 3, UNIX V6, John Lions, and the ports of UNIX to the Interdata 7 and the Interdata 8 were mentioned. We're about to move on to V7. But first, a note about names and dates.

Names. By and large, Unix users refer to "Sixth Edition" and "V6" interchangeably. At Bell Labs, there was a continually changing version of Unix running. Only when Doug McIlroy caused the first "UNIX PROGRAMMER'S MANUAL" to be written, did there appear to be a fixed form. So, the manuals were listed by "Edition," and the system referred to was the "Version."

Dates. Every AT&T manual carried a date. This is the set.

- First Edition - November 3, 1971
- Second Edition - June 12, 1972
- Third Edition - February 1973
- Fourth Edition - November 1973
- Fifth Edition - June 1974
- Sixth Edition - May 1975
- Seventh Edition - January 1979
- Eighth Edition - February 1985
- Ninth Edition - September 1986
- Tenth Edition - October 1989

V7. The wonder of V6 was that it was under 10,000 lines of code. (But it had no full-screen editor, no windowing system, etc.) V7 was much larger. And it contained much more.

V7 accommodated large filesystems; it did not restrict the number of user accounts; it had improved reliability. Steve Johnson referred to V7 as "the first portable Unix. It also had a large number of new commands. Among these were:

COMMANDS: at, awk, calendar, cb, cd, cpio, cu, derooff, expr, f77, find, lex, lint, m4, make, refer, sed, tail, tar, touch, uucp, uux

SYSTEM CALLS: ioctl

SUBROUTINES: malloc, stdio, string

GAMES: backgammon

awk (=Aho-Weinberger-Kernighan; a pattern scanning and processing language; gawk [GNU AWK] is the most employed version), lint (Johnson, a C program verifier), make (Feldman, a utility to simplify the maintenance of other programs), and uucp (Lesk) would have been enough, but there was much more. The V7 manual had grown to over 400 pages, with two 400-page supplementary volumes. V7 contained a full Kernighan-Ritchie C compiler; a far more sophisticated shell (sh), the Bourne shell; Dick Haight's find, cpio and expr; and a large number of include files. Dated "January 1979," the title page of the Seventh Edition *UNIX PROGRAMMER'S MANUAL* bore neither Dennis' nor Ken's name. It was headed: **UNIX™ TIME-SHARING SYSTEM.**

Along with all this, V7 came with a major drawback: its performance was poorer than most V6 systems, especially those that had been "tuned." The users went to work.

Bill Joy (at Berkeley) changed the size of the data blocks on the VAX 11/780. (Jeff Schriebman ported that to the PDP-11/70 in April 1980, having gone to UniSoft from Berkeley.) In December 1979, Ed Gould moved the buffers out of kernel address space. Joy changed the stdio library on the VAX and Tom Ferrin (at UCSF) ported those changes to the PDP-11. Tom London (in Holmdel, NJ) improved the movement of output characters from user space to kernel space. John Lions (at UNSW in Australia) proposed a new procedure for directory pathnames. UNSW also provided new code for process table matches. Bruce Borden (at the RAND Corporation) provided the symorder program. Ferrin also rewrote parts of copyseg() and clearseg(). (The entire set of improvements was made available to the community on a PDP-11 distribution: 2.8.1BSD -- it was announced by Ferrin at the USENIX Conference in January 1982 in Santa Monica, CA.

The users had enhanced V7's performance dramatically.

But I've gone too far ahead.

USENIX. Ron Baecker of the University of Toronto hosted the USENIX Association meeting in Toronto, June 20-23. There were about 400 attendees.

Al Arms, an AT&T lawyer, announced the new licensing terms for V7. In addition to a fee schedule that increased the costs of the various licenses, the academic/research license no longer automatically permitted classroom use. As Greg Rose (who was one of John Lions' students at UNSW and is now with Qualcomm) told me:

It seems that in so many stages of Unix's evolution, an action that AT&T took in order to stifle something actually caused the opposite to happen.

Suppressing Lions' commentary led to wholesale Xeroxing. Here it led to something truly unforeseen.

Andrew S. Tanenbaum of the Vrije Universiteit in Amsterdam had been using V6 in his classes on operating systems. It now appeared that he would not be able to employ V7. "It was the new licensing restrictions," he told me.

When AT&T released Version 7, it began to realize that UNIX was a valuable commercial product, so it issued Version 7 with a license that prohibited the source code from being studied in courses, in order to avoid endangering its status as a trade secret. Many universities complied by simply dropping the study of UNIX, and teaching only theory. [Operating Systems (1987), p. 13]

Andy was not to be deterred: he created an operating system like Unix for use on Intel's new x86 architecture: Minix. 1

MIT. In 1974, Robert Greenblatt at the MIT AI Lab began the Lisp machine project. His first machine was called CONS (1975). 2 This was improved into a version called CADR (1977). 3 CADR was the direct ancestor of both Lisp Machines, Inc., of which Greenblatt was the founder and president, and Symbolics.

And Symbolics, in several ways, forced Richard Stallman to form the Free Software Foundation and the GNU Project. I'll discuss these in Chapter 8.

Berkeley. Though I will go into detail concerning the Computing Systems Research Group (CSRG) at Berkeley in the next chapter, I think it important to note here that 3BSD, the first Berkeley release for the VAX, came out at the end of 1979. But this was not based on V7, but on 32V. Berkeley utilities and modifications from 2BSD (including the C shell) had been added, as well as a virtual memory system done by a number of Berkeley graduate students, including Bill Joy.

It was an exciting year.

32V. Before going on to Berkeley, I should say something about 32V, the UNIX port to the VAX.

Nearly four years passed between V6 and V7. But, as I noted, things didn't stand still at the Labs. In fact, what we think of as V7 was available internally nearly a year prior to the publication of Seventh Edition.

Things hadn't remained static at DEC, either, and the first VAX (Virtual Address eXtension), a 32-bit computer, was pre-announced in 1977 and went on sale in 1978. Dennis, Ken and Steve Johnson felt alienated by DEC, for a variety of reasons. And Dennis and Steve were working on the port to the Interdata 8. So when DEC offered them a VAX, they just said "no." DEC then turned to Bell Labs in Holmdel. I spoke to Charlie Roberts, who was to manage the project.

DEC came to us in Holmdel. We were clearly the second-string. Tom London and John Reiser were interested and so was Ken Swanson, and we got the VAX in early '78. I didn't do any of the technical work. In fact, I devoted a lot of energy and time to getting the management to let us do it. It wasn't research you see. However, they let us take the time. And in about three months my small group ported Version 7 to the VAX. We got the machine in January; they had it running in April; and by August it really worked. By then folks knew what we were doing. I had had calls from about six universities -- Brown, UCLA, Berkeley, Waterloo. I can't recall the others. So I went to Roy Lipton and Al Arms in Patents and Licensing about getting it out. After a lot of back-and-forth, they decided that we could give it to one university for research purposes and that Al would set up a "special research agreement" with that institution.

I had met [Bob] Fabry at a couple of conferences and I had been out in Berkeley and given a paper and talked to [Domenico] Ferrari as well as Emmanuel Blum and Bill Joy. So, with the blessings of BTL area 11 management, we sent 32V to Berkeley. It was in October or November 1978.

With that background, let's go cross country to Berkeley.

¹ I'll return to Minix in Chapter 9. But it's important to note that some universities went along with AT&T. The first of these was Carnegie Mellon. Clem Cole told me that he and Dan Klein had gone on a hunger strike until CMU purchased a "full" \$20,000 license.

² This was a pun on the list construction operator in Lisp.

³ In Lisp, cadr returns the second element in a list.

Chapter 7. BSD and the CSRG

Professor Robert Fabry was on the program at the SOSP where Ken and Dennis delivered the first paper on Unix. He was full of enthusiasm when he returned to Berkeley. When he got back to Cory Hall, his first chore was to attempt to assemble a joint purchase (by Computer Science, Math and Statistics) of a PDP-11/45.

He then ordered a tape from Thompson and, in January 1974, Keith Standiford (a graduate student) installed Unix. As Kirk McKusick tells the tale, Ken had been involved in all the early installations -- but not this one, "though his expertise was soon needed to determine the cause of several strange system crashes." Thompson would phone Keith in the machine room, the phone would be inserted in the 300 baud acoustic modem, and Ken would "remotely debug crash dumps from New Jersey."

The next problem arose because Math and Statistics wanted to run DEC's RSTS.¹ So slices of each day were allocated: eight hours for Unix, followed by 16 hours for RSTS. One of the undergraduates introduced to Unix at that time was Eric Allman.²

"I was taking an introductory OS course," he told me, "and they had been using something called the Toy Operating System on the [CDC Cyber] 6400. But they wanted to get off it and on to Unix on the 11/40,³ where we could only work eight hours a day, and a different eight hours each day. I recall having difficulties. I was reading the manual, and I remember not understanding why you would ever want the echo command. Totally bizarre. Now, of course, I know. Of course, 4th Edition was pretty flaky. It was a system that only a researcher could love. It was slow. It didn't have a lot of tools. Then I got hired by the Ingres project."

The Ingres project of Professors Michael Stonebraker and Eugene Wong was one of the first projects moved to Unix. As they were dissatisfied with the way time had been allocated, they bought their own PDP-11/40 in spring 1974. But there still wasn't enough student time available on the 11/45. In June 1974, Stonebraker and Fabry set out to get two instructional 11/45s for Computer Science. The money was obtained early in 1975, just about the time DEC announced the 11/70, which seemed more suitable (it was announced in March). So they pooled the money. The 11/70 arrived in fall 1975, just when Ken Thompson arrived for his one-year sabbatical as visiting professor. Thompson, Bob Kridle and Jeff Schriebman brought up V6 on the newly-installed 11/70. The "50 bugs" tape (see Chapter 3) was a side-effect.

That same autumn, two new graduate students arrived on the Berkeley campus: Chuck Haley and Bill Joy. They were fascinated by V6 on the 11/70 and began working on the Pascal system that Ken had hacked together. In fact, Haley and Joy improved the Pascal system to the point that it became the system of choice for the students. And when the Model 33 Teletype terminals were replaced by ADM-3 screen terminals, they felt that ed just wasn't good enough. So they developed ex from Coulouris' em (see Chapter 4).

Meanwhile, in a different part of the universe, Kirk McKusick, at that time a "JCL hacker at Cornell, was exposed to Unix by a friend who was studying at the University of Delaware.

"He showed me how you could play games on it," McKusick told me, "so I really didn't have any exposure to Unix till I got out to Berkeley in 1976. I came out in the spring because I was

looking at graduate programs to go to, and they were on spring break, so there weren't a lot of people around. But I came across Bill Joy in the computer room where he was busily hacking away and he said 'Hi, I'm Bill Joy. This is what we're working on here - a Pascal system that runs under Unix.' And I said, 'What can you do with Unix?' So he said, 'You can edit files and compile files and you can play chess. Let me log you in.' So I proceeded to play chess the rest of the afternoon. And I decided that I liked Berkeley and that's where I'd go."

When Thompson returned to BTL at the end of the summer of 1976, Joy and Haley turned their interests toward the kernel. But word of the Pascal compiler got around. And at the beginning of 1978, Joy began producing the Berkeley Software Distribution (BSD). He offered a copy to Tom Ferrin (at UCSF) in a letter dated 9 March 1978. The "license" was on one side of a sheet of paper. Tom signed it on 13 March. The tape (in "standard tp format, 800 bpi...") consisted of:

- a) Unix Pascal system
- b) Ex text editor
- ...
- created by
 - a) W.N. Joy, S.L. Graham, C.B. Haley, K. Thompson
 - b) W.N. Joy

Bill Joy, acting as "distribution secretary," sent out about 30 "free" copies of BSD in early 1978. By mid-year enough had been done (the Pascal system had been made more robust, the system could be run on the 11/34,⁴ and vi and termcap⁵ had been written), that a "Second Berkeley Software Distribution" was put on tape.

Bill answered the phone, put together the distributions, and incorporated user feedback into the system. He also mailed out nearly 75 copies of 2BSD. 3BSD appeared in December 1979.

The CSRG.

Kirk McKusick has told the tale of the CSRG.

In the fall of 1979, Bob Fabry responded to DARPA's interest in moving towards Unix by writing a proposal suggesting that Berkeley develop an enhanced version of 3BSD for the use of the DARPA community. Fabry took a copy of his proposal to a meeting of DARPA image processing and VLSI contractors, plus representatives from Bolt Beranek and Newman, the developers of the ARPAnet. There was some reservation whether Berkeley could produce a working system; however, the release of 3BSD in December 1979 assuaged most of the doubts.

With the increasingly good reputation of the 3BSD release to validate his claims, Bob Fabry was able to get an 18-month contract with DARPA beginning in April 1980. This contract was to add features needed by the DARPA contractors. Under the auspices of this contract, Bob Fabry set up an organization which was christened the Computer Systems Research Group, or CSRG for short. He immediately hired Laura Tong to handle the project administration. Fabry turned his attention to finding a project leader to manage the software development. Fabry assumed that since Joy had just passed his Ph.D. qualifying examination, he would rather concentrate on completing his degree than take the software development position. But Joy had other plans. One night in early March he phoned Fabry at home to express interest in taking

charge of the further development of Unix. Though surprised by the offer, Fabry took little time to agree.

The project started promptly. Tong set up a distribution system that could handle a higher volume of orders than Joy's previous distributions. Fabry managed to coordinate with Bob Guffy at AT&T and lawyers at the University of California to formally release Unix under terms agreeable to all. Joy incorporated a variety of changes and 4BSD was released in October 1980.⁶

(The last CSRG system was 4.4BSD-Lite, Release 2, June 1995. Soon thereafter, the CSRG was disbanded. Another era was over.)

¹ Resource Sharing Time Sharing Extended, a multi-user OS for the PDP-11 series, generally programmed in "Basic Plus."

² Inventor of sendmail.

³ The 11/40 was a higher performance version of the 11/20, introduced in January 1973. The 11/45 was a faster, microcoded 11/20, introduced in June 1972.

⁴ The 11/34 was a cost-reduced version of the 11/35.

⁵ vi was the result of the arrival of the ADM 3a terminals; termcap was born when Joy decided to consolidate screen management, using an interpreter to redraw the screen.

⁶ from Marshall Kirk McKusick, "20 Years of Berkeley UNIX," in Open Sources (O'Reilly, 1999).

Chapter 8. "Free as in Freedom"

Richard M. Stallman, though a freshman at Harvard, began working for Russ Noftsker at the MIT Artificial Intelligence Lab in 1971. While still in high school (The Adams School through junior year, senior year at Louis D. Brandeis on West 84th Street) in New York, he had worked briefly at the IBM Science Center and at Rockefeller University.

As he put it,

I became part of a software-sharing community that had existed for many years. Sharing of software was not limited to our particular community; it is as old as computers, just as sharing of recipes is as old as cooking. But we did it more than most.

The AI Lab used a time-sharing operating system called ITS (the Incompatible Timesharing System) that the Lab's staff hackers had designed and written in assembler language for the Digital PDP-10... As a member of this community, an AI Lab staff system hacker, my job was to improve this system.

We did not call our software "free software," because that term did not yet exist, but that is what it was. Whenever people from another university or a company wanted to port and use a program, we gladly let them. If you saw someone using an unfamiliar and interesting program, you could always ask to see the source code, so that you could read it, change it, or cannibalize parts of it to make a new program.¹

Less than a decade later, everything changed for the worse. "It was Symbolics that destroyed the community of the AI Lab," rms told me. "Those guys no longer came to the Lab. In 1980 I spent three or four months at Stanford and when I got back [to Tech Square], the guys were gone. The place was dead." (Sam Williams says that Symbolics hired 14 AI Lab staff as part-time "consultants." Richard was truly the "last of the hackers.")

We see here what Richard wanted: a cooperative community of hackers, producing software that got better and better.

"In January '82 they [Symbolics] came out with a first edition," rms continued. They didn't share. So I implemented a quite different set of features and rewrote about half of the code. That was in February. In March, on my birthday [March 16], war broke out. Everyone at MIT chose a side: use Symbolics' stuff, but not return source for development. I was really unhappy. The community had been destroyed. Now the whole attitude was changing."

In the essay cited above, rms continued:

When the AI Lab bought a new PDP-10 in 1982, its administrators decided to use Digital's non-free timesharing system instead of ITS.

The modern computers of the era, such as the VAX or the 68020, had their own operating systems, but none of them were free software: you had to sign a nondisclosure agreement even to get an executable copy.

This meant that the first step in using a computer was to promise not to help your neighbor. A cooperating community was forbidden. The rule made by the owners of proprietary software was, "If you share with your neighbor, you are a pirate. If you want any changes, beg us to make them."

The idea that the proprietary software social system--the system that says you are not allowed to share or change software--is antisocial, that it is unethical, that it is simply wrong, may come as a surprise to some readers. But what else could we say about a system based on dividing the public and keeping users helpless? Readers who find the idea surprising may have taken the proprietary social system as given, or judged it on the terms suggested by proprietary software businesses. Software publishers have worked long and hard to convince people that there is only one way to look at the issue. ...

I have quoted Richard at length, because I think that his "voice" should be heard. He has frequently said that "Software wants to be free." But in 1982 and 1983 his was a single, lonely voice. He duplicated the work of the Symbolics programmers in order to prevent the company from gaining a monopoly. He refused to sign non-disclosure agreements, and he shared his work with others in what he still regards as the "spirit of scientific collaboration and openness."

In September 1983, rms announced the GNU project. In January 1984 he resigned from his job at MIT.

He has written:

I began work on GNU Emacs in September 1984, and in early 1985 it was beginning to be usable. This enabled me to begin using Unix systems to do editing; having no interest in learning to use vi or ed, I had done my editing on other kinds of machines until then.

At this point, people began wanting to use GNU Emacs, which raised the question of how to distribute it. Of course, I put it on the anonymous ftp server on the MIT computer that I used. (This computer, prep.ai.mit.edu, thus became the principal GNU ftp distribution site; when it was decommissioned a few years later, we transferred the name to our new ftp server.) But at that time, many of the interested people were not on the Internet and could not get a copy by ftp. So the question was, what would I say to them?

I could have said, "Find a friend who is on the Net and who will make a copy for you." Or I could have done what I did with the original PDP-10 Emacs: tell them, "Mail me a tape and a SASE, and I will mail it back with Emacs on it." But I had no job, and I was looking for ways to make money from free software. So I announced that I would mail a tape to whoever wanted one, for a fee of \$150. In this way, I started a free software distribution business, the precursor of the companies that today distribute entire Linux-based GNU systems.

That's it. In September 1983, the first draft of the *Manifesto* announced Richard's intent; just over a year later, his \$150 GNU Emacs initiated an innovative business model.

Thanks to Patrick Henry Winston, director of the MIT AI Lab from 1972-1997, Richard's resignation didn't have the expected consequences. Winston allowed rms to continue to have office and lab space at Tech Square. The AI Lab's computing facilities were also available for Richard's use.

In his *Defence of Poesy* (1595), Sir Philip Sidney contrasts the historian, who is obliged to be faithful to recorded events, to the poet, who is capable of depicting ideals, employing imaginative fictions. To Sidney, the poet's superiority lies with clarity of moral vision, whereas the details of events may result in the blurring of the historian's vision. Spenser (1552-1599), referring to himself as a "Poet historical," views historians as being forced to follow orderly chronology, where poets can move back and forth in time. All of this is to attempt to excuse my moving ahead to 1984, perhaps illustrating my drift between historian and "Poet historical."

Let me now move back in time and across the Atlantic.

¹ From Free Software, Free Society (FSF, 2002), p. 15.

试用水印

Excursus: Hardware

J. C. R. Licklider wrote that computers were communication devices, not calculating devices. Tomlinson's creation of email (1970) was a step demonstrating that. The continued expansion of the Internet provided the medium. UUCP and Usenet provided further impetus.

In January 1976, there were 63 hosts on the Internet. Five years later, there were just over 200. In August 1982, there were 235. For nearly 20 years thereafter, the number of Internet host sites doubled yearly: 562 in August 1983; 1024 in October 1984; 2308 in February 1986; 28,174 in December 1987; 727,000 in January 1992. But the slope of that curve has flattened. We no longer talk about hosts; we talk about users. But that's hard to estimate. How many people own a domain ... or several domains? And how many users are "on" some major domains? And what's the average? So it's largely guesswork. Around 2002, there appeared to be about 300 to 500 million users of the Internet. Right now, perhaps a sixth of mankind has access.

Two factors drove this (in my opinion): the development of the modem and the affordable personal computer. [This Excursus is not intended to be a complete history; my aim is to provide background so that subsequent chapters will be more intelligible.]

A modem (**mod**ulator-**dem**odulator) sends and receives data between two computers. The first commercial modem, the Bell 103, was built by AT&T in 1962. It had full duplex transmission, frequency-shift keying and operated at 300 bits per second. Things got a bit speedier over time: 1200 bps, 9600 bps, 14,400bps. Robert Lucky invented the "automatic adaptive equalizer" at Bell Labs in 1965. Brent Townshend, a Quebec inventor, created the 56K pulse-code-modulated modem in 1996. Things had gotten a lot better.

Personal workstations and personal computers are not new, either. I can recall the LINC of 1963, for example. But these machines, as well as others, though of historical importance, had little impact on where we are now.

But we do need to look back at the 1976 Apple 1, running at 1MHz with 8kB RAM (max. 32KB); the Apple II (1977), the IIe (1983), and the Macintosh SE (1987). And they had their beginnings in the Altair. (Aspects of all these were developed a decade earlier, many by Doug Engelbart.)

In 1973 and 1974 a small company in Albuquerque, NM, called MITS (Micro Instrumentation Telemetry Systems), which had been producing inexpensive calculators, was seeking a new product. Texas Instruments had just taken over its calculator market. Ed Roberts, with the help of Les Solomon, decided to build a computer kit. Assisted by two hardware engineers, William Yates and Jim Bybee, they developed the MITS Altair 8800, which was featured on the cover of the January 1975 Popular Electronics.¹ The magazine called the Altair the "World's First Minicomputer Kit to Rival Commercial Models." It shipped for \$400, but the purchaser had to assemble it, get it to work, and write the necessary software. It wasn't all fun, but it sold.

Among others, it sold to a Harvard freshman, Bill Gates, and his friend Paul Allen. They compiled a version of BASIC to run on the Altair. (Roberts offered Allen the post of Director of Software at MITS -- he was the only person in the department. Gates joined him upon leaving Harvard.)

The Apple][e ran on a SynerTek 6502 board. The SE was 8MHz, 256kB, ran on the Motorola 68000, and had a serial port into which a modem could be plugged. It sold for just under \$3000. (While I was Executive Director of USENIX in 1987, Telebit gave me a QBlazer and I connected the SE from home to the office using Red Rider.) Real power. And there were nearly 300 groups on Usenet.

But in 1975, IBM had tried to enter the "small machine" market with its 5100 -- a 50-pound "portable" computer, priced at \$9000 to \$20,000. It was a dismal flop. ("There is no reason anyone would want a computer in their home," Ken Olsen, founder of DEC, said in 1977.)

IBM then contemplated buying up Atari, but instead set up an "independent business unit" in Boca Raton, FL, to build the "Acorn."² The team of a dozen engineers was headed by William C. Lowe. They made a number of unusual decisions, the most notable of which was that the PC would have open, rather than proprietary architecture. They also decided to save time by purchasing an operating system, as well as making hardware components open to competitive bids. Finally, the PCs were to be sold through retail channels.

Released in August 1981, the original PC ran on a 4.77MHz Intel 8088. It had 16kB of memory, expandable to 256kB. It was priced at \$1565 and launched through a brilliant advertising campaign featuring a Charlie Chaplin look-alike. By 1985, IBM's sales had overtaken Apple's and IBM had 40% PC market share.

However, the same open architecture that made the IBM PC a success, led to its decline. Open architecture meant that others could clone it, and the first of these was Compaq, coming out with an 80386-based machine in 1986. Others followed. By 1995, IBM's market share had dropped to 7.3% and in 2003 it was 6%.

In 1981, the Osborne 1, the first true portable had been released. It had 64K RAM; twin 5.25", 91K drives; and ran on a Zilog Z80 at 4MHz. Though the company eventually failed, the machines were definite landmarks.³

Smaller size; lower price; access via the ubiquitous telephone. That did it.

¹ The Altair was named by Solomon's 12-year-old daughter, after a Star Trek episode.

² See Pugh, Building IBM, pp. 313f., and P.E. Ceruzzi, A History of Modern Computing, 2nd Ed. (MIT Press, 2003), pp. 268-273.

³ Adam Osborne (1939-2003) was a fascinating individual, but a poor businessman. See his Hypergrowth: The Rise and Fall of the Osborne Computer Corporation (1984).

Chapter 9. MINIX

Like Richard Stallman, Andy Tanenbaum was born in New York. After graduating from high school in White Plains (just north of New York City), he went to MIT, and subsequently received his doctorate from Berkeley. Since 1971 he has lived in the Netherlands (his wife is Dutch) where he's Professor of Computer Science at the Vrije Universiteit [Free University, VU].

After AT&T's 1979 announcement of the V7 licensing restrictions, precluding the use of the code in classrooms, Andy decided that the solution lay in his "helping himself."

Together with Sape Mullender, of the CWI [Center for Mathematics and Computer Science], Andy had originated the Amoeba project. Amoeba was one of the earliest attempts at a distributed operating system, contemporary with Roger Needham's work in Cambridge, and preceeding LOCUS, CHORUS, V, and Mach. [Tanenbaum & Mullender, in Operating Systems Review 13(2): 26-32 (1981)] That same year, Andy's valuable Computer Networks (Prentice-Hall) appeared.

And, while doing research, teaching classes, supervising graduate students, and writing, Andy worked on his own system. "It took about two years," he told me. He called it MINIX.

MINIX was a micro-kernel UNIX clone. While it emulated UNIX, it contained no AT&T code -- not in the kernel, not in the compiler, not in the utilities. It was 1986. The next year, 1987, Operating Systems: Design and Implementation came out, the book's title reflecting the VU course that Andy was teaching. At first, the code (v1.1) was only available on diskettes from Prentice-Hall. Soon it was available without the book.

MINIX 1.3 was on five 1.2M floppies and cost \$60; MINIX 1.5 (1991) came on 12 720K diskettes at \$169. 1.5 contained the accumulated bug fixes; it was V7 system call compatible; it would run on the IBM PC, PC AT, PC XT, PS/2, and 286/386 as well being "available" for the Atari ST, the Macintosh, the Amiga, and the Sun SparcStation 1, 1+, or IPC. 1.5 had a K&R compatible C compiler, a Bourne-like shell, five editors (ed, ex, an Emacs subset, a vi clone, and "a simple screen editor"), and a great deal of other goodies.

MINIX was intended as a teaching tool, and it was far from freely redistributable. It was under copyright by Prentice-Hall, but with rather liberal copy and revise/extend restrictions. However, it was certainly not free. As I understand it, the Prentice-Hall lawyers were to blame here, not Andy. MINIX now is free: the license has been redone and made retroactive. MINIX 2 is freely redistributable software.

Chapter 10. SUN and gcc

The company we think of as Sun Microsystems began with Andreas Bechtolsheim and some other graduate students at Stanford emulating Motorola's 68000 CPU cheaply. Stanford licensed a single board: the Stanford University Network board -- SUN.

Soon companies began licensing the board: Codata, Fortune, Dual, Cyb, Lucasfilm, and others. Machines began appearing. Each was "just another workstation" -- JAWS.

The first UNIX workstation had been the Z8000 ONYX, hardly a VAX on a chip. John Bass demo-ed it at the USENIX Conference in Boulder, CO, 29 January to 1 February 1980.

The system we took to Boulder was on three boards about 15 by 22 inches [Bass told me]. Its performance and architecture was more like a PDP-11/45 or 11/70 ... segmented memory, no paging. ... That aside, the ONYX was the first table-top system designed to run UNIX. With eight serial ports [users] and at under \$25k, it made a great short-term alternative to PDP-11 UNIX systems.

But then came those JAWS -- some of them at under \$10k. And all of them ran AT&T's System III or 4.1BSD.

System III was AT&T's commercial variety of V7. Though its official release date was 31 October 1981, it reached some of the purchasers earlier and the general public in 1982.

Berkeley¹ issued 4BSD in October 1980. It included a faster file system, job control, auto reboot, delivermail (soon to be renamed sendmail), and Franz Lisp. In June 1981, 4.1BSD, which had autoconfiguration and some minor improvements, was issued. Just why it was 4.1 (leading to 4.1a, 4.1b, 4.1c, 4.2, 4.3, and 4.4) is another of the silly consequences that licensing restrictions forced upon developers.

Bill Joy made a 10-day visit to DEC in early 1981, working with Armando Stettner (who had gone to DEC from Bell Labs as "a sort of UNIX ambassador") on porting 4BSD to the VAX.

We made a dump tape [Armando told me] and Bill packed up and went back to California. [Bill] Shannon and I took the disk pack and brought it up on decvax -- a 780, our main system. Bill Joy called a couple of days later and said, "Hey, there's going to be a lot of hassle with the license if we do another release. So why don't we call it 4.1BSD?"

4.1a, 4.1b, and 4.1c were all "test releases." 4.1a included TCP/IP and the socket interface and was sent to a number of ARPANET sites. 4.1b included the new "fast file system" and new networking code. It was only used on the Berkeley campus once, in a graduate OS class. 4.1c was almost 4.2BSD, lacking only the new signal facility. It was sent to about 100 sites.

System III was distributed by AT&T without source. It was the first version of UNIX to be issued that way. But those customers who were unhappy merely obtained V7 from Western Electric or used the Berkeley editions...which came with source.

The June 1982 issue of ;login: (the newsletter of the USENIX Association) carried an article headlined:

**Interesting Developments:
Bill Joy of UCB moving to Sun Microsystems**

Bill took a tape of 4.1cBSD with him. It became the basis for SunOS.² 4.2BSD became DEC's Ultrix.

For half-a-dozen years, improvements in BSD were incorporated into subsequent versions of SunOS. But in 1988 AT&T announced a major investment in Sun Microsystems and thereby startled the UNIX community. (Ostensibly, the purpose was to merge the AT&T and Berkeley strains of UNIX. Most saw a far darker purpose.) I'll return to this in the next chapter (11).

gcc

The GNU C Compiler (gcc) was Richard Stallman's first free software "hit." There were many C compilers available (at least four or five of them written by Whitesmiths, P.J. Plauger's software company,³ but they were all proprietary. Stallman's was unencumbered -- and it worked well. (gcc now stands for GNU Compiler Collection, and comprises compilers for C, C++, Objective-C, Fortran, Java, and Ada, and a large number of libraries; a two-CD set still costs only \$45.)

Remember, the USENIX community had been issuing free distribution tapes for a decade, and Rich Morin, one of the founders of the Sun User Group, had emulated this practice. When Stallman's compiler came out in 1987, Morin recognized that the hassles he had encountered in getting permissions from contributors were resolved by the GPL. And he recognized that the GPL made what he was engaged in a possible business. Morin's "service" became Prime Time Freeware.

In 1990 I became Executive Director of the Sun User Group. That December I headed for San Jose for SUG's Eighth Annual Conference and Exhibit. It was a very tense meeting. In the first few hours I was at the hotel and the Convention Center, I became aware of the fact that there were two separate (though overlapping) sets of irate users.

One of these was made up of those who had bought a Sun 386i, Sun's sole venture into the Intel world. Though it was a business failure, the decision to end support for the machine was not greeted with huzzahs. (At the "Meet the Executives" session, Ed Zander explained that Sun wasn't "abandoning" the users and that an external firm would support the 386i for (as I recall it) "up to five years." The faithful were not appeased.

The second group was irate because Sun had "unbundled" its software. That is, rather than getting all of Sun's developer tools together, they had to be purchased separately. And of course, they cost more this way.

But wait. Why purchase the C compiler from Sun, when you could get a better one for less money from the FSF? That's what a large number of Sun's users asked themselves. And the net result was a real jump in CD sales at the FSF. (Several years later, when I organized the Freely Redistributable Software Conference [February 1996] and then was Vice President of the FSF, I realized more fully just how much Sun had benefited the FSF. I'm certain this was not a foreseen consequence.)

The GNU C compiler was not the first piece of freely redistributable software, but it was the first

widely circulated product of Stallman's project.

¹ I owe most of the BSD version chronology to Keith Bostic.

² Sun was always a UNIX company, but the Sun-1 ran on UniSoft's v7.

³ See my interview with Plauger, Quarter Century of UNIX, pp. 174-176.

试用水印

Excursus: UUNET

By the mid-1980s there were several commercial networks in operation. But they were limited in service and, generally quite high in price. None was what we would think of as an ISP.

In the Autumn of 1985, Rick Adams (then at "seismo"), approached Debbie Scherrer, vice-president of USENIX, with a plan for a centralized site, accessed via Tymnet by subscribers, supplying Usenet access.¹ In an email dated December 6, 1985, Debbie expressed interest in this.

Rick attended the October 1986 Board meeting in Monterey, CA, where reaction was mixed, one director asking why folks would pay for access that could be obtained free. But the Board agreed to entertain a proposal. Rick brought a brief plan to the January 1987 (Washington, DC) meeting.

A majority of the USENIX Board liked the plan, but it really wasn't much of a "business plan," and Rick was asked to fill out the plan, with the participation of Board members John Quarterman and Wally Wedel, and return.

By late March 1987 (in New Orleans), Rick was back with a full plan and the Board approved it enthusiastically. I was authorized to spend up to \$35,000 for an experimental period.²

UUNET was born. "As people moved from universities and corporations where they had email and Usenet access to jobs where they had no access," Rick told me, "a need developed for a service that could provide email and Usenet access. UUNET was created in response to that need."

When the word got out, the demand far exceeded expectations. For example, Rick and Mike O'Dell had forecast 50 subscribers by the "end of summer." They topped 50 by mid-June 1987. Five years later, they had several thousand customers. UUNET reincorporated as a for-profit and then had its IPO. There is a long and interesting history; but this is not the place for it.

The important thing is that UUNET initiated commercial delivery of USENET and the Internet.

¹ Tymnet was an early proprietary network, first set up parallel to the ARPAnet by Tymshare, Inc., using Interdata 7/32s as nodes. In 1979, Tymnet was spun off by Tymshare and bought up by McDonnell-Douglas in 1984. In 1989, BT North America bought Tymnet from McDonnell-Douglas. In 1993, MCI bought Tymnet from BT North America for stock. Tymnet survived MCI's acquisition by WorldCom, but was finally closed down in 2004.

² At that time I was Executive Director of the USENIX Association. I handled the UUNET application for not-for-profit status, the liaison with the lawyer, and signed all the checks for about 14 months. Over that period, the USENIX Board increased its "advance" to over \$100,000. In only a few years, UUNET repaid all its debt.

Chapter 11. OSF and UNIX International

In 1987, AT&T purchased a sizable percentage of Sun Microsystems and there was a joint announcement that they would be involved in a grand merger of System V and BSD. Moreover, AT&T announced that Sun would receive "preferential treatment" as AT&T/USL [UNIX Systems Laboratories] developed new software. Sun announced that its next operating system would not be a further extension of SunOS.

The scientific community felt that Sun was turning its back on them. The manufacturers felt that the special relationship would mean that Sun would get the jump on them. Great cries of praise did not go up from the computer manufacturers.

"When Sun and AT&T announced the alliance," Armando Stettner told me, "we at Digital were concerned that AT&T was no longer the benign, benevolent progenitor of UNIX . . . Sun was everyone's most aggressive competitor. We saw Sun's systems were direct replacements for the VAX. Just think: the alliance combined our most aggressive and innovative competitor with the sole source of the system software -- the balance shifted."

On 7 January 1988 there was a meeting at DEC's Western Offices in Palo Alto, CA. There were participants from Apollo, DEC, Gould Electronics, Hewlett-Packard, Honeywell-Bull, InfoCorp, MIPS, NCR, Silicon Graphics, UniSoft, and Unisys. The group (called "the Hamilton Group," because DEC's building was at 100 Hamilton Avenue) sent a telegram to James E. Olson, CEO of AT&T, requesting a meeting "during the week of January 25" with Vittorio Cassoni (Senior VP of AT&T's Data Systems Division).

Larry Lytel of HP called a preliminary meeting of the group at the JFK Marriott for the evening of 27 January. The meeting with Cassoni took place the next day. There was a follow-up meeting of the Hamilton Group in Dallas on 9 February. The meeting with Cassoni had had no positive effect where the Group was concerned. (It's not clear whether AT&T took the Group seriously. It appears that Cassoni just thought of this as jockeying for commercial advantage.) In March, IBM was invited to join.

Apollo, DEC, HP, IBM, Bull, Nixdorf, and Siemens held semi-secret meetings and in May 1988, the formation of the Open Software Foundation was announced. (The Wall Street Journal for May 18 noted that no one present at the launch of OSF could recall ever seeing Ken Olsen sharing a stage with an IBM chief executive.)

Ken Thompson was in Australia at the time. When Ritchie told him what had transpired, he said: "Just think, IBM and DEC in one room and we did it!"

The seven companies listed above were joining hands to produce a new UNIX kernel and a new user interface. Their "temporary" headquarters would be in Lawrence, MA. A delegation of executives (loaned to OSF from their various corporations) attended the USENIX Conference in San Francisco in June.

It didn't take long for AT&T, Sun and their coterie to form a counter-consortium, UNIX International, dedicated to the marketing of SVR4.

OSF quickly named its executive team, including David Tory (Computer Associates) as President; and Roger Gourd (DEC), Ira Goldstein (HP), and Alex McKenzie (IBM) among the Vice Presidents.

UI appointed Peter Cunningham (ICL) as President.

By the end of 1989, Gourd's engineering team had come out with a new user interface, Motif, which was well-received, and Goldstein's research team had chosen Mach as the underlying kernel for the OS. OSF also increased its number of sponsors, adding Hitachi and Philips. However, as HP swallowed up Apollo and Siemens bought Nixdorf, at year end there were still seven sponsors.

Both OSF and UI ran membership drives and gave out pens and badges and stickers. Each ended up with about 200 members.

In 1991-92 the worldwide economy worsened. Bull, DEC, IBM, and the computer side of Siemens all lost money. AT&T resold its share of Sun. The fierce mudslinging appeared to be over. (At one point there was even a rumor of OSF and UI merging, for the good of UNIX.)

It hardly seemed to matter: Sun had adopted Motif; in 1993 USL sold UNIX to Novell, whereupon UI disbanded; OSF abandoned several of its previously announced products (shrink-wrapped software and the distributed management environment); Bull, Philips and Siemens withdrew from sponsorship of OSF.

Armando remarked to me: "It's not clear whether there's any purpose to OSF any more."

In 1984 a group of UNIX vendors had formed a consortium, X/Open, to sponsor standards. It was incorporated in 1987 and based in London. In 1996 OSF merged with X/Open to become The Open Group.

X/Open owned the UNIX trademark, which passed on to The Open Group. The Group also took on Motif and the Common Desktop Environment (CDE).

But the Open Group maintained its concern with standards, and is the sponsor of the Single UNIX Specification. It has also taken on sponsorship of other standards including CORBA and the Linux Standard Base.

Chapter 12. GNU, the GPL and Cygnus

In September 1983, there was an "announcement." Then, in 1984, Richard Stallman wrote *The GNU Manifesto*.¹ In my opinion, it marks the true beginning of the GNU Project.

For several years, the Manifesto was updated repeatedly in minor ways, but it remained the primary document through which rms would "ask for participation and support."

Stallman and his small group of programmers had been working on rms' stated goal of free versions of all the UNIX applications and tools. By mid-1984, there were the following:

an Emacs text editor with Lisp for writing editor commands, a source-level debugger, a yacc-compatible parser generator, a linker, and around 35 utilities. A shell (command interpreter) is nearly completed. A new portable optimizing C compiler has compiled itself and may be released this year. An initial kernel exists but many more features are needed to emulate UNIX. . . We will use TeX as our text formatter, but an nroff is being worked on. . . .

It was pretty impressive.

I'm not going to quote much more of the Manifesto, but there is one part, "Why I must write GNU," that has been my "favorite" for twenty years:

I consider that the golden rule requires that if I like a program I must share it with other people who like it. Software sellers want to divide the users and conquer them, making each user agree not to share with others. I refuse to break solidarity with other users in this way. I cannot in good conscience sign a nondisclosure agreement or a software license agreement. For years I worked within the Artificial Intelligence Lab to resist such tendencies and other inhospitalities, but eventually they had gone too far. I could not remain in an institution where such things are done for me against my will.

So that I can continue to use computers without dishonor, I have decided to put together a sufficient body of free software so that I will be able to get along without any software that is not free. I have resigned from the AI lab to deny MIT any legal excuse to prevent me from giving GNU away.

Take this seriously. Just over twenty years ago, rms was talking about a political and social movement. He was talking about "solidarity" and "conscience." Stallman was more interested in "freedom than in having a better program."²

Hey! Not a bad outlook! I really like, say, the most recent Tom Clancy, so I lend it to a friend. This week's cartoon in *The Economist* strikes me, so I clip it and send it to a pal. But, if it's software? Fuggedaboutit!

Stallman also produced the GPL -- the GNU Public License, now the GNU General Public License.³

The GPL grew out of a real need for legal documentation.

James Gosling, then a graduate student at Carnegie Mellon, wrote a C-based version of Emacs which used a simplified Lisp, MOCKLISP. In order to construct GNU Emacs on Lisp, rms freely borrowed Gosling's innovations. (Stallman had been told by other CMU developers that Gosling had assured their work on GOSMACS and the Lisp interpreter would remain available. But Gosling put GOSMACS under copyright and sold the rights to UniPress. UniPress, in turn, threatened to sue rms.)

As Sam Williams put it:

Once again, Stallman faced the prospect of building from the ground up.

In the course of reverse-engineering Gosling's interpreter, Stallman would create a fully functional Lisp interpreter, rendering the need for Gosling's original interpreter moot. Nevertheless, the notion of developers selling off software rights -- indeed, the very notion of developers having software rights to sell in the first place -- rankled Stallman.⁴

GNU Emacs was released in 1985, but rms had come to realize just how important it would be for GNU software to have a "legal foundation" to stand upon. The first version of the GPL was the direct result. Richard had realized that one needed to actually bestow an absolute right on users. He had spoken with Mark Fischer, a Boston IP lawyer, and to Jerry Cohen, another lawyer, but wrote his own license. Only a few years later, GPL version 2 was released. This was in 1991. Just about 15 years later, we still use that version, though version 3 is in the offing.

If you're curious, look at the copyright notice in the README file of trn [= threaded read news], written by Larry Wall, prior to the creation of Perl. It says:

Copyright (c) 1985, Larry Wall

You may copy the trn kit in whole or in part as long as you don't
try to make money off it, or pretend that you wrote it.

Yep.

In 1991 there was no Web.

In 1991 there was no Linux.

In 1991 KDE, Gnome, Apache, Netscape, hadn't even been thought of.

In 1991 we were still waiting for the Hurd.

But John Gilmore, employee number 5 at Sun Microsystems, was aware of the importance of GNU and of the GPL. Gilmore was part of the Usenet community. He was more than just a reader of net news. In November 1986, Gilmore suggested that rms "remove 'EMACS' from the license and replace it with 'SOFTWARE' or something." Version 1.0 of the GPL was officially released in 1989, a year after the release of the GNU debugger, which carried the 1985 draft (emended innumerable times).

Another individual taken by the GNU philosophy was Michael Tiemann. In an essay published in 1999,

Tiemann looked back at the Manifesto: "It read like a socialist polemic, but I saw something different: I saw a business plan in disguise."⁵ He dropped out of the Ph.D. program at Stanford to pursue that plan.

Tiemann wrote the GNU C++ compiler and the first native-code C++ compiler and debugger. He is now Vice President for Open Source Affairs at RedHat.

In 1989, Gilmore, Tiemann and David Henkel-Wallace co-founded Cygnus Solutions, the first "open source" business. (I put "open source" in quotation marks to differentiate it from free source. Stallman notes that "open source" stresses the technical side of the software, excellence through code sharing, whereas "free software" emphasizes the moral and ethical, technical excellence being a desirable byproduct.) Gilmore ceased working at Cygnus in 1995 and stepped down from its Board in 1997.

Cygnus was founded on the theory that "There is great value in having good people working on software whose precedents will set the standards of tomorrow. We believed at the beginning that people would understand this value proposition, and would value the opportunity to pay us to create high-quality, open-source programs that would become the de facto standard of the software world."⁶

Cygnus began by selling the GNU compiler and debugger as shrink-wrapped software. Gilmore sent out email telling folks that he'd be the debugger maintainer and integrator. Gilmore and Gumby [D.V. Henkel-Wallace] hacked and Tiemann sold contracts.

Cygnus became a success: it demonstrated that money could be made through service, packaging and distributing source that was otherwise free.

All the vision of Stallman and the hard work of Gilmore, Gumby and Tiemann bore fruit in Cygnus' GNUPro Developers Kit, which contained:

- GCC
- G++
- GDB
- GAS [GNU Assembler]
- LD [GNU Linker]
- Cygwin [UNIX Environment for Windows]
- Insight [GUI for GDB]
- Source-Navigator

By the early 1990s, the world was beginning to change.

¹ GNU Manifesto.

² Interview with Michael Gross, "early 1999"; in *The More Things Change* (Harper Collins, 2000).

³ The GPL

⁴ Sam Williams, *Free as in Freedom* (O'Reilly, 2002), pp. 104f.

⁵ In *Open Sources* (O'Reilly, 1999), p. 139.

⁶ Tiemann in *Open Sources*.

Chapter 13. USL v The Regents of the University of California

From late 1986 on, Keith Bostic would stand up at each USENIX Conference and announce the progress of his -- the CSRG's -- project: about 35% of the code was AT&T license free; about 55%; about 77%... The progress may have seemed slow at times, but there was always some progress.

And there were always loud cheers and resounding applause.

AT&T's lawyers had started off on the wrong foot in the mid-1970s; the fee structure made AT&T-license-free UNIX a financial necessity. Among others, John Gilmore nudged the CSRG to produce a license-free version. After all, it was clear that AT&T hadn't objected to Minix, which was a UNIX-like system with no AT&T code.

In November 1988, at the BSD Workshop in Berkeley, Keith, Mike Karels and Kirk McKusick announce the completion and availability of BSD Networking Release 1.

NET 1 was a subset of the then-current Berkeley system. It was quite similar to 4.3-Tahoe, including source code and documentation for the networking portions of the kernel, the C library and utility programs. It was available without evidence of any prior license (AT&T or Berkeley), and was (re)distributed via anonymous FTP. The source carried a Berkeley copyright notice and a legend that allowed redistribution with attribution.

(The Berkeley license was, and still is, different from the GPL. Keith and rms had debated the various aspects of the licenses repeatedly, without convergence. I will discuss this later.)

In June 1991, at the USENIX Conference in Nashville, BSD Networking Release 2 was available. NET 2 contained far more than just networking code and, like NET 1, was available with no prior license. The new features included a new virtual memory system (derived from Carnegie-Mellon's Mach system, which had been ported at the University of Utah) and a port to the Intel 386/486.

It was nearly the end of the line for the CSRG. Karels left Berkeley in 1992; Bostic and McKusick, in June 1993. But, NET 2 was a US-Russia collaboration, with contributions by Bill Jolitz, Donn Seeley, Trent Hein, Vadim Antonov, Mike Karels, Igor Belchinsky, Pace Willisson, Jeff Polk, and Tony Sanders. It was turned into a commercial product, known as BSDI (= Berkeley Software Design, Inc.) and was complete by the end of 1991 and released to the public on April 10, 1993 as 1.0, the long delay being the consequence of UNIX Systems Laboratories (USL) filing suit to prevent BSDI from shipping its product.

USL filed for an injunction barring distribution of "BSD/386, pending trial, arguing that BSD/386 infringed USL's copyright in its 32V software and misappropriated USL trade secrets."¹ The Court denied USL's request for a preliminary injunction on March 3, 1993, ruling that USL had "failed to show a likelihood of success on either its copyright claim or its trade secret claim."

On March 30, 1993, Judge Dickinson Debevoise of the US District Court of New Jersey reaffirmed his

denial of USL's motion for a preliminary injunction against BSDI. The Court found that the 32V source code had been distributed without a copyright notice. The Court rejected USL's argument that the publication of 32V was sufficiently limited to avoid a forfeiture, and thus found that USL had failed to establish that BSD/386 contained any USL trade secrets.

USL subsequently filed a motion for reconsideration, asking the District Court to hold a new hearing on whether USL had published 32V without a copyright notice. USL argued that the Court's prior ruling was based on an incorrect finding as to the number of copies distributed. (USL's motion for reconsideration did not challenge the Court's ruling that USL had failed to establish trade secret misappropriation.)

The Court denied USL's motion for reconsideration. Although the Court amended its prior factual finding as to the number of copies distributed, it found that the number was not critical to its ruling on the issue of publication without notice.

It was just under 20 years since Ken had delivered that first UNIX paper at SOSP and began receiving requests for the software.

It was 15 years since *UNIX NEWS* became *;login:* and the "UNIX Users Group" turned into USENIX.

And, through all of this, Western Electric, AT&T, and now USL had learned nothing about the nature of the user community.

What BSDI (and others, like mt Xinu) were trying to do was ensure the continued development, growth and use of the UNIX operating system. The suit by USL was an attempt to protect something of value. But the value had been discovered too late. Maybe Ritchie and Thompson had handled the system carelessly in the mid-1970s; maybe BTL employees inadvertently gave UNIX to the public without any significant restriction. But the users, like O'Dell and Kolstad, Coulouris and Lions, saw the value.

BSDI had distributed pre-production releases of BSD/386 (Beta version). It now began international distribution. Full source was priced at \$1000. (In the January/February 1994 *;login:*, Lou Katz wrote: "It works! It works!").

Because of the way the licenses had been granted to the University of California, the Regents of the University had been included in USL's suit. In June 1993, the Regents struck back, filing suit against USL for not complying with agreements made with the CSRG.

In the meantime, Novell acquired USL. A decade later, it is still unclear exactly what USL believed it was selling, or what Novell thought it was buying.

To take but one example, the "PROPRIETARY NOTICE" on the System V `errno.h` file reads:

This source code is unpublished proprietary information constituting, or derived under license from AT&T's UNIX (r) System V. In addition, portions of such source code were derived from Berkeley 4.3 BSD under license from the Regents of the University of California.

USL had filed suit in New Jersey, the home of AT&T and Bell Labs (and USL), clearly thinking that this would yield an advantage in their suit. When the USL v BSDI and Regents suit was dismissed, the

Regents filed suit in California.

On Friday, February 4, 1994, Novell and the University of California agree to drop all relevant suits and countersuits. BSDI immediately announced the availability of a release based on "4.4BSD-Lite."

"We are delighted to settle with USL so that we can devote our full efforts to creating products and serving our customers," Rob Kolstad, president of BSDI, said to me. (I will not go into the corporate history of BSDI here.)

There is no question that The Open Group "owns" the UNIX trademark. There are many questions as to whether any one entity "owns" UNIX. In 1995, Novell executed an "Asset Purchase Agreement" with the Santa Cruz Operation, though exactly what was purchased is unclear a decade later.

While both The SCO Group (a successor to the Santa Cruz Operation) and Novell claim rights to the UNIX source code, many elements of that code carry copyright notices from the University of California and other companies and individuals.

All of this is important because of the various suits filed by The SCO Group and the ownership claims made. It is clear to me that while AT&T owned a good portion of several versions of System V, as a result of the incorporation of code from the CSRG and other sources, they never had full rights to all the source of any of the versions. Exactly what was conveyed to Novell, from Novell to the Santa Cruz Operation, and from the latter to The SCO Group, may be adjudicated in one of the current suits.

¹ The story of 32V is in Chapter 6.

Chapter 14. BTL after UNIX: Plan 9 and Inferno

In July 1990, I flew from Boston to London for the UKUUG Conference. (I was to give a talk on UNIX standards and specifications.) But there were three talks on the program that blew me away.

They concerned "Plan 9" a new OS being worked on at Bell Labs. It was named Plan 9 from Bell Labs after "Plan 9 from Outer Space," perhaps the worst science fiction movie ever filmed.

Plan 9 is a UNIX clone. But it presents a consistent interface which is easy to use. I am not going to go into it at any length. But, it was the successor to UNIX, which, Rob Pike said, was dead: "It's been dead for so long it doesn't even stink any more."¹

Rob delivered the keynote address at the UKUUG: "Plan 9 from Bell Labs." He's now at Google.

Dave Presotto then spoke about "Multiprocessor Streams for Plan 9." He's at Google, too.

Tom Duff talked about "Rc -- A Shell for Plan 9 and UNIX Systems." Tom's now at Pixar, the proud owner of parts of several Oscars.

Fifteen years later, what had been the UNIX group (1127) has been dispersed. In addition to Rob, Dave and Tom,

- Ken Thompson retired to California;
- Brian Kernighan is a Professor at Princeton;
- Phil Winterbottom is CTO at Entrisphere;
- Gerard Holzmann is at NASA/JPL Laboratory for Reliable Software;
- Bob Flandrena is at Morgan Stanley;
- Sean Dorward is at Google;

Dennis Ritchie and Howard Trickey remain at Lucent/BTL.

But, before it disappeared, the "1127 group" made yet another contribution to OS development: Inferno.

Inferno is a compact OS designed for building "cross-platform distributed systems." It can run on top of an existing OS, or as a stand-alone. The nomenclature owes much to Dave Presotto, who founded it firmly in Dante. The company marketing Inferno is Vita Nuova; the communications protocol is Styx; applications are written in type-safe Limbo, which has C-like syntax.

The 4th edition of Inferno was released in 2005 as free software, but under a mixture of licenses.

¹In the July 2005 issue of IEEE Spectrum, there's an article "The End of AT&T" with the blurb:

Once the world's largest company, Ma Bell will soon vanish. But its innovations -- from the transistor to communications satellites to laser cooling-live on. By Michael Riordan
--

Note what's important. CS isn't.

Chapter 15. Commercial UNIXes to BSDI

In the 15 years following the release of V6 (April 1976), Berkeley was not the only place where versions and clones of UNIX sprouted. While I doubt whether I can even enumerate all of them, the following will give an image of the geography of the field. To me, the most significant UNIX releases were:

- November 1976: Interactive Systems IS/1
- March 1978: 1BSD
- November 1978: Cromemco CROMIX
- November 1978: Technical Systems Consultants UniFLEX
- December 1978: V7
- April 1979: AT&T UNIX 32/V
- May 1979: 2BSD
- February 1980: 3BSD
- August 1980: MicroSoft XENIX OS
- September 1980: 4BSD
- 1980: Idris
- May 1981: 4.1BSD
- October 1981: System III
- November 1981: IBM CPIX
- November 1981: ULTRIX
- August 1982: 4.2BSD
- December 1982: AT&T System V
- March 1983: XENIX 3.0
- May 1983: Mark Williams Coherent
- September 1983: SCO Xenix
- March 1984: AT&T SVR2
- November 1984: SCO Xenix System V/286
- January 1985: BTL Eighth Edition
- November 1985: Apple A/UX
- November 1985: AT&T SVR3.0
- November 1985: AIX/RT 2
- November 1985: HP-UX 1.0
- May 1986: 4.3BSD
- August 1986: BTL Ninth Edition
- November 1986: Minix 1.0
- November 1986: AT&T SVR3.2
- November 1986: HP-UX 1.2
- September 1987: SCO Xenix System V/386
- November 1987: NeXTStep
- November 1987: Acorn RISC Unix
- January 1988: AIX 1.0
- May 1988: 4.3BSD-Tahoe

- October 1988: BSD Net/1
- September 1989: BTL Tenth Edition
- November 1989: Coherent 3.0
- November 1989: AT&T SVR4
- February 1990: SunOS 4.1
- October 1990: Solaris 1
- November 1990: Novell UnixWare Personal Edition 1.1
- January 1991: Trusted Xenix 2.0
- May 1991: BSD Net/2

I've mentioned Minix and the AT&T, BTL and BSD releases earlier. But several of the others are worth devoting a vignette to them.

Interactive Systems

Interactive was founded by Peter Weiner in 1977. (Weiner had been Brian Kernighan's Ph.D. advisor at Princeton.) In 1978, Heinz Lycklama joined him in Santa Monica. Lycklama had just written LSX, a version of V6 UNIX for the LSI-11 microprocessor. Interactive's product was called IS/1 and ran on most PDP-11s. Interactive's UNIX was an important product for nearly a decade. In 1985, Interactive's IN/ix became the basis for AIX (announced 21 January 1986). Some of the later modifications to AIX were developed by Interactive under contract to IBM.

Cromix

Cromix was a proprietary UNIX clone of CROMEMCO. The CROMEMCO 100 ran on a Xilog 80 and had 512K of RAM, 50M of hard disk, and an XPU processor, enabling 32-bit processing. Founded in the early 1970s by Roger Melen and Harry Garland, Stanford students who lived in **CRO**thers **MEM**orial Hall, it was incorporated in 1976. In 1985, it was bought up by Dynatech, and disappeared. But Cromix was the first UNIX clone. The CROMEMCO 100 and 300 ran both Cromix and System V. The 300 ran a 68000 timesliced with a Z80 coprocessor to enable multiuser CP/M WordStar.

TSC UniFLEX

Technical Systems Consultants wrote a drive for the then-new 5.25" drives in 1976: DOS MiniFLEX. It was superseded by FLEX for the 6800 a few months later. FLEX was adopted by virtually all of the 68xx SS-50-based computers (even the Tandy Color Computer and the UK Dragon). TSC now turned to producing a UNIX-like multi-user for the 6809: UniFLEX. It was a failure because of the introduction of 16-bit processors and the PC.

Microsoft XENIX

Microsoft licensed 7th Edition from AT&T in 1979. On 25 August 1980 they announced that XENIX would be available for 16-bit processors (Microsoft couldn't license the name, "UNIX"). XENIX wasn't identical to 7th Edition because Microsoft incorporated several features from BSD.

Microsoft didn't sell XENIX: it was licensed to manufacturers who were responsible for the porting. The first ports were to the Zilog Z8001, a 16-bit processor. Altos shipped one in early 1982. Tandy shipped one for 68000 systems in January 1983 and SCO released their port to the 8086 in September 1983. The license had been for V7, XENIX was based on System III.

XENIX 2.0 (1985) was based on System V, and added support for 80286. However, Microsoft apparently lost interest in XENIX after signing an agreement with IBM to develop OS/2. In 1987 Microsoft transferred ownership of XENIX to SCO in exchange for 25% of the company. That same year, SCO ported Xenix to the 386 and Xenix 2.3.1 supported SCSI and TCP/IP.

Xenix became SCO UNIX in 1989.

Idris

P.J. [Bill] Plauger received his Ph.D. in Nuclear Physics from Michigan State in 1969. From 1969 to 1975 he was a Member of Technical Staff at Bell Labs. Together with Brian Kernighan, he wrote *Elements of Programming Style* (1974) and *Software Tools* (1976). He also writes science fiction, and won the 1975 John W. Campbell Award as the best new SF writer of 1975.

It was while writing *Software Tools* that Plauger left the Labs. He told me this:

I ended up leaving the Labs. I felt I didn't have a future there and that I'd better move on before they told me to move on. And I was able to get a job at Yourdon . . .

After a few years of traveling all over the world lecturing, I felt that I wanted to get back to programming. Ed [Yourdon] had an opportunity to get a contract to do a commercial C compiler, and I talked him into doing it. I worked around the clock for a week. . . .

Plauger went on to form a three-man company, Whitesmiths.

I think we started on August 1st, '78. We were going to sit down and write a C compiler from scratch -- my third C compiler, I guess. I paid a lot of attention to not having any notes from my Lab days or my Yourdon days. . . I wrote like a fiend and by the end of November, we had a compiler.

Whitesmiths' first compiler was for Fisher and Porter in Philadelphia. It was for the PDP-11. "We gave them an 8080 compiler by the middle of '79; a VAX compiler by the end of that year; and we gave them a 68000 compiler in the middle of 1980," he said. "And we were doing Idris at the same time."

Idris was a UNIX-like multi-user multi-tasking operating system, written by Plauger and M. S. Krieger. Originally, Idris ran only on the PDP-11. But it was soon ported to the VAX, the 68000 and the 8086. In 1986, Atari hired Computer Tools International to port Idris to the Atari ST. Whitesmiths was sold to Intermetrics in 1988.

Mark Williams Coherent

The Coherent Operating System from Mark Williams was a UNIX-like OS for PCs. It was introduced in 1983. As I knew that several former University of Waterloo students had worked on it, I asked Tom Duff. Here it is, in his own words:

I was at Mark Williams from roughly August 1 to October 31 of 1980. After leaving the NYIT Graphics Lab, I had 6 months free (later reduced to 3 months) before I was scheduled to start at Lucasfilm. Mark Williams CEO Bob Swartz heard that I was available and asked if I'd like to

work in Chicago for a while.

When I arrived, they had a working C compiler, assembler and loader and a version of ed, written by Dave Conroy, hosted on RSX.

Randall Howard was doing most of the kernel work. Johann George, David Levine and Bob Welland were also there, but I'm not sure what they were working on -- Johann was probably doing kernel stuff.

Dave Conroy, Randall, Johann and I were all friends at Waterloo in the early '70s.

This was an amazing crew: Dave Conroy most recently was in charge of engineering the Mac Mini, Randall founded MKS, Johann founded Sourcelight Technologies (Randall and Johann are both semi-retired VCs now), David Levine wrote a legendary early video game called Ballblaze and Bob worked on the design of a bunch of important Amiga hardware.

When I arrived, it was pretty clear that the kernel was pretty much taken care of (though it wouldn't be running well enough for daily use until after I'd left), but nobody was working on user-space stuff. So I opened the 6th edition manual to page one and started implementing commands. In the three months that I was there, I think I did A through M, (As I remember, I started at make, then jumped back to ar and just plowed through. I remember make, diff and dc being a lot of fun.

And, I did units, because the library research required to dig up the more obscure quantities seemed interesting.

While I was there, Ciaran O'Donnell (another friend from Waterloo) visited for two weeks during which he wrote, in a feat of coding acrobatics such as I have never seen before or since, a complete, functioning YACC clone, working just from Aho and Johnson's 1974 Computing Surveys paper.

Coherent eventually ran on most 286, 386 and 486 boxes. It actually had support for X11.

The Mark Williams Company went bankrupt in 1995.

A/UX

A/UX was Apple's entry to the world of UNIX in 1988. It was based on SVR2.2 with element of SVR3 and SVR4 as well as 4.2BSD and 4.3BSD. It is POSIX and SVID compliant. From A/UX v2 on, it included TCP/IP. The last version (3.1.1) was released in 1995.

NeXTSTEP

A UNIX-like kernel based on Mach (CMU) with many BSD features and display PostScript with a windowing engine lay at the heart of NeXTSTEP. Previewed several times beginning in 1986, it was released on 18 September 1989. The last release (3.3) came out in early 1995.

There are, of course, many other UNIX-like things one could talk about, but I never found Trusted Xenix nor the RISC version nor Compaq's NonStop-UX very interesting.

Chapter 16. The Hurd and BSDI

The Hurd

Richard Stallman had long wanted a GNU replacement for the UNIX kernel. A first pass, Trix, barely got going in the late 1980s. This changed, however, when Thomas Bushnell came to Boston and joined the GNU effort.

Thomas was born in Albuquerque, NM. He attended Carnegie-Mellon University for a year (1985-86), the University of New Mexico for nearly two years, worked, enrolled at the University of Massachusetts Boston, and received a B.A. summa cum laude in 1999 in philosophy and classics. Thomas is a brother in the Brotherhood of St. Gregory, an Episcopal order. He received his M.A. in Philosophy from the University of California, Irvine, in 2003 and is currently a Ph.D. candidate there.

Thomas worked as an Assistant Systems Administrator at UNM from 1986-89 and for the FSF from 1990-1998. He told me:

I wrote a BASIC interpreter as a demonstration that I could code before I was hired. My interpreter had a feature that would let you dynamically load math functions out of the C library -- before shared libraries existed.

I worked on GNU tar as well, before my main work was the Hurd.

The GNU Hurd is the GNU project's replacement for the UNIX kernel. The Hurd is a collection of servers that run on the Mach microkernel to implement file systems, network protocols, file access control, and other features that are implemented by the UNIX kernel or similar kernels (such as Linux). Thomas told me:

RMS was a very strong believer -- wrongly, I think -- in a very greedy-algorithm approach to code reuse issues. My first choice was to take the BSD 4.4-Lite release and make a kernel. I knew the code, I knew how to do it. It is now perfectly obvious to me that this would have succeeded splendidly and the world would be a very different place today.

RMS wanted to work together with people from Berkeley on such an effort. Some of them were interested, but some seem to have been deliberately dragging their feet: and the reason now seems to be that they had the goal of spinning off BSDI. A GNU based on 4.4-Lite would undercut BSDI.

So RMS said to himself, "Mach is a working kernel, 4.4-Lite is only partial, we will go with Mach." It was a decision which I strongly opposed. But ultimately it was not my decision to make, and I made the best go I could at working with Mach and doing something new from that standpoint.

This was all way before Linux; we're talking 1991 or so.

Currently, the Hurd runs on IA32 machines. The Hurd should, and probably will, be ported to other hardware architectures or other microkernels in the future.

According to Thomas:

'Hurd' stands for 'Hird of Unix-Replacing Daemons'. And, then, 'Hird' stands for 'Hurd of Interfaces Representing Depth'. We have here, to my knowledge, the first software to be named by a pair of mutually recursive acronyms.

The FSF states: "The Hurd, together with the GNU Mach microkernel, the GNU C Library and the other GNU and non-GNU programs in the GNU system, provide a rather complete and usable operating system today. It is not ready for production use, as there are still many bugs and missing features. However, it should be a good base for further development and non-critical application usage."

Unfortunately, the Hurd is late. By 1995, Linux had many users. By 2000, it was a well-understood and popular system. By 2005, Linux had millions of users and the support of IBM. It was seen as a threat by Microsoft. The Hurd, unfortunately, is still "not ready for production use."

BSDI

BSDI was the first company to offer a full version of BSD Unix for the Intel platform.

Despite the fact that everything was in the public eye and exposed at the USL vs. BSDI trial, there appears to be confusion as to the history of BSDI.

I think Thomas was right, to a certain extent.

While several Berkeley developers were involved in the formation of BSDI in 1990-91, none left the University of California to join Berkeley Software Design, Inc. at the outset. BSDI was founded by Rick Adams, who told me: "It was my idea and my funding. I also handled the logistics (via UUNET) and the little matter of the lawsuit."

Donn Seeley related:

The first organizational meeting occurred at a bar in Boulder during the Boulder Berkeley Workshop in October 1990. I was invited to the meeting without any advance warning and to my surprise, I was offered a job. My recollection is that Rick, Mike, Kirk, Keith, and Bill J[olitz] were present at the meeting. I believe that a more formal meeting was held in early December 1990 at Kirk's house [in Berkeley], where we voted to go ahead with the proposal. I think this meeting was when we came up with the name BSDI.

We decided to work under UUNET's wing for a while so that we would not alert any potential competition; that continued until the summer of 1991. I was to start work as soon as possible; I took an extended vacation from my job at the University of Utah, and set up shop in my parents' basement in Bellingham, WA, with a PC provided by Rick, running mt Xinu Mach/BSD (I think). (I don't remember exactly when I gave notice at Utah, but I set things up so that my employment terminated shortly before the Winter Usenix [21-25 January 1991; Dallas].) I couldn't actually work directly on the OS, since it still contained licensed code at that point.

The BSD distribution was still hung up on the issue of certain possibly licensed files, so my job

was to work on freely available software. I started with the compiler toolchain (based on GCC 1). Once it was clear that there would be missing files, I went ahead and wrote a replacement for the Berkeley init(8) program. I'm not sure whether Bill was working on BSDI-related stuff at this point, but I'm pretty sure that he had started by the time of the 1991 Winter Usenix, where we all met again.

At that time Kirk McKusick was President of USENIX, Rick was in Dallas to report on UUNET and recruit, Trent Hein was chairing the session on File Systems, and Keith Bostic and Mike Karels were part of the CSRG. It wasn't hard to call a meeting.

Trent was a student at the University of Colorado, where he was a co-author of both the UNIX and the Linux system administration handbooks. He worked on the 4.4BSD port to the MIPS architecture. More recently, he was co-founder of Applied Trust Engineering. He said:

I can concretely say that the original engineering team "hired" by BSDI (Spring, 1991) consisted of Donn Seeley, Bill Jolitz and myself. Bill left BSDI later that year. Rob Kolstad joined the BSDI team much later. [Kolstad was Secretary of USENIX at that time.]

Mike Karels told me:

I'd say that the founders were Rick Adams, Keith Bostic, Kirk McKusick, me, Bill Jolitz and Donn Seeley, in approximately that historical order. This group was involved at the time of formation. Bill and Donn were the first two full-time employees, and Trent started about the same time at just less than full-time. They worked for UUNET temporarily until the company started operations, which I think was about July 1991. Bill left at the end of November '91, and Rob [Kolstad] started December 1. The proximity of the dates is not a coincidence. I started February 1, 1992, at which time two Russians had also started, and also Jeff Polk. My departure from Berkeley and position at BSDI were announced at USENIX in January '92 [San Francisco], at which Bill made a famous appearance.

I asked Rick to clarify and he affirmed:

The first employees were Donn Seeley and Bill Jolitz. Peter Collinson signed on very early for European sales and Bob Kummerfeld for Australia.

We picked up Vadim Antonov and Igor Belchiski from USSR that fall (1991). Rob Kolstad came on as president in December 1991.

Donn Seeley provided yet more detail.

Bill believed that he deserved a larger role as systems architect, press contact and marketer. His coding contributions mainly came before he started working for UUNET/BSDI, by porting to PCs the drivers we'd written at Utah for HP 68k systems, and writing the locore assembly source and related files. As for Bill's departure, the straw that broke the camel's back was an issue with Bill's unauthorized expenses for a trip to Europe, if I recall correctly, but it was clear long before this point that Bill was not happy. Rick was BSDI's original president, but he was asked to separate UUNET from BSDI by UUNET's first big investors, so he enlisted Rob to replace him.

[There is a long and complex tale concerning Jolitz' departure and his appearance at the January 1992 USENIX meeting. I do not think it relevant to this narrative. One view may be found [here](#).]

Insofar as Keith Bostic was concerned, he said:

I joined much later than Mike and the founders, though. I stayed at UC Berkeley for quite some time after BSDI was founded.

Another person mentioned by Rick was Peter Collinson. In 1980-81, Collinson (then at the University of Kent in Canterbury) was offered a USENET feed by Teus Hagen at the CWI in Amsterdam. They couldn't dial out, but the CWI would dial in, via a modem brought into the UK by Martin Levy. In April 1982, he was instrumental in the formation of EUnet.

"I think it was the Fall of 1993 that Rick asked me to sell things in Europe," Collinson told me.

The earliest date on a file that I have is September 1993. I think I was at a BSDI meeting at the Usenix conference in San Francisco in January 1994 [January 21-24].

When did I leave? -- we were forced out by the sales department at the end of 1995 -- we had the fax in September -- we settled and were gone by January 1996.

We in Europe did OK -- but we were not that good at Sales -- and would have had to think hard about Sales-led sales rather than Techy-led sales very soon anyway.

In 2000, BSDI merged with Walnut Creek CDRom and then with Telenet Systems. The next year, Wind River Systems purchased the software business. Renaming itself iXsystems with plans to specialize in hardware, the server business was acquired by Offmyserver in 2002. I asked Collinson why he thought BSDI had failed.

BSDI didn't really fail. It allowed Linux to flourish unhindered by lawsuits; but it was not really technically viable. BSDI couldn't move quickly enough to keep up with the technical changes -- and Linux could because of the customer base which was a new generation of UNIX hackers and lovers.

Excursus: The GPL and Other Licenses

A license is formal permission for something. In general, this is written permission. Historically, such things began with (unwritten) permission to "go upon or use the land of another" -- to cross a lord's manor or forest, for example -- as a "personal privilege, not constituting an interest in the land."

Copyright is a form of an exclusive license. It is a license granted by the state to the "author," granting them a monopoly on certain activities for a time. The first copyright statute was the Statute of Anne, usually cited as 1709. (As with almost everything, the date is ambiguous: the Statute was introduced in 1709, but "passed into law" on 10 April 1710. The actual publication of the Act [BL 8 Anne c. 19, 261] is just headed "Anno Octavo." But, as the eighth year of Anne's reign terminated on 17 March 1710, contemporaries must have thought of the statute as dating from 1709. On the other hand, Adam Budd [TLS 15 July 2005, p. 24] calls it "The Copyright Act of 1710.")

At any rate, the Statute required booksellers to record their titles in the Stationer's Register to gain protection for their "literary property." All registrations stem from this.¹

In addition to copyright, the law recognizes patents and trade secrets (as well as trademarks, service marks, and "trade dress"), insofar as intellectual property rights are concerned. Until a few years ago, it was generally held that patents needed to be capable of physical instantiation. In 1997, patents on "business methods" were recognized. The next year, Amazon.com filed for a patent on "one-click" purchasing. In 2005, the battle over software patents is still being fought, while interpretations of what is allowed seem to be expanding.

There is no formal filing where trade secrets are concerned, although there are recognized steps you must take to protect your secrets, and most lawsuits have concerned improprieties: breach of contract and outright theft, for example. Attempting to keep the knowledge secret (think of the famed Coca-Cola recipe) is basic here.

To all effects and purposes, software licensing began when UNIX was ported to the Interdata 7 (in Australia) and the Interdata 8 (at BTL). Prior to that, an OS ran only on the machine with which it was sold or, in the case of UNIX, on the PDP-11. (And, if you were running UNIX and had the misfortune to call DEC service, they would tell you they "didn't know" about the alien system.)

The first "open source" license was that of the University of California at Berkeley (I will not distinguish between "free" and "open" source at this time). It came about through an active belief in academic freedom -- the right to examine and investigate.

As has been related earlier, the specific impetus given to Stallman to write the original version of the GPL was the refusal of LISP Machines to share their code.

Though there are now over 500 variants of Free and Open Source licenses, more than fifty of them recognized by the [Open Source Initiative \(OSI\)](#), they all trace their roots to the BSD and the GPL licenses. And, actually, I see both of these as having a common origin in the concept of "the right to tinker."

My guess is that there is no one reading this whose childhood is not littered with the parts of alarm clocks, toy trains, toasters, robots, radios, etc. (Yes, there is always an extra screw, gear or spring.) Part of learning how things work is taking them apart and reassembling them (or attempting to). Part of learning is destructive analysis.

And that tinkering leads to improvements.

Both the BSD and the GPL are founded in the notion that opening source code to examination leads to extensions and improvements that can subsequently be reincorporated into future code.

But proprietary software (and hardware) doesn't permit that.

In fact, as [Professor Edward Felten](#) has repeatedly pointed out, the [Digital Millennium Copyright Act](#) (DMCA), specifically criminalizes tinkering. It has been alleged it also criminalizes, for example, access to no-longer-current Web pages. The Wayback Machine has just been [sued](#) by Healthcare Advocates for retaining archived pages. While the courts may dismiss the suit, its very presence is chilling.

In Prokofiev's "Peter and the Wolf," the wolf, in his haste, swallowed the duck whole. In their headlong run to "protect" everything, the US Congress more or less did the same with the DMCA Title V, sections 1301-1332. The result has been extreme interpretations, leading to vexatious lawsuits, such as the attempt, happily failed, to extend the DMCA to cover garage door openers.

While the excesses of the large media-producing companies are many, this "slop-over" into other areas was, I'm sure, unintentional. However, the DMCA in the US and its proposed parallel in Canada, and its support by WIPO, are having truly stultifying effects on research and development.

For details on the variety of Open Source licenses, see Rosen's superb book on the subject.² However, I'd like to mention some things I believe make a few licenses beyond [the GPL](#) and the [New BSD](#) licenses important.

- The [MIT license](#) includes a "right to sublicense".
- The [Apache license](#) protects the Apache trademark, acknowledging the importance of trademarks to open source projects.
- The [Mozilla Public License](#) distinguishes "files containing derivative works," as opposed to just "derivative works."
- The Open Group now has a "[Test Suite License](#)."
- The W3C has "without fee or royalty" permissions language in its [W3C® Software Notice and License](#).

"Free as in freedom."

¹ Anyone interested in the history should read Ronan Deazley's *On the Origin of the Right to Copy* (Oxford, 2004).

² Lawrence Rosen, *Open Source Licensing*, Prentice Hall PTR, 2005.

Chapter 17. The Web

Just what will inspire invention is infinitely variable.

Ted Nelson says that his notion of hypertext was inspired by Vannevar Bush's "As We May Think"¹ and by S.T. Coleridge's poem "Kubla Khan" (1798, published in 1816).

Sir Timothy John Berners-Lee says that in his childhood home there was a book entitled Enquire Within upon Everything, a "musty old book of Victorian advice." What we now think of as the Web, was originally called "Enquire."²

The son of two mathematicians, Tim Berners-Lee took a degree in physics from Queens College Oxford and then worked for Plessey Telecommunications and D.G. Nash, prior to going to CERN as an independent contractor in 1980.

At CERN, Berners-Lee felt a need for researchers to locate and share information. Having read Ted Nelson's work, he determined that hypertext was the appropriate model to use. With the aid of Robert Cailliau, he set out to build a prototype system -- Enquire. But Berners-Lee left CERN at the end of 1980 to work for Image Computing Systems.

In 1984, Berners-Lee returned to CERN as a fellow and immediately went to work on CERN's Internet site, which by 1989 was the largest single site in Europe. He jumped at the opportunity of "marrying" the notion of hypertext and the Internet.

In Chapter 5, I outlined Lesk's development of uucp (1976) and the evolution of Netnews and the search engines (Gopher, archie, WAIS). What Berners-Lee was creating was the logical product of this decade's work by a variety of people. "I just had to take the hypertext idea and connect it to the TCP and DNS ideas and -- ta-da! -- the World Wide Web."³

Berners-Lee envisaged knowledge as an immense reticulum, and so he named his creation the World Wide Web. To navigate within the Web, he designed and built the first browser (WorldWideWeb) and developed (on NextStep). The first server was called httpd (hyper text transfer protocol daemon). The new proposal for this was written on November 12, 1990; work was begun the next day. The tools were written over Christmas holiday 1990-91. The world learned about it on August 6, 1991, when Berners-Lee posted a summary of his project on alt.hypertext.

The Web is an information space in which items of interest ("resources") are tagged with global identifiers (Uniform Resource Identifiers [URIs]). The Web is not the Internet, it is a service operating on the Internet.

And on April 3, 1993, CERN announced that the code would be free, with no fee. This last was crucial, for the University of Minnesota had succeeded in dashing the enthusiasm for Gopher through the cold water of a fee.

The Internet was free. TCP/IP was free. UUCP was free. Gopher had no chance. The World Wide Web now did.

I'm certain that Vannevar Bush had no notion of the inspiration his 1945 article would provide: to Doug Englebart and Ted Nelson; to Tim Berners-Lee; to innumerable others. But what has been salient over these 60 years has been the notion of building on the previous constructs, which have been freely accessible.

Hypertext (in the sense most of us use it) has little to do with what Ted Nelson wrote about in the late 1960s and the 1970s. I asked Ted about the Web:

"Berners-Lee came to my office in 1992 and showed me what he'd done," he told me. "I was polite, didn't say I thought it was stupid, and took him to lunch. That was the extent of our interaction."

He continued:

"The web has nothing whatsoever to do with my notion of hypertext, and I am still fighting for what I believe in. Real Soon Now, I hope this month, I'll be announcing a new spec called Transliteration. Watch for it.

"What would I have to do with http?"

But hypertext was Ted's concept. It has been refashioned into something very different.

And I can't even buy a bar of soap that doesn't have a URL on it.

¹ First published in The Atlantic, January 1945

² Enquire Within... was one of the very many Victorian compendia. It was originally published in 1859 and went through over 100 printings and editions, the most recent of which was published in New York in 1978.

³ For a truly personal view of this history, see Berners-Lee, Weaving the Web (1999).

Chapter 18. "Just for Fun"

I frequently point to August 1969 as the "birthmonth" of UNIX. A few days later, the ARPANET (soon to become the Internet) was born. And, on 28 December 1969, Linus Torvalds was born.

- Murray Hill, NJ
- Los Angeles, CA
- Helsinki, FI

I believe each step grew from the earlier one(s).

Where Free Software is concerned, the geographical spread is equally interesting:

- Richard Stallman, New York and Cambridge, MA
- Tim Berners-Lee, Oxford, UK, and Geneva, CH
- Linus Torvalds, Helsinki, FI

Linus was born into the Swedish minority of Finland (about 5% of the 5,000,000 Finns). Linus was a "math guy" throughout his schooling. Early on, he "inherited" a Commodore VIC-20 (released in June 1980) from his grandfather; in 1987 he spent his savings on a Sinclair QL (released in January 1984, the "Quantum Leap," with a Motorola 68008 running at 7.5MHz and 128kB of RAM, was intended for the small business and the serious hobbyist). It ran Q-DOS. And it was what got Linus involved:

One of the things I hated about the QL was that it had a read-only operating system. You couldn't change things...

[Linus] bought a new assembler ... and an editor ... Both ... worked fine, but they were on the microdrives and couldn't be put on the EEPROM. So I wrote my own editor and assembler and used them for all my programming. Both were written in assembly language, which is incredibly stupid by today's standards...
[Just for Fun (2001), p. 45]

That was the beginning. A high school student, interested in bettering his system, wrote the tools he wanted.

During his first year at the University, Linus tells us that he did little programming, and at the end of that year, he enlisted in the Finnish army to fulfill his obligation. He was 19. He "got out" on 7 May 1990.

In the fall of 1990, the University of Helsinki installed its first Unix machine, a MicroVAX running Ultrix. But Linus was "eager to work with Unix by experimenting with what I was learning in Andrew Tanenbaum's book" (p. 53). Linus read all 700-odd pages of Operating Systems. The book "lived on my bed."

One of the things that struck Linus about Unix was its openness. Another was its simplicity. And then came a bolt from the blue: in early 1991, Lars Wirzenius dragged Linus to the Polytechnic University of Helsinki to hear Richard Stallman. "I don't remember much about the talk," Linus says. "But I guess

something from his speech must have sunk in. After all, I later ended up using the GPL for Linux."

But on 5 January 1991, Linus got his father to drive to a "mom and pop" computer store, where he had ordered a no-name 4-meg, 33MHz, 386 box, so he could get it home. He was 21. The box came with DOS, but Linus wanted Minix, and ordered it. It took a month to find its way to Finland. But it arrived. And Linus fed the 16 diskettes to the machine. And then he began "playing" with it. The first thing he wrote was a terminal emulator: "That's how Linux got started. With my test programs turning into a terminal emulator."

Because Linus was truly dependent upon the Internet and (specifically) the comp.os.minix newsgroup, we can date events far more accurately than in earlier decades.

We know that Linus' first posting to comp.os.minix, asking about the POSIX standard, was 3 July 1991. And we can see his posting about "doing a (free) operating system (just a hobby, won't be big and professional like gnu) ... This has been brewing since April ...," of 25 August 1991.

There was a reasonable expression of interest. We thus know that Linus put what we would now call Linux 0.01 up on the University of Helsinki ftp site on 17 September 1991. "No more than one or two people ever checked it out," he said.

The following January there was discernible growth in the Linux community, leading (I think) to the attack on Linux by Andy Tanenbaum on 29 January 1992 [see next chapter]. Perhaps more important, in the spring Orest Zborowsky ported X-windows to Linux.

The number of Linux users continued to grow, as did the versions of the software. .01 was 63KB compressed. Only a few weeks later, Linus posted .02 on 5 October. On 19 December, v.11 was posted; and on 5 January 1992, v.12 -- 108KB compressed -- appeared. On 7 March, there was v.95 and on 25 May 1992, v.96, with support for X, and taking up 174KB compressed.

It was barely a year since Linus' first posting, but in 1992 SuSE was formed, in February Bruce Perens released MCC Linux, and on 8 December Yggdrasil alpha was released.

1993 began with Yggdrasil beta's release (18 February) and went on to RedHat's being set up by Mark Ewing. August 1993 brought us Debian (from Debbie and Ian Murdoch).

And, on 5 November 1993, Linus spoke at the NLUUG (Netherlands UNIX Users' Group).

On 12 March 1994, Linus released Linux 1.0, basically v0.99, patch level 157. It was the first stable kernel distribution.

I don't want to go into extensive detail here. But I think that there are a number of important points to be made:

- The birth, growth and development was totally unorganized.
- It was well-distributed, geographically.
- It was conducted via the Internet.

Ted Ts'o was one of the first Linux users in the US. I spoke to him over dinner in Atlanta.

I was working as an undergraduate staff person at MIT -- I was planning to go to graduate school, but I got caught up in projects. So I've got some courses, but no grad degree.

I worked at Athena for three years. For Dan Geer and Jeff Schiller, who were not yet at Kerberos. In '91 I was working on a help desk application, and in the midst of this I discovered Linux.

It was via Usenet. It think .08 or .09 had been cross-posted.

I'll return to Ted's recollections later. The important thing was that thanks to the Internet and to Usenet, the work of a hobbyist in Finland could be picked up elsewhere in Europe, in Australia, and in the US.

"There was fairly strong social cohesion," Ted told me. "Linux was the first big project to succeed in a distributed fashion."

试用水印

Chapter 19. Tanenbaum and Torvalds

Linus posted his queries, his information and his work on comp.os.minix beginning in mid-1991. But on 29 January 1992, Andy Tanenbaum posted a note with the line:

Subject: LINUX is obsolete¹

After a few introductory paragraphs, Tanenbaum got to his real criticism of Linux:

As a result of my occupation, I think I know a bit about where operating systems are going in the next decade or so. Two aspects stand out:

Microkernel vs Monolithic System

Most older operating systems are monolithic, that is, the whole operating system is a single a.out file that runs in 'kernel mode.' This binary contains the process management, memory management, file system and the rest. Examples of such systems are UNIX, MS-DOS, VMS, MVS, OS/360, MULTICS, and many more.

The alternative is a microkernel-based system, in which most of the OS runs as separate processes, mostly outside the kernel. They communicate by message passing. The kernel's job is to handle the message passing, interrupt handling, low-level process management, and possibly the I/O. Examples of this design are the RC4000, Amoeba, Chorus, Mach, and the not-yet-released Windows/NT.

While I could go into a long story here about the relative merits of the two designs, suffice it to say that among the people who actually design operating systems, the debate is essentially over. Microkernels have won. The only real argument for monolithic systems was performance, and there is now enough evidence showing that microkernel systems can be just as fast as monolithic systems (e.g., Rick Rashid has published papers comparing Mach 3.0 to monolithic systems) that it is now all over but the shoutin'.

MINIX is a microkernel-based system. The file system and memory management are separate processes, running outside the kernel. The I/O drivers are also separate processes (in the kernel, but only because the brain-dead nature of the Intel CPUs makes that difficult to do otherwise). LINUX is a monolithic style system. This is a giant step back into the 1970s. That is like taking an existing, working C program and rewriting it in BASIC. To me, writing a monolithic system in 1991 is a truly poor idea. . . .

Linus responded the same day with: "Well, with a subject like this, I'm afraid I'll have to reply. Apologies to minix-users who have heard enough about linux anyway. I'd like to be able to just 'ignore the bait', but ... Time for some serious flamewesting!" and a long (somewhat intemperate, but this is a 23-year old student) response.

There was a good deal of going back and forth, and even Brian Kernighan put in a few lines. But the result was that Andy remains to this day a committed microkernel devotee and Linus has continued with a largely monolithic system. (Of course, this generalization is inaccurate, but it serves.)

And, on a certain level, there is no question in my mind but that Andy's position is right: microkernels are "better" than monolithic systems. But, on the other hand, I find both Andy's original posting unnecessarily rebarbative and Linus' "serious flametesting" inappropriate.

Over a decade later, I find it hard to discern any anger or resentment on either side. I asked Andy about the exchange, but he just shrugged me off. "In a straight test," he later remarked, "Linux loses by about 8% in speed." That may well be true. But it's not much of an epitaph.

However, I think the microkernel (as evidenced in Mach [and in the Hurd], in Chorus, in Amoeba) is superior to the monolithic kernel, as long as the message-passing is efficient.

I guess I'll now be subject to a flame war.

¹ A large collection of the correspondence -- or at least that of the "major" contributors -- can be found here. As I am interested in discussing this, I will refrain from extensive citation. A version of much of the discussion is available as Appendix A of *Open Sources: Voices from the Open Source Revolution* (O'Reilly, 1999; ISBN 1565925823).

Chapter 20. Proliferating Penguins

From the early 1980s on, the big gripe about Unix was that it had split and resplit, that there were just too many variants. The fact that they had a common base was irrelevant to the critics -- and many (if not most) of those critics were selling VMS or MVS or DOS or...

Following Linus' postings of 1991, there soon were what we have come to call "distributions." And, rather than utilizing ftp, they came on CD-ROM.

The first of these was Adam Richter's Yggdrasil (in the Old Norse Edda, Yggdrasil is the "world ash," from a branch of which Odin/Wotan made his spear). Yggdrasil alpha was released on 8 December 1992. It was called LGX: Linux/GNU/X -- the three components of the system. Recall that Gilmore, Tiemann and Henkel-Wallace formed Cygnus in 1989. Richter spoke to Michael Tiemann about setting up a business, but was "definitely uninterested in joining forces with Cygnus."

Yggdrasil beta was released the next year. Richter's press release read:

The Yggdrasil beta release is the first UNIX(R) clone to include multimedia facilities as part of its base configuration. The beta release also includes X-windows, networking ... an easy installation mechanism, and the ability to run directly from the CD-ROM.

The beta was priced at \$50; the production release was \$99.

SuSE was formed in 1992 also, as a consulting group (SuSE was originally S.u.S.E., which stood for "Software-und-System-Entwicklung," Software and System Development), but did not release a Linux distribution for several years. The next distribution -- and the oldest still in existence -- was Patrick Volkerding's Slackware, released 16 July 1993, soon after he graduated from Minnesota State University Moorhead. It, in turn, was the basis for SuSE's release "Linux 1.0" of SLS/Slackware in 1994. (SLS was "Softlanding Linux System," Peter McDonald's 1992 distribution, on which parts of Slackware were based.) SuSE later integrated Florian La Roche's Jurix distribution, resulting in a unique distribution: SuSE 4.2 (1996).

The next year, Mark Bolzern was trying to sell a Unix database from Multisoft, a German company. He encountered difficulties because it was relatively expensive to set up the Unix system. Then he came across Gnu/Linux and realized that he now had a real solution. He convinced Multisoft to port Flagship (the db) to Linux and "that was the first commercial product released on Linux," Bolzern said.

"People were always having trouble installing Linux," he continued, "and then Flagship wouldn't run right because something had changed." Bolzern decided that what was needed was a release that wouldn't change for a year, so he "picked a specific distribution of Slackware" and "the name Linux Pro." Soon he was selling more Linux than Flagship: "we're talking hundreds per month."

And when Red Hat came out, Bolzern picked that up.

Mark Ewing had set up Red Hat in 1993. Mark Ewing has said: "I started Red Hat to produce a development tool I thought the world needed. Linux was just becoming available and I used [it] as my development platform. However, I soon found that I was spending more time managing my Linux box

than I was developing my software, and I concluded that what the world really needed was a good Linux distribution..."¹

In 1993, Bob Young was working for Vernon Computer Rentals. He told me: "I knew the writing was on the wall for my future with that company." He continued:

Red Hat the company was legally incorporated in March of 1993 in Connecticut under the name: ACC Corp. Inc. It changed its name to Red Hat Software, Inc. in early 1996, and changed its name a last time to simply Red Hat, Inc. just before going public in June of 1999.

ACC Corp. Inc. bought the assets, including all copyrights and trademarks (none were registered at the time) relating to Marc Ewing's sole proprietorship business venture in January 1993. Marc's Red Hat project was not incorporated but was run out of Marc's personal checking account. Marc received shares in ACC Corp, Inc. in return for the Red Hat name and assets.

In 1995 Red Hat packaged Linux, some utilities and initial support for \$50. Also in 1995, Bryan Sparks (with funding from Ray Noorda, former CEO of Novell) founded Caldera and The Apache Foundation released Apache, which would become the most widespread Web server. But Red Hat soon became the most popular Linux release. This was unexpected: Linus had said that he expected Caldera to be the top supplier, because it was "kind of a step beyond," in that it was targeting the office market. "I think what's interesting about Caldera is they based their stuff on Red Hat and then they added a commercial kind of approach."

When Red Hat became a "success," Bob Young and family moved from Connecticut to North Carolina (Ewing lived in Durham).

It was the end of July 1996. Just in time for Hurricane Fran, the first hurricane to visit Raleigh since hurricane Hazel in 1954. Yes, "the" hurricane Hazel that is the only hurricane to make it to southern Ontario still categorized as a hurricane that I know of.

(Before abandoning this, I should point out that Young is from Hamilton, ON, and attended the University of Toronto. During the night of October 18, 1954, "Hurricane Hazel pelted Toronto with rain and killed 81 people. On one street alone, Raymore Drive, 35 neighbors were drowned." Environment Canada, Canadian Hurricane Centre, Storms of 1954.)

ACC, Young's company, sold Linux/Unix software and books. Young had been introduced to the Unix world in 1988, when he was with Vernon Leasing and Rentals, and began publishing New York UNIX as well as catalog sales. This led to his being the founding editor of *Linux Journal*, a post he held for two issues in 1994, before "splitting the losses" with Phil Hughes, who is still the publisher of LJ.

In the summer of 1995, I was approached by Lisa Bloch, then the Executive Director of the FSF, as to the feasibility of a conference on "Freely Redistributable Software." I was enthusiastic, but had my qualms about profitability. Richard, at our meeting, was quite understanding: FSF would bankroll the affair, but he hoped we could turn a small profit.

Lisa and I put together a committee (Bob Chassell, Chris Demetriou, John Gilmore, Kirk McKusick, Rich Morin, Eric Raymond, and Vernor Vinge) and we posted a Call for Papers on several newsgroups.

Thanks to "maddog" (Jon Hall), Linus agreed to be a keynote speaker, Stallman was the other. We had a day of tutorials and two days of papers. February 3-5, 1996 at the Cambridge Center Marriott. Everything ran smoothly. By the end, I was a nervous wreck. And the FSF ended up making a tiny profit.

¹ See Glyn Moody, *Rebel Code* (Perseus Publishing, 2001), p. 97.

试用水印

Chapter 21: Daemonology

The daemon image, what Kirk McKusick calls the "beastie," dates from 1976. Created by comic artist Phil Foglio and first used by Mike O'Brien, here's the story as told by Mike:

I was a bonded locksmith. Phil's roommate had unexpectedly split town, and he was the only one who knew the combination to the wall safe in their apartment. This apartment was the only one I'd ever seen that had a wall safe, but it sure did have one, and Phil had stuff locked in there. I didn't hold out much hope, since safes were far beyond my competence, but I figured "no guts, no glory" and told Phil that I'd give it a whack. In return, I requested T-shirt art. Phil readily agreed.

Wonder of wonders, this safe was vulnerable to the same algorithm to which Master locks used to be susceptible. I opened it after about 15 minutes of manipulation. It was my greatest moment as a locksmith and Phil was overjoyed. I went down to my lab and shot some Polaroid snaps of the PDP-11 system on which I was running UNIX, and gave them to Phil with some descriptions of the visual puns I wanted: pipes, demons with forks running along the pipes, a "bit bucket" named /dev/null, all that.

What Phil came up with is the artwork that graced the first decade's worth of "UNIX T-shirts", which were made by a Ma-and-Pa operation in a Chicago suburb. They turned out transfer art using a 3M color copier in their basement. Hence, the PDP-11 is reversed (the tape drives are backward), but since Phil left off the front panel, this error was hard to detect. His trademark signature was photo-reversed, but was recopied by the T-shirt people and "re-forwardized"-- which is why it looks a little funny compared to his real signature.

The art was used on the USENIX Tenth Anniversary shirt in 1985 (I still have one). McKusick comments: "About 1 year after Usenix produced the Portland conference T-shirts, they paid Phil for the artwork. Thus, Usenix currently holds title to the copyright." I was the individual who actually paid Phil.

As I said [earlier](#),

NET 1 was a subset of the then-current Berkeley system. It was quite similar to 4.3-Tahoe, including source code and documentation for the networking portions of the kernel, the C library and utility programs. It was available without evidence of any prior license (AT&T or Berkeley), and was (re)distributed via anonymous FTP. The source carried a Berkeley copyright notice and a legend that allowed redistribution with attribution.

In June 1991, at the USENIX Conference in Nashville, BSD Networking Release 2 was available. NET 2 contained far more than just networking code and, like NET 1, was available with no prior license. The new features included a new virtual memory system (derived from Carnegie-Mellon's Mach system, which had been ported at the University of Utah) and a port to the Intel 386/486.

But all was not happy in Eden. While BSDI's version of the release was complete by the end of 1991, it was only released to the public on April 10, 1993 as 1.0, the long delay being the consequence of USL's

filing suit to prevent BSDI from shipping its product.

BSDI had distributed pre-production releases of BSD/386 (Beta version). It now began international distribution. Full source was priced at \$1000. (In the January/February 1994 ;login:, Lou Katz wrote: "It works! It works!").

On Friday, February 4, 1994, Novell and the University of California agreed to drop all relevant suits and countersuits. BSDI immediately announced the availability of a release based on "4.4BSD-Lite."

In the meantime, several groups of coders had begun work on other releases (daemons multiplying nearly as rapidly as penguins). The earliest of these was 386BSD, by Lynne and Bill Jolitz, though others have proven to be of greater importance.

William Jolitz had had considerable experience with prior BSD releases while at Berkeley (2.8BSD, 2.9BSD) and he and Lynne Jolitz contributed code to Berkeley developed at Symmetric Computer Systems during the 1980s. The public 386BSD releases beginning in 1992 were based on portions of the NET 2 release plus with additional code written by the Jolitizes.

FreeBSD derived from the 386BSD 0.1 release. It was the first free software organization founded on BSD.

Initial development of FreeBSD was started in 1993, taking its source from 386BSD. As a consequence of concerns about the legality of some of the code, NetBSD, like FreeBSD, was derived from 4.3BSD via NET 2 and 386BSD. The project began as a response to the FreeBSD unified patchkit and 386BSD, with its slow development process and focus on the i386 platform. The four founders of the project, Chris Demetriou, Theo de Raadt, Adam Glass and Charles Hannum, felt that a more BSD-like development focus would be more beneficial to the project.

de Raadt suggested the name "NetBSD" and it was readily accepted by the other founders. Although Demetriou inadvertently made a premature release of the code, NetBSD 0.8 was the first official release in May 1993.

In 1994, one of the founders, Theo de Raadt, was forced out of the project. He later founded a new project, OpenBSD, from a forked version of NET 1. The first release of OpenBSD was in October 1995.

The current release of NetBSD is version 2.1 (October 2005).

I had dinner with Poul-Henning Kamp in Copenhagen and asked him about his work on FreeBSD.

"I've been involved with FreeBSD from before it existed," he told me.

You know, the basic BSD kernel suffers from a number of Ph.D. theses. I see the split as between the academic (OpenBSD) and the practical (FreeBSD). Chris Demetriou is towards the more academic side.

I love having choice. Diversity in OSes is as important as is diversity in crops.

Chapter 22: Yet More Penguins

Debian Linux, as I stated in Chapter 20, was created by Ian Murdock. He officially founded the "Project" on August 16, 1993. From November 1994 to November 1995, the Debian Project was sponsored by the FSF.

In November 1995, Infomagic released an experimental version of Debian which was only partially in ELF format as "Debian 1.0." On December 11, Debian and Infomagic jointly announced that this release "was screwed." Bruce Perens, who had succeeded Murdock as "leader," said that the data placed on the 5-CD set would most likely not even boot possibly.

The real result was that the "real" release, *Buzz*, was 1.1 (June 17, 1996), with 474 packages. Bruce was employed by Pixar and so all Debian releases are named after characters in *Toy Story* (1995).

- 1.2 Rex, December 12, 1996 (848 packages)
- 1.3 Bo, June 5, 1997 (974 packages)
- 2.0 Hamm, July 24, 1998 ("over 1500 packages")
- 2.1 Slink, March 9, 1999 ("about 2250 packages")
- 2.2 Potato, August 15, 2000 ("more than 3900 binary packages")
- 3.0 Woody, July 19, 2002 (8500 binary packages)
- 3.1 Sarge, June 6, 2005 (15,400 packages)

Buzz fit on one CD. *Slink* went to two. *Sarge* is on 14 CDs in the official set. It was released fully translated to over 30 languages and contains a new *debian-installer*. *Slink* had also introduced ports to the Alpha and Sparc. In 1999, Debian also began a Hurd port.

Though Debian carried the burden of being tough to install for several years, *Sarge* has changed that. The new installer with automatic hardware detection is quite remarkable.

I introduced Red Hat in Chapter 19, and I will return to the company again, but at this point I'd like to introduce Mandrake, a Linux distribution based on Red Hat 5.1 and KDE. It was created by Gael Duval, a graduate of Caen University, in July 1998. From 1998 to early 2004, Mandrake was reasonably successful, notable for its high degree of internationalization as well as the variety of chips it would run on. However, in February 2004 MandrakeSoft lost a suit filed by the Hearst Syndicate which claimed invasion of their trademarked "Mandrake the Magician." Starting with 10.0, there was a minor name change. Then, in April 2005, Mandrakesoft announced that there was a merger with Conectiva, and that the new name would be Mandriva.

Joseph Cheek founded Redmond Linux in 2000. In 2001 it merged with DeepLinux. In January 2002 the company was renamed Lycoris and its Desktop/LX was based on Caldera's Workstation 3.1. In June 2005, Lycoris was acquired by Mandriva.

I've gone through all this to show just how complex the tale of Linux distributions can be. And, as of this writing, there appear to be well over 100 distributions. I will neither enumerate nor elaborate on most of them. However, the most "popular" appear to be:

- Red Hat
- Fedora

- Debian
- Gentoo
- Knoppix
- SuSE/SUSE (Novell)
- Slackware
- TSL
- Yellow Dog
- Mandriva
- College Linux
- Ubuntu
- Kubuntu
- Puppy

It might be a full-time job to track all the distributions and their origins. For instance, Kanotix is a Debian derivative. It is also a Knoppix derivative, as it is a live CD. And it is solid as a rock.

Knoppix was created by Klaus Knopper, a freelance IT/Linux consultant. It has achieved popularity because it is easily run from the CD, without installation and because it can be readily employed to fix corrupted file systems, etc. It was the first Linux on a live CD.

In 1996, Bob Young and Red Hat moved corporate headquarters to North Carolina. In January 1997, Greylock and August Capital invested \$6.25 million in Cygnus Solutions, becoming the first VCs to invest in a free software business. In July, Red Hat 4.2 was released and in December, 5.0 was announced.

These are important events, as could be seen in November 1998 when a Microsoft lawyer waved a Red Hat box in the air to "refute" the US Justice Department charge that Microsoft has a monopoly on the desktop operating system market.

While Red Hat may not have been the most innovative company, they had already become the iconic Linux enterprise.

In August of 1999, Red Hat had its IPO, the eighth largest first day gain in Wall Street history. (On 9 December 1999, VA Linux had its IPO.) And in November 1999, Red Hat acquired Cygnus, creating the largest "open source" company in the world.

Just how successful Linux and some Linux companies had become was made obvious at the outset of the new millennium:

- * In January 2001 Scott McNealy said that Linux is a "better NT than NT"; and
- * In February 2001 Steve Ballmer called Linux "a cancer" and "an intellectual property destroyer."

Oh, boy!

Chapter 23: Oceans of FUD

When Gene Amdahl coined the word "FUD" (for fear, uncertainty and doubt) in the mid-1970s, his ire was aimed at Frank Cary, chairman of the Board at IBM, who was waging a no-holds-barred attack on Amdahl, Intel, Control Data, and the other small companies that were selling machines that competed with the IBM 360/168. According to Robert Sobel:¹

The campaign began in a conventional fashion. IBM salesmen and executives visited clients who were thought to be considering plug-compatible machines, to warn them of problems that might arise should Amdahl or National Semiconductor leave the business. There was talk of reduced maintenance on IBM peripheral equipment hooked onto other mainframes, of software changes to eliminate or reduce compatibility, and of alterations in hardware that could make the Amdahls less compatible than advertised.

Sound familiar?

By the end of 1997, Eric Raymond had delivered "The Cathedral and the Bazaar" at least twice: at Linux Kongress in May and at the Perl Conference in November. It appeared on First Monday online in 1998, on paper in Matrix News in three installments (June, July and August 1998), and in book form in 1999. It does not seem to have been read in Redmond, WA.²

In the May 1999 issue of Microsoft Internet Developer, Douglas Boling wrote:

While free distribution is a great marketing tool (think about all those samples you get in the mail), what does it say about the product itself? Frankly, it says that the product (or the effort that went into making the product) has no value. Is that what you software engineers out there want?

... If ... you gave away all software, how would you pay the creators of that software?

Boling goes on, but I'll spare my readers. I was also going to cite Microsoft's "Linux myths," but that page is no longer accessible at <http://www.microsoft.com/ntserver/nts/news/msnw/LinuxMyths/asp>

It was there I read that there were "hundreds of UNIX vendors with no 'standard' flavor of UNIX" [take that, POSIX!]. And that "Windows NT 4.0 Outperforms Linux on Common Customer Workloads" [inability of the Linux stack to handle multiple network cards on SMP machines adequately was the vital "Customer Workload," incidentally].

But in August 1999 Red Hat had its IPO and, by Christmas, it had acquired Cygnus and VA Linux Systems had had its IPO. Free software was becoming big business.

But then, so was nearly everything else. Pets.com, boo.com, and a variety of other fantasies blew hundreds of millions of dollars.

In 1999 we were nearly at the peak of what was (retrospectively) known as the Dot-Com bubble: the dot-bomb. But a look at history is needed.

NASDAQ was begun in 1951. By 1990 it was a good-sized marketplace with a large number of new

and recently-formed corporations holding their IPOs (Initial Public Offerings). Many of these were hi-tech; many were in areas previously untested -- selling products over the Internet, setting up and using Web sites, indulging in e-commerce rather than selling products in shops. Some of the new companies were actually involved in Information Technology, rather than using it. But they all seemed to show great promise, and folks didn't want to miss the boat.

After the stock market crash of 1987 (in which the Dow-Jones average dropped 22.6% and lost about \$500 billion on October 19), the markets around the world continued their bullish ways. In the early 1990s, the personal computer was becoming a household object and the advent of the Web made access yet more user-friendly. 1994 saw the business world "discover" the Internet as a commercial opportunity, and yet more companies were formed. Amazon began in 1994; eBay in 1995. On December 5, 1996, Alan Greenspan warned of "irrational exuberance" as evidenced by the rapidly rising stock prices.

In 1997, NASDAQ announced a new listing standard: it would base new listings on market capitalization alone, basically telling the world that accounting regulations hindered many new firms, preventing them from listing. There was a surge of registrations. In fact, "nearly 50% of the new listings between Aug 1997 and June 2000 entered under the market capitalization standard."³

Greenspan's warning didn't count. From 1996 to 2000, NASDAQ went from 600 to 5000. And then it crashed. Within six weeks, NASDAQ dropped from 5000 to 2000, then to 800 (in 2002). MicroStrategy, a soaring business-software provider, fell from \$3500 per share to \$4, the victim of an accounting scam. On December 14, 2000, it was at \$15.19. The emperor had no clothes. On March 10, 2000, the NASDAQ fell from its peak of 5132.52. Over five years later, at the end of 2005, it had climbed back to 2200. Even Microsoft dropped from over \$60/share to \$20/share in 2000-2001, losing two-thirds of its (paper) value. It closed 2005 at under 50% of its peak.

Looking at Klein and Mohanram again, "367 non-financial firms [were] listed under the Type 3 criteria between ... August 1997 and the end of the hi-tech IPO boom in June 2000. Without this alternative, none of these 367 firms would have entered the NNM [NASDAQ National Market] on their entry date." Moreover, "over a four-year event-time window, Type 3 firms earn significantly less than other NNM new listings..."

Klein and Mohanram illustrate that the inflation of the bubble (and its bursting) were not merely "irrational exuberance," but specifically an "irrational exuberance" concerning barely-understood yet extensively hyped hi-tech ventures. No idea was too bizarre to invest in.

These last few paragraphs are a background. The rise and ebb of FUD has consistently followed the rise and fall of the stock market or the rise and fall of (perceived) commercial threats. Thirty years ago, the rise of "other" mainframes worried IBM. The collapse of its stock price worried Microsoft. So did the fact that new offerings were not really in the offing.

One of the useful forms of past FUD had been "preannouncement" -- press concerning wonders of the future. Following what appeared in The Register, we can find:

- July 27, 2001: "an intermediate release ... dubbed 'Longhorn' will ... slip out late next year or early 2003."
- August 7, 2001: "the next release of Windows Server, codenamed Longhorn and due in mid 2003..."
- October 24, 2001: "the wheels have come off the Windows rollout wagon..."

- May 8, 2003: "It will assuredly be stuff that's in Longhorn ... but we detect bits that must currently be missing, and that will be hard, if not impossible, to execute by 2004."
- August 27, 2004: "Microsoft project managers have demanded that features be jettisoned in order for the next major version of Windows to ship as projected by 2006..."
- May 19, 2005: "Gartner says the first Longhorn client could slip into 2007..."

But then, on July 22, 2005, Microsoft issued a [press release](#):

Media Alert: Microsoft Unveils Official Name for "Longhorn" and Sets Date for First Beta Targeted at Developers and Professionals.

The date of release was August 3, 2005.

What's the function of this?

Let's suppose that you're the CIO or CTO reporting to the CEO of a Fortune 1000 company. Microsoft targets its marketing pitch at that CEO. Your company is going to invest lots of cash, dollars, yen, pounds, euros. Do you take a chance on the unknown (Mandriva, SuSE, Red Hat) or stay with the familiar (known to your CEO) and wait? Remember: No one ever got fired for buying {IBM, XEROX}.

Preannouncement is one tactic; planting "news" is another; questioning bonae fides is a third.

As an illustration, here are some data from 2003:

- March 2003. Caldera (d/b/a The SCO Group) files suit against IBM in 3rd Judicial District, Salt Lake County, court.
- March 25, 2003. The case is removed to Federal jurisdiction.
- May 29, 2003. Chris Sontag, SCO Group's "senior vice president and general manager of SCOSource Division," tells Patrick Thibodeau of Computerworld: "There is no mechanism in Linux to ensure [the legality of the] intellectual property being contributed by various people. ... I would suspend any new Linux activities until this is all sorted out."
- The 1Q2003 Caldera filing with the SEC reveals nearly \$10 million income from two license sales: to Microsoft and Sun.
- October 16, 2003: Press Release: "\$50 Million Private Investment Transaction Led by BayStar Capital Provides SCO With Funding for ... and the Protection of the Company's Intellectual Property Assets." (This was later altered to "from Two Investors including BayStar Capital..." The SEC 8K and purchase agreement reveals the second (larger) PIPE investor to be the Royal Bank of Canada.)

License fees, private equity investments. Shoring up confidence in a company; raising questions for potential customers; stalling for time when an OS is delayed; paying the lawyers (Boies, Schiller received a \$31 million fee).

Nearly a year earlier, in 2002, the Alexis de Tocqueville Institution issued a white paper by Kenneth Brown using, according to Richard Forno, "'terrorism' and 'national security' [in] shameful attempts to use fear, uncertainty and doubt to push Microsoft's monopolistic agenda."⁴

In May 2004, the Institution and Brown resurfaced. This time, Brown put out a "study" which claimed that Linus Torvalds wasn't the father of Linux at all. Here's a part of [the press release](#):

In one of the few extensive and critical studies on the source of open source code, Kenneth Brown, president of AdTI, traces the free software movement over three decades -- from its romantic but questionable beginnings, through its evolution to a commercial effort that draws on unpaid contributions from thousands of programmers.

Among other points, the study directly challenges Linus Torvalds' claim to be the inventor of Linux.

Brown's account is based on extensive interviews with more than two dozen leading technologists in the United States, Europe, and Australia, including Richard Stallman, Dennis Ritchie, and Andrew Tanenbaum.

"The report," according to Gregory Fossedal, a Tocqueville senior fellow, "raises important questions that all developers and users of open source code must face.

"One cannot group all open source programmers together. Many are rigorous and respectful of intellectual property. Others, though, speak of intellectual property rights -- at least when it comes to the property of others -- with open contempt."

Linus responded, saying it was true -- he had been found out, "The true fathers of Linux are Santa Claus and the Tooth Fairy."⁵

Andy Tanenbaum was less easygoing: "Brown is not the sharpest knife in the drawer," [he posted](#).

While listed as a 124-page E-Book, Mr. Brown's opus is "not yet available," over 18 months after the press release.

In general, FUD has quite limited utility. In the 1970s it could be somewhat effective. The growth of the Internet has reduced that: reality moves around at the speed of light. And while Don Basilio was right about rumors (in Rossini's "Barber of Seville"), technology has caught up with him.

¹ *IBM: Colossus in Transition* (Times Books, 1981), chapter 15

² There will be further discussion in a future chapter on "literature."

³ I am indebted to the extensive analysis of April Klein and Partha Mohanram, "They Came, they Conquered, they Collapsed" (March 2005).

http://www.lerner.udel.edu/finance/Seminar_Papers/listingrequirements_7.pdf

⁴ "[Alexis de Tocqueville Serves Up a Red Herring](#)," *Security Focus* June 19, 2002. The paper, "Opening the Open Source Debate," is available at <http://www.adti.net/ip/opensource.pdf>.

⁵ "[Linus Discloses "Real" Fathers of Linux](#)," *LinuxWorld* May 17, 2004.

Chapter 24: The Documents of Freedom

Richard Stallman wrote [The GNU Manifesto](#) in 1984. As I said in Chapter 12, "it marks the true beginning of the GNU Project." One part of the Manifesto, "Why I must write GNU," has been a "favorite" of mine for over twenty years. Let me quote it again.

I consider that the golden rule requires that if I like a program I must share it with other people who like it. Software sellers want to divide the users and conquer them, making each user agree not to share with others. I refuse to break solidarity with other users in this way. I cannot in good conscience sign a nondisclosure agreement or a software license agreement. For years I worked within the Artificial Intelligence Lab to resist such tendencies and other inhospitalities, but eventually they had gone too far. I could not remain in an institution where such things are done for me against my will.

So that I can continue to use computers without dishonor, I have decided to put together a sufficient body of free software so that I will be able to get along without any software that is not free. I have resigned from the AI lab to deny MIT any legal excuse to prevent me from giving GNU away.

The various GNU tools were a success: the FSF's versions of AWK, the C compiler, Emacs, yacc (-> Bison), etc., were used widely. With the advent of Linux, "free" software took off.

But 40 years earlier, Joseph Schumpeter (1883-1950) had published [Capitalism, Socialism and Democracy](#) (1942; 2nd Ed., 1944; 3rd Ed., 1950) in which he coined the phrase "[creative destruction](#)." Schumpeter felt that capitalism would not survive, not because it would be overthrown by a Marxist proletariat, but because it would be destroyed by its successes. He believed that capitalism would give rise to a large intellectual class that subsisted by attacking the bourgeois system of private property and freedom which were necessary for the class's existence.

Schumpeter lauded the entrepreneur, who figures out how to use inventions, who introduces new means of production, and new forms of organization. It is innovation by the entrepreneur, he claimed, which leads to "creative destruction" as old inventories, ideas, technologies, skills become obsolete.

The growth of an IT-driven society has resulted in a world-wide intellectual class (Brazil, Chile, China, and India are good examples outside of Europe and North America). Globalization and outsourcing are obsoleting the fixed industrial plant and local labor employment practices of 1800-1960.

And FLOSS has become the force that drives the "gales" of creative destruction today.

Among those who realized that FLOSS was a disruptive technology were Eric Raymond and "Doc" Searls. While there are many other authors and books, I will use them as my victims.

In chapter 23 I wrote:

By the end of 1997, Eric Raymond had delivered "The Cathedral and the Bazaar" at least twice:

at Linux Kongress in May and at the Perl Conference in November. It appeared on First Monday online in 1998, on paper in Matrix News in three installments (June, July and August 1998), and in [book form](#) in 1999.

In 1994, [Linux Journal](#) appeared. Produced by Bob Young and Phil Hughes, the first two issues were simple two-color print jobs, but it soon became a multicolor slick. After the second issue, Hughes was alone. Young went on to Red Hat.

But there was now a serious organ for Linux and for GNU tools. But let me return to Raymond.

As I have read it for nearly a decade, the point of Raymond's metaphor is that centralized design (the cathedral -- the unitary, proprietary, corporate method) is an ineffective, inflexible method when compared with the efficacy of the open source process (the bazaar of a multiplicity of vendors).

The bazaar is an open market where all are free to evaluate the merchandise (software) and decide to use or improve it. The cathedral refers to closed, proprietary structure (programming where the software is kept pure of outside influences, developed by a small team, usually with a hierarchical organization.)

Where proprietary software is concerned, a small team's production tends to be buggy and frequently does not correspond to the (potential) customer's expectations, wishes or needs. Open software is produced in (semi-)public, is exposed to examination and (frequently) destructive testing by thousands of pairs of eyes, encourages access, and thereby generally supplies precisely what is needed.

Freely redistributable software with accessible code means that a company or an individual with special needs or unique hardware can adjust, adapt or extend the software to suit those requirements. With proprietary code, a request can be made, but if you aren't a valued customer, your needs may remain unfilled.

"Doc" Searls spent years as a high tech marketer and has been Senior Editor of [Linux Journal](#) for a decade. Together with Chris Locke, Rick Levine, and David Weinberger, he composed and posted 95 theses on a web site (www.cluetrain.com) in 1999.¹

The essays in [The Cluetrain Manifesto](#) are important: but they are repetitious, as are the theses. Nonetheless, they are worth reading, and the first two theses are of real value:

1. Markets are conversations.
2. Markets consist of human beings, not demographic sectors.

Pause a moment to consider just how important these are. How many times have you answered the phone to find someone you don't know calling you by your first name? And it may not be what you are usually called. Sales/marketing folk appear to think that this opens up their campaign. I hang up on them.

In the market place, the *agora*, Raymond's bazaar, one engages in conversations -- walking to a specific stall or shop, greeting the vendor or clerk, discussing and selecting the merchandise, paying and receiving. Real conversations; real exchanges.

Among many things, *The Cluetrain Manifesto* suggests that the strategem that usually accompanies

buying and selling should be replaced by a true attempt at satisfying the needs, wants and desires of those on both sides of the equation. Despite their long digressions, the authors occasionally succeed in making solid, clever points that reveal fundamental flaws in the structure of traditional businesses. Consider this comment about business hierarchies: "First they assume--along with Ayn Rand and poorly socialized adolescents--that the fundamental unit of life is the individual. This despite the evidence of our senses that individuals only emerge from groups."

Their Sixth Thesis counsels "The Internet is enabling conversations among human beings that were simply not possible in the era of mass media," then business is warned by the Seventh Thesis: "Hyperlinks destroy hierarchy." Hierarchies rank people and restrict information flow because information access is a function of rank. Hyperlinks democratize information flow, nullifying the main offensive weapon that hierarchies depend on to remain hierarchies. (This is, of course, what governments are beginning to discover: the Internet is an anarchy machine. Information is a destabilizing force.)

Most leaders in Old Economy hierarchies see the Internet as just a new product distribution channel (effectively, both amazon.com and eBay are examples of this). They don't realize that the Internet is a new conversation channel that greatly amplifies the voices in the marketplace (blogs like Daily Kos and Groklaw are examples of this). Cluetrain's First Thesis states, "Markets are conversations."

Chapter 4 of *The Cluetrain Manifesto*, by Searls and Weinberger, is an excellent exposition of how today's businesses "have to figure out how to enter this global conversation rather than relying on the old marketing techniques of public relations, marketing communications, advertising, and other forms of propaganda. We don't want messages at all, we want to speak with your business in a human voice."

¹The "95 theses" were a specific allusion to the 95 Theses which Martin Luther posted (!) on October 31, 1517, which condemned "greed and worldliness," among other things. Pope Leo X dismissed him as "a drunken German."

Chapter 25: The URL on Your Cereal Box

In Chapter 17, I limned the creation and development of the Web. In a subsequent chapter, I'll talk about the geographical spread of Linux. But first, I want to look at the spread of the Internet and the Web that depends on it.

The ARPAnet became functional in 1969: at the end of that year, there were four nodes. In January 1976, there were 63 (so much for 5- or 6-bit addressing). Five years later, in August 1981, Host Table #152 listed 213 hosts. In May 1982, Host Table #166 listed 235.

The great switch to the domain system occurred on January 1, 1983. It was none too soon. The 562 hosts of August 1983 just wouldn't have been feasible under the older protocols and the older scheme.

Here's the growth over the next few years:

- 10/84 1024
- 10/85 1961
- 02/86 2308
- 11/86 5089
- 12/87 28174
- 07/88 33000
- 10/88 56000
- 01/89 80000
- 10/89 159000
- 10/90 313000
- 10/91 617000

I'll stop there for a while, for several reasons, not least because this marks the advent of both the Web and of Linux. But this marks several other things as well.

On the political front, the Department of Defense relaxed its notion that only government, academic and research sites could connect. This was partially the result of recognizing the expansion of private networks (like IBM's VNet and Prodigy) and of networks distributing news and mail (UUNET, Bitnet, Fidonet), as well as the recognition that the network of networks was already vastly larger than the US and its "allies."

On the "engineering" side, the advent of the desktop machine and the commercial modem, meant that individuals could have a computer at home and plug in to their telephone lines. All for under \$3000!

In 1971, BBN estimated 10 users per host. It was extrapolating from 617000 to one million and multiplying that gave the Internet Society the 10,000,000 users that TIME magazine claimed in 1992. In an informal poll a decade ago, I found that I knew individuals who owned a dozen domains and who were sole user on most (0.1 user/domain?) and that IBM's T.J. Watson Research Center had over 4000 users and one domain. I didn't know then and don't know now what the "correct" ratio is. Craig

Partridge estimated 5/domain in 1994. John Quarterman estimated 3.5. These sound more reasonable than double or triple those ratios.

The advent of NAT (Network Address Translation) makes all the ratios and estimates yet more unreliable: we have no way of determining in any accurate fashion how many desktops on a LAN are lurking behind a single address.

But in 2003 (the last dates for which the numbers seem feasible), there were about 175 million domains. Using the growth rates of 1995-2002 for 2002-2006 would mean 600-700 million domains at the end of 2006, and about 2,000,000,000 users worldwide. (This counts students in schools and people in libraries, of course.) Just under a third of the world's population has access to an Internet-connected device.

There were about two dozen in January 1970.

In 1990, when there were about 300,000 hosts, three students at McGill University in Montreal set about writing a search engine that would poll FTP archives on a regular basis. Alan Emtage, Bill Heelan and Peter Deutsch called itarchie. Soon there were other search engines, too: jughead, veronica, WAIS [Wide Area Information Server].

They were very useful, for a brief period of time. The University of Minnesota, where gopher was developed, wanted to profit from it. Tim Berners-Lee offered the Web free. Though he's now Sir Timothy, he's not a millionaire. But, by 1994, the World Wide Web had swept away all the other browsers.

Why is this important? Well, in 1973, when Dennis M. Ritchie and Ken Thompson gave the first UNIX paper, there were about 200 people in the room. There were just over 40 hosts on the Net. Word of mouth and then the CACM paper were how the word got out. When Linus Torvalds posted his note on the comp.os.minix newsgroup, there were about 200 groups. There were over 600,000 hosts connected to the Internet. The potential audience for Linux was enormous. And it was virtually instantaneous.

I wonder what my long-ago Toronto colleague, Marshall McLuhan, would say; today's Internet and Web are "hot" media far beyond his notions.

As I wrote in Chapter 17, try buying a candy bar that doesn't have a URL on it.

Some Errata

One of the "problems" of writing is that your readership can be quite notable. I have received clarifying comments from two of the major "participants."

With regard to Chapter 2 (UNIX), Dennis Ritchie has pointed out the following (as well as a few minor points):

1. "'AT&T Bell Labs' is correct, I guess, but BL didn't really start using that name until divestiture, and AT&T wanted the prominent branding. Of course we were a part of AT&T (via WEC), so it's not wrong."
2. "Doug wasn't heavily involved in Multics at the time, and he was always pursuing other things as well (including, course, being Ken's boss)."
3. "I don't think 'UNICS' was ever committed to paper... A couple of years ago I checked again with Brian and Peter N, and Brian admits to Unix, Peter denies it. Like you I thought that it must have been Peter because of the punsterish tendencies, and may have written that, but it seems not to be true. I suppose leaving the air of mystery has some virtue, though."
4. "Bob [Morris] retired several years ago from NCSC."

I also received two pieces of mail from Richard Stallman, one pointing out that: "It was only in senior year [in high school] that I was using a computer."

Moreover, I wrote "He [RMS] has frequently said that 'Software wants to be free'."

Richard writes: "I don't believe I ever said those particular words. ... What I say is that software should be free; that is to say, its users should have freedom."

I have received other comments and addenda from several notables (e.g. Rob Kolstad and Mike O'Dell). All will be incorporated into the next full version.

I am really gratified that the most eminent workers in the field have been reading my work and have taken the time to write to me.

PHS