



哈爾濱工業大學
HARBIN INSTITUTE OF TECHNOLOGY



操作系统

Operating Systems

刘川意 教授

liuchuanyi@hit.edu.cn

哈尔滨工业大学（深圳）

2021年10月



Module 6: I/O与存储

1. **I/O devices**
2. Hard Disk Drives
3. **RAID**
4. 其他I/O设备

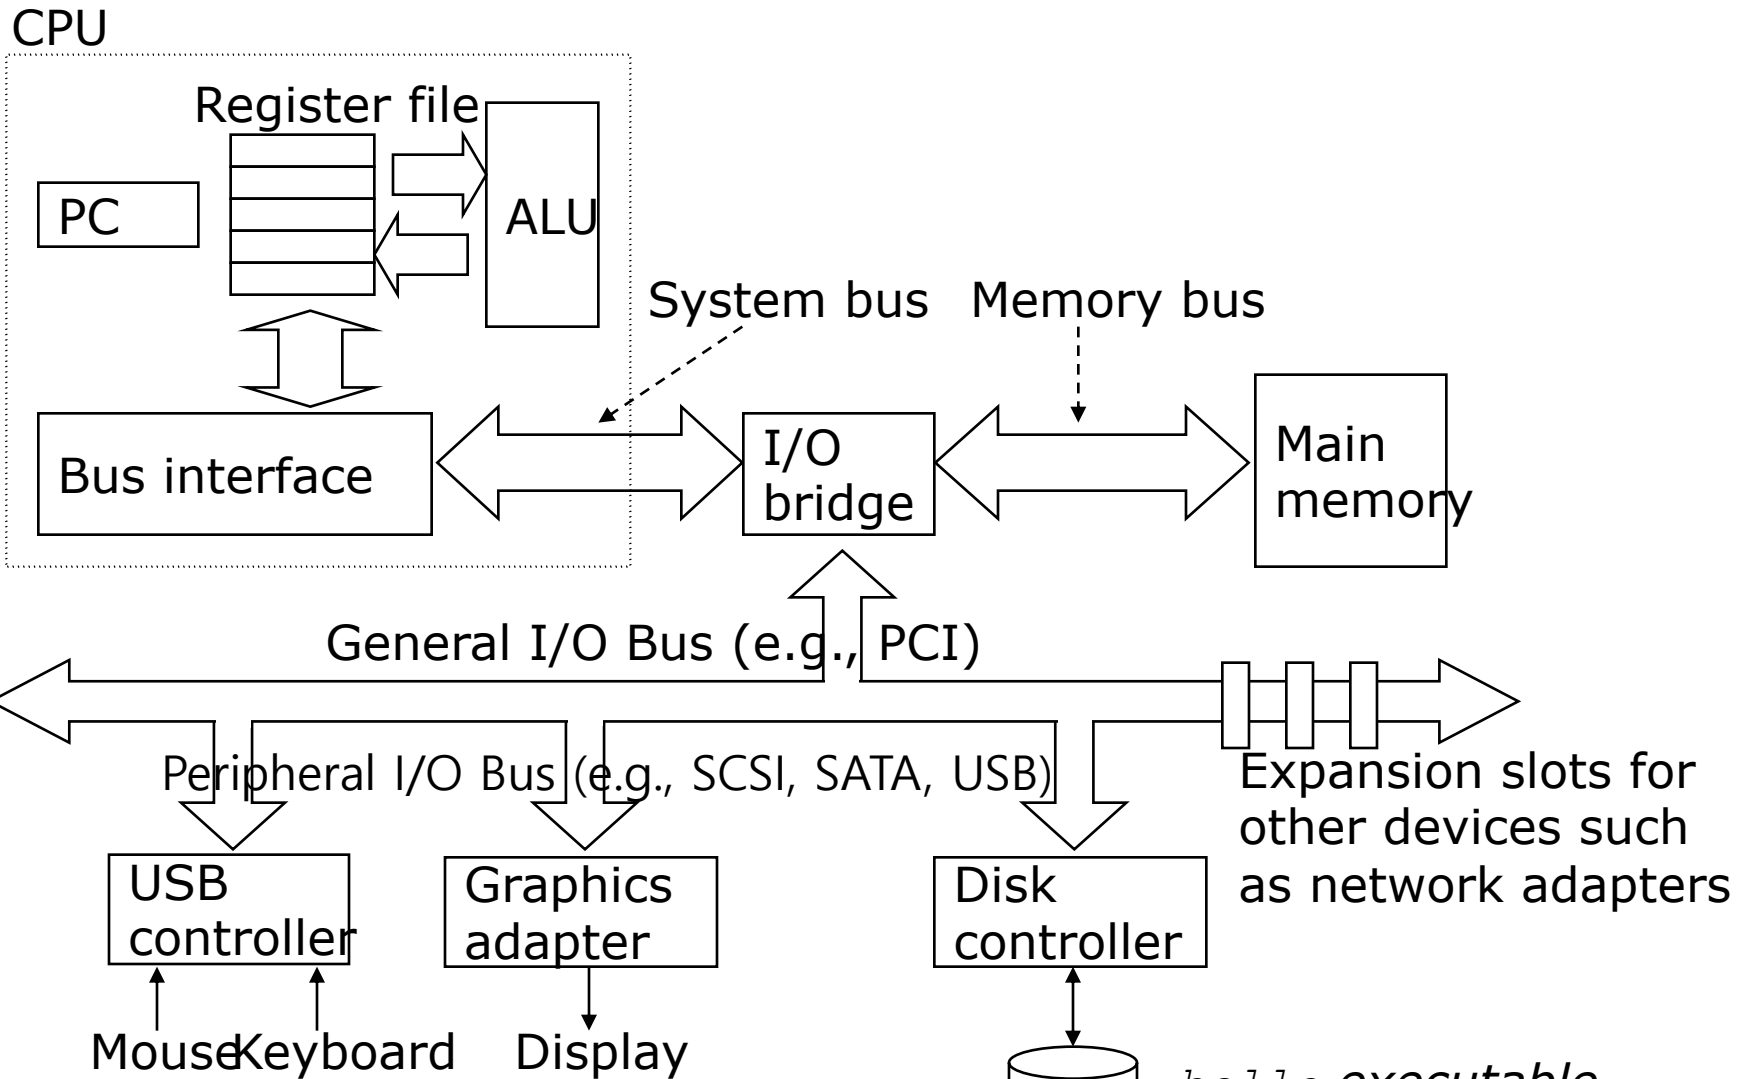


I/O Devices

- I/O is **critical** to computer system to **interact with systems**.
- I/O需要解决的重要问题：
 - How should I/O be integrated into systems?
 - What are the general mechanisms?
 - How can we make the efficiently?



Structure of input/output (I/O) device



CPU is attached to the main memory of the system via some kind of memory **bus**.

Some devices are connected to the system via a general **I/O bus**.

I/O Architecture

■ Buses

- Data paths that provided to enable information between CPU(s), RAM, and I/O devices. 不同的I/O设备(如键盘、鼠标、磁盘等)需要通过相应的接口电路与总线相连接, 这些接口电路由“控制器”或“适配器”提供(后面统称为“设备控制器”)。不同的设备控制器能够支持不同的接口协议

■ I/O bus

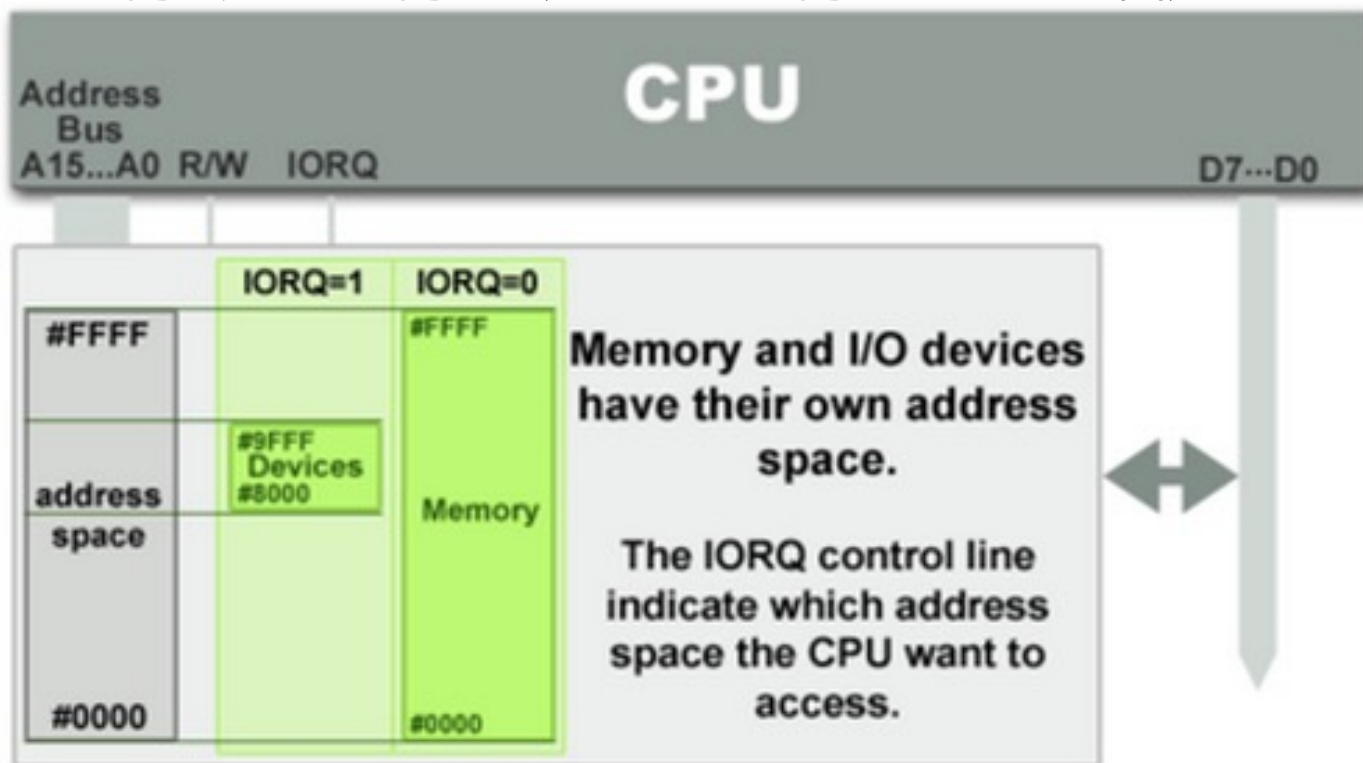
- Data path that connects a CPU to an I/O device.
- I/O bus is connected to I/O device by three hardware components: I/O ports, interfaces and device controllers.
- 根据接口协议的性能区别, 现代计算机对I/O总线进行了分层。在上图中, 图像或者其他高性能的I/O设备通过常规的I/O总线连接到系统, 在许多现代系统中会是PCI或它的衍生形式。而一些相对较慢的I/O设备则通过外围总线(peripheral bus)连接到系统, 比如使用SCSI、SATA或者USB等协议的I/O设备

Device interaction

- How the OS communicates with the **device**? 主机对I/O设备进行访问的目标是I/O设备的寄存器或者内存。常见的I/O设备都只提供寄存器供主机访问, 对于低速外设这样的模式是足够的, 但是对于需要大量、高速数据交互的外设(如显卡、网卡), 就需要主机能够直接访问外设的内存了
- Solutions: 现代计算机提供了两种方式来访问I/O设备, 它们分别是PMIO和MMIO
 - PMIO: 端口映射I/O (Port-mapped I/O)。将I/O设备独立看待, 并使用CPU提供的专用I/O指令访问; **I/O instructions**: a way for the OS to send data to specific device registers.
 - ▶ Ex) `in` and `out` instructions on x86
 - MMIO: 内存映射I/O (Memory-mapped I/O)。将I/O设备看作内存的一部分, 不使用单独的I/O指令, 而是使用内存读写指令访问; **memory-mapped I/O**
 - ▶ Device registers available as if they were memory locations.
 - ▶ The OS `load` (to read) or `store` (to write) to the device instead of main memory.

PMIO (Port-mapped I/O)

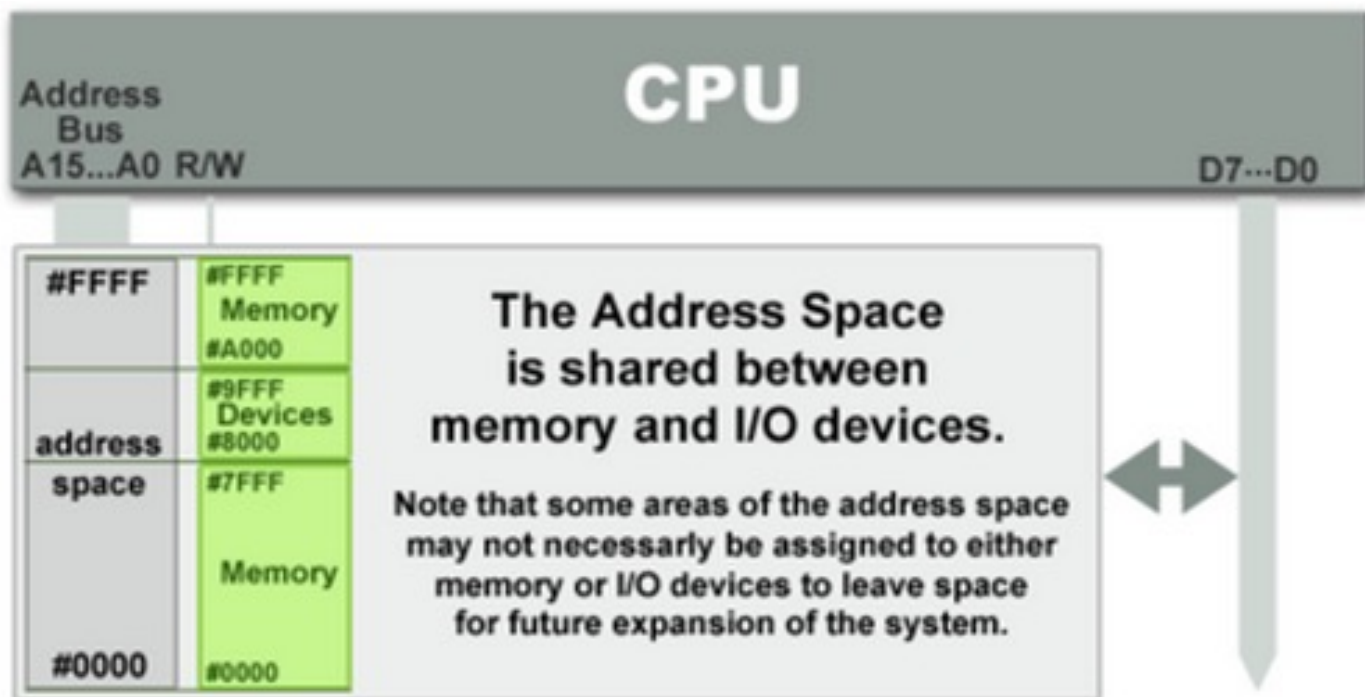
- 端口映射I/O, 又叫做被隔离的I/O (isolated I/O), 它提供了一个专门用于I/O设备“注册”的地址空间, 该地址空间被称为I/O地址空间, 最大寻址范围为64K



- 为了使I/O地址空间与内存地址空间隔离, 要么在CPU物理接口上增加一个I/O引脚, 要么增加一条专用的I/O总线。因此, 并不是所有的平台都支持PMIO, 常见的ARM平台就不支持PMIO。支持PMIO的CPU通常具有专门执行I/O操作的指令, 例如在Intel-X86架构的CPU中, I/O指令是in和out, 这两个指令可以读/写1、2、4个字节(outb, outw, outl)从内存到I/O接口上。由于I/O地址空间比较小, 因此I/O设备一般只在其中“注册”自己的寄存器, 之后系统可以通过PMIO对它们进行访问

MMIO (Memory-mapped I/O)

- 在MMIO中，物理内存和I/O设备共享内存地址空间(注意，这里的内存地址空间实际指的是内存的物理地址空间)



- 当CPU访问某个虚拟内存地址时，该虚拟地址首先转换为一个物理地址，**对该物理地址的访问，会通过南北桥(现在被合并为I/O桥)的路由机制被定向到物理内存或者I/O设备上**。因此，用于访问内存的CPU指令也可用于访问I/O设备，并且在内存(的物理)地址空间上，需要给I/O设备预留一个地址区域，该地址区域不能给物理内存使用。
- MMIO是应用得最为广泛的一种I/O方式，由于内存地址空间远大于I/O地址空间，I/O设备可以在内存地址空间上暴露自己的内存或者寄存器，以供主机进行访问

PCI设备

- PCI及其衍生的接口(如PCIE)主要服务于高速I/O设备(如显卡或网卡), 使用PCI接口的设备又被称为PCI设备。与慢速I/O设备不同, 计算机既需要访问它们的寄存器, 也需要访问它们的内存。
- 每个PCI设备都有一个配置空间(实际就是设备上一组连续的寄存器), 大小为256byte。配置空间中包含了6个BAR(Base Address Registers, 基址寄存器), BAR中记录了设备所需要的地址空间类型、基址以及其他属性

Memory Space BAR Layout

31 - 4	3	2 - 1	0
16-Byte Aligned Base Address	Prefetchable	Type	Always 0

I/O Space BAR Layout

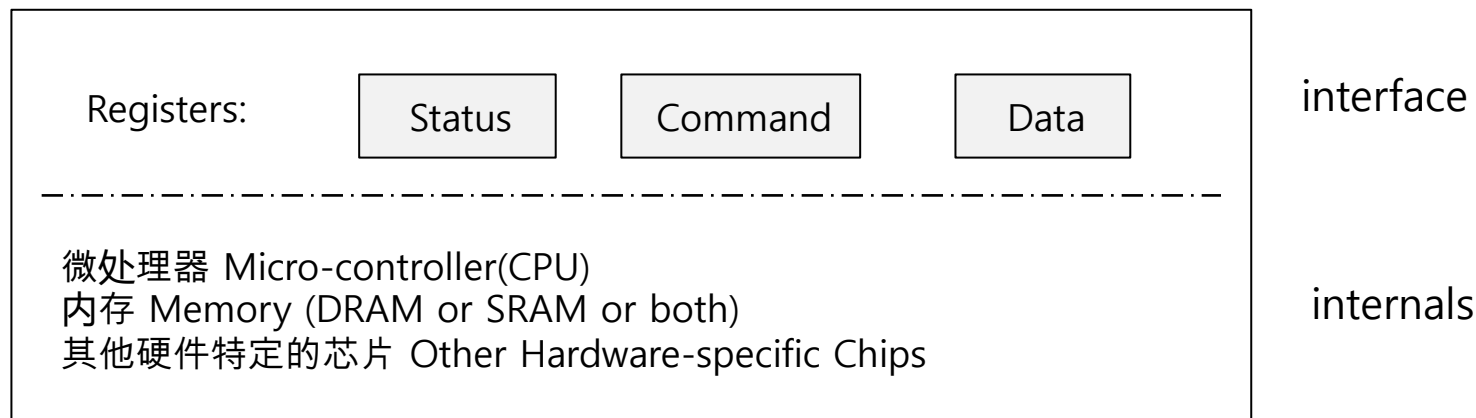
31 - 2	1	0
4-Byte Aligned Base Address	Reserved	Always 1

PCI设备 (cont.)

- 可以看到, PCI设备能够申请两类地址空间, 即内存地址空间和I/O地址空间, 它们用BAR的最后一位区别开来。因此, PCI设备可以通过PMIO和MMIO将自己的I/O存储器(Registers/RAM/ROM)暴露给CPU(通常寄存器使用PMIO, 而内存使用MMIO的方式暴露)。
- 配置空间中的每个BAR可以映射一个地址空间, 因此每个PCI设备最多能映射6段地址空间, 但实际上很多设备用不了这么多。PCI配置空间的初始值是由厂商预设在设备中的, 也就是说, 设备需要哪些地址空间都是其自己定的, 这可能会造成不同的PCI设备所映射的地址空间冲突, 因此在PCI设备枚举(也叫总线枚举, 由BIOS或者OS在启动时完成)的过程中, 会重新为其分配地址空间, 然后写入PCI配置空间中。
- 在PCI总线之前的ISA总线是使用跳线帽来分配外设的物理地址, 每插入一个新设备都要改变跳线帽以分配物理地址, 这是十分麻烦且易错的, 但这样的方式似乎我们更容易理解。能够分配自己总线上挂载设备的物理地址这也是PCI总线相较于I2C、SPI等低速总线一个最大的特色

标准外设(Canonical Device)

- Canonical Devices has two important components.
 - **Hardware interface** allows the system software to control its operation. 硬件接口本质就是I/O设备提供的各式寄存器，系统软件通过与这些寄存器进行交互，达到控制I/O设备的目的
 - **Internals** which is implementation specific. 实现硬件接口提供的功能，不同的I/O设备具有不同功能，因此它们的内部实现和包含的元器件也不尽相同
 - 使用I/O设备的目的是为了交互数据，不管是网卡、磁盘，亦或是键盘，总归要将数据进行输入输出



Canonical Device

Hardware interface of Canonical Device

■ status register

- See the current status of the device

■ command register

- Tell the device to perform a certain task

■ data register

- Pass data to the device, or get data from the device

**By reading and writing above three registers,
the operating system can control device behavior.**

■ Typical interaction example

```
while ( STATUS == BUSY)
    ; //wait until device is not busy
write data to data register
write command to command register
    (Doing so starts the device and executes the command)
while ( STATUS == BUSY)
    ; //wait until device is done with your request
```

Polling

- Operating system waits until the device is ready by **repeatedly** reading the status register. **Device driver通过轮询读取设备状态**
- 一般来说, 主机与I/O设备要进行数据交互, 会经过这样一个过程:
 - CPU通过I/O设备的硬件接口(以下简称I/O接口)获取设备状态(即状态寄存器的值), 只有“就绪”状态的设备才能进行数据传输。
 - CPU通过I/O接口下达交互指令: 如果是读数据, 则向I/O接口的命令寄存器输入要获取的数据在I/O设备的内部位置以及读设备指令; 如果是写数据, 则向I/O接口的命令寄存器输入要存放的数据在I/O设备的内部位置、写设备指令, 以及向数据寄存器写入数据。
 - I/O设备内部根据I/O接口中寄存器的值, 开始执行数据传输工作。
 - CPU在I/O设备完成工作后, 执行其他操作, 完成数据传送。

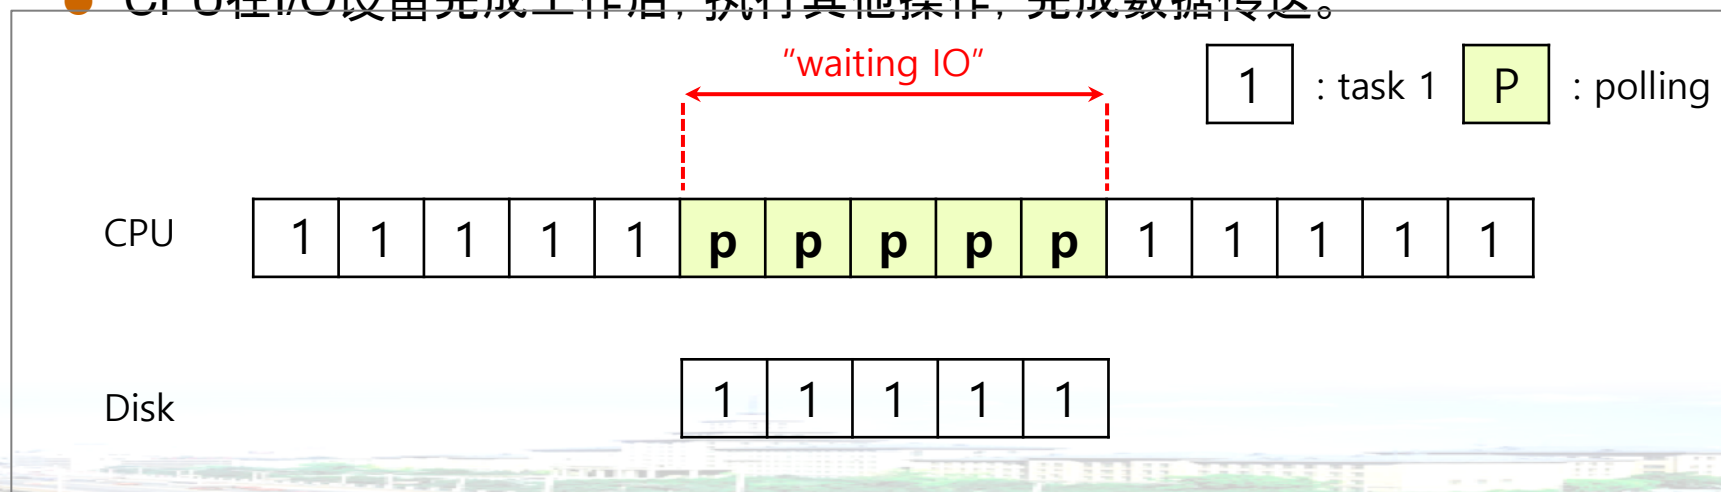



Diagram of CPU utilization by polling

Polling Evaluation

- 标准交互流程实现起来比较简单, 但是难免会有一些低效和不方便。第一个问题就是轮询过程比较低效, 在等待设备是否满足某种状态时浪费大量CPU时间(下图描述的就是磁盘在执行数据传输过程中, CPU不能执行其他任务, 只能等待传输完成), 如果此时操作系统可以切换执行下一个就绪进程, 就可以大大提高CPU的利用率。
 - Positive aspect is simple and working.
 - **However, it wastes CPU time just waiting for the device.**
 - ▶ Switching to another ready process is better utilizing the CPU.

interrupts

- 有了中断机制, CPU向设备发出I/O请求后, 就可以让对应进程进入睡眠等待, 从而切换执行其他进程。当设备完成I/O请求后, 它会抛出一个硬件中断, 引发CPU跳转执行操作系统预先定义好的中断处理程序, 中断处理程序会挂起正在执行的进程, 同时唤醒等待I/O的进程并继续执行
- **Put the I/O request process to sleep** and context switch to another. **OS调度程序让执行I/O request的进程进入sleep状态, 并切换成另一个进程去执行**
- When the device is finished, wake the process waiting for the I/O by **interrupt**. **外设通过中断机制唤醒处于sleep状态的等待进程**
 - Positive aspect is allow to **CPU and the disk are properly utilized.**



在磁盘执行进程1的I/O过程中, CPU同时执行进程2, 并且在I/O请求执行完毕后, 回过头来再次执行进程1

Polling vs interrupts

- However, “interrupts is not always the best solution”
 - If, device performs very quickly, interrupt will “slow down” the system.
 - Because **context switch is expensive** (switching to another process)

If a device is fast → **poll** is best.
If it is slow → **interrupts** is better.

CPU is once again over-burdened

- CPU wastes a lot of time to copy the *a large chunk of data* from memory to the device.

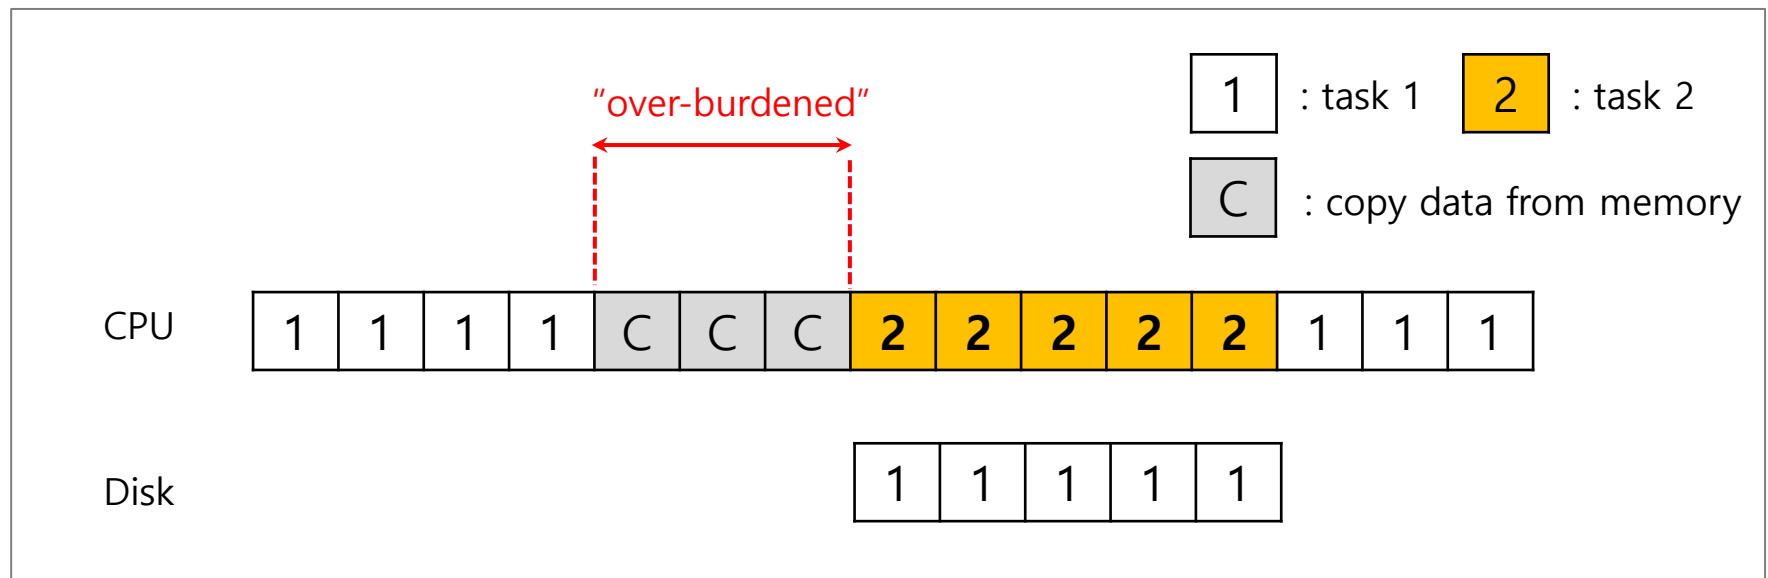


Diagram of CPU utilization

DMA (Direct Memory Access)

- **Copy data** in memory by knowing “where the data lives in memory, how much data to copy”
- When completed, DMA raises an interrupt, I/O begins on Disk.




Diagram of CPU utilization by DMA

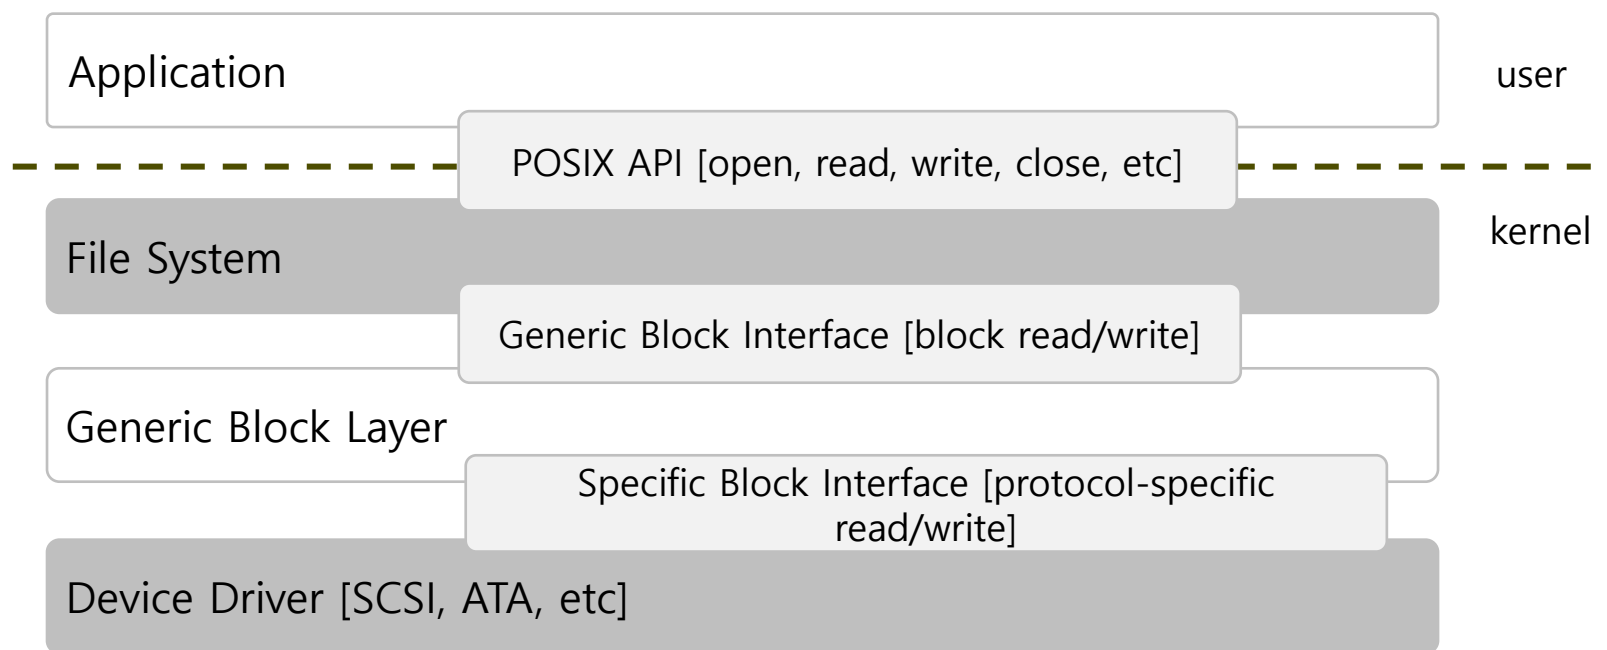
Device interaction (Cont.)

- How the OS interact with **different specific interfaces**?
 - Ex) We'd like to build a file system that worked on top of SCSI disks, IDE disks, USB keychain drivers, and so on.
- Solutions: **Abstraction**
 - Abstraction encapsulate **any specifics of device interaction**.



File system Abstraction

- File system **specifics** of which disk class it is using.
 - Ex) It issues **block read** and **write** request to the generic block layer.



The File System Stack

Problem of File system Abstraction

- If there is a device having many special capabilities, these capabilities **will go unused** in the generic interface layer.
- Over 70% of OS code is found in device drivers.
 - Any device drivers are needed because you might plug it to your system.
 - They are primary contributor to **kernel crashes**, making **more bugs**.

A Simple IDE Disk Driver

- Four types of register
 - Control, command block, status and error
 - Memory mapped IO
 - in and out I/O instruction



IDE Device Interface

■ Control Register:

Address 0x3F6 = 0x80 (0000 1RE0): R=reset, E=0 means "enable interrupt"

■ Command Block Registers:

Address 0x1F0 = Data Port

Address 0x1F1 = Error

Address 0x1F2 = Sector Count

Address 0x1F3 = LBA low byte

Address 0x1F4 = LBA mid byte

Address 0x1F5 = LBA hi byte

Address 0x1F6 = 1B1D TOP4LBA: B=LBA, D=drive

Address 0x1F7 = Command/status

■ Status Register (Address 0x1F7):

7	6	5	4	3	2	1	0
BUSY	READY	FAULT	SEEK	DRQ	CORR	IDDEX	ERROR

■ Error Register (Address 0x1F1): (check when Status ERROR==1)

7	6	5	4	3	2	1	0
BBK	UNC	MC	IDNF	MCR	ABRT	T0NF	AMNF

● BBK = Bad Block

● UNC = Uncorrectable data error

● MC = Media Changed

● IDNF = ID mark Not Found

● MCR = Media Change Requested

● ABRT = Command aborted

● T0NF = Track 0 Not Found

● AMNF = Address Mark Not Found

OS跟设备交互的典型协议

- **Wait for drive to be ready.** Read Status Register (0x1F7) until drive is not busy and READY.
- **Write parameters to command registers.** Write the sector count, logical block address (LBA) of the sectors to be accessed, and drive number (master=0x00 or slave=0x10, as IDE permits just two drives) to command registers (0x1F2-0x1F6).
- **Start the I/O.** by issuing read/write to command register. Write READ—WRITE command to command register (0x1F7).
- **Data transfer (for writes):** Wait until drive status is READY and DRQ (drive request for data); write data to data port.
- **Handle interrupts.** In the simplest case, handle an interrupt for each sector transferred; more complex approaches allow batching and thus one final interrupt when the entire transfer is complete.
- **Error handling.** After each operation, read the status register. If the ERROR bit is on, read the error register for details.

Wait for drive to be ready

```
1. static int ide_wait_ready()  
2. /* ensure the drive is ready before issuing a request to it */  
3. {  
4.     while (((int r = inb(0x1f7)) & IDE_BSY) ||  
5.            !(r & IDE_DRDY))  
6.         ; // loop until drive isn't busy  
7. }
```

Write parameters to command registers

```
1. static void ide_start_request(struct buf *b)
2. /* send a request (and perhaps data, in the case of a write) to the disk, in
   and out x86 instructions are called to read and write device registers */
3. {
4.     ide_wait_ready();
5.     outb(0x3f6, 0); // generate interrupt
6.     outb(0x1f2, 1); // how many sectors?
7.     outb(0x1f3, b->sector & 0xff); // LBA goes here ...
8.     outb(0x1f4, (b->sector >> 8) & 0xff); // ... and here
9.     outb(0x1f5, (b->sector >> 16) & 0xff); // ... and here!
10.    outb(0x1f6, 0xe0 | ((b->dev&1)<<4) | ((b->sector>>24)&0x0f));
11.    if(b->flags & B_DIRTY)
12.    {
13.        outb(0x1f7, IDE_CMD_WRITE); // this is a WRITE
14.        outsl(0x1f0, b->data, 512/4); // transfer data too!
15.    }
16.    else
17.    {
18.        outb(0x1f7, IDE_CMD_READ); // this is a READ (no data)
19.    }
20.}
```

IO interface

```
1. void ide_rw(struct buf *b)
2.   // queues a request (if there are others pending)
3.   // or issues it directly to the disk (via ide_start_request())
4. {
5.   acquire(&ide_lock);
6.   for (struct buf **pp = &ide_queue; *pp; pp=&(*pp)->qnext) ; // walk queue
7.   *pp = b; // add request to end
8.   if (ide_queue == b) // if q is empty
9.     ide_start_request(b); // send req to disk
10.  while ((b->flags & (B_VALID|B_DIRTY)) != B_VALID)
11.    sleep(b, &ide_lock); // wait for completion
12.  release(&ide_lock);
13. }
```

Handle interrupts

```
1. void ide_intr()  
2. /* invoked when an interrupt takes place; it reads data from the device (if  
   the request is a read, not a write), wakes the process waiting for the I/O to  
   complete, and (if there are more requests in the I/O queue), launches the next  
   I/O via ide_start_request() */  
3. {  
4.     struct buf *b;  
5.     acquire(&ide_lock);  
6.     if (!(b->flags & B_DIRTY) && ide_wait_ready(1) >= 0)  
7.         insl(0x1f0, b->data, 512/4); // if READ: get data  
8.     b->flags |= B_VALID;  
9.     b->flags &= ~B_DIRTY;  
10.    wakeup(b); // wake waiting process  
11.    if ((ide_queue = b->qnext) != 0) // start next request  
12.        ide_start_request(ide_queue); // (if one exists)  
13.    release(&ide_lock);  
14. }
```

总结一些I/O系统要完成的工作!



```
write(buf, 10);
```

OS需要提供系统调用接口

```
DMA.addr = buf;  
DMA.count = 10;  
.....  
sleep_on(Disk);
```

查一下手册就可以找到该写什么命令?该向哪里写?

让出**CPU**?

需要写中断处理程序!

```
do_write_end()//中断处理  
{  
    wakeup(Disk);  
}
```

■ **总的感觉:** 很简单
处理流程是很简单, 复杂的是一些**细节问题**, 如滚屏



I/O设备管理总结

- 如何实现交互? \Rightarrow 首先需要了解I/O的工作原理
- 从用户如何I/O开始 \Rightarrow 用户发送一个命令(read)
- 系统调用read \Rightarrow 被展开成给一些寄存器发送命令的代码
- 发送完命令以后... \Rightarrow CPU轮询, CPU干其它事情并等中断
- 中断方案最常见 \Rightarrow 相比其他设备, CPU太快了
- 实现独享设备的共享 \Rightarrow 假脱机系统 (SPOOLING)



Module 6: I/O与存储

1. I/O devices
2. **Hard Disk Drives**
3. **RAID**
4. 其他I/O设备



Tape is Dead Disk is Tape Flash is Disk RAM Locality is King



Jim Gray

Microsoft

认识一下磁盘



认识一下磁盘



机械臂杆




盘片高速旋转产生气流非常强，足以使磁头托起，并与盘面保持一个微小的距离。

现在的水平已经达到 $0.005\mu\text{m} \sim 0.01\mu\text{m}$ ，这只是人类头发直径的千分之一。


认识一下磁盘

- 画一个示意图：



- 所以，磁盘被称为块设备！

- 看看俯视图：




扇区是磁盘的寻址单位、访问单位

- 磁盘的数据单位是扇区
- 扇区大小：512字节
- 扇区的大小是传输时间和碎片浪费的折衷

Interface

- The only guarantee is that a single 512-byte write is **atomic**.
- Multi-sector operations are possible.
 - Many file systems will read or write 4KB at a time.
 - **Torn write:**
 - ▶ If an untimely power loss occurs, only a portion of a larger write may complete.
- Accessing blocks in **a contiguous chunk** is the fastest access mode.
 - A sequential read or write
 - Much faster than any more random access pattern.

Basic Geometry



A Disk with Just A Single Track (12 sectors)

- **Platter** (Aluminum coated with a thin magnetic layer)
 - A circular hard surface
 - Data is stored persistently by inducing magnetic changes to it.
 - Each platter has 2 sides, each of which is called a **surface**.

Basic Geometry (Cont.)

■ Spindle


- Spindle is connected to a motor that spins the platters around.
- The rate of rotations is measured in **RPM** (Rotations Per Minute).
 - ▶ Typical modern values : 7,200 RPM to 15,000 RPM.
 - ▶ E.g., 10000 RPM : A single rotation takes about 6 ms.

■ Track

- Concentric circles of sectors
- Data is encoded on each surface in a track.
- A single surface contains many thousands and thousands of tracks.




A Simple Disk Drive



■ Disk head (One head per surface of the drive)

- The process of *reading* and *writing* is accomplished by the **disk head**.
- Attached to a single disk arm, which moves across the surface.


Single-track Latency: The Rotational Delay



A Single Track Plus A Head

- **Rotational delay:** Time for the desired sector to rotate
 - Ex) Full rotational delay is R and we start at sector 6
 - ▶ Read sector 0: Rotational delay = $\frac{R}{2}$
 - ▶ Read sector 5: Rotational delay = $R-1$ (worst case)

Multiple Tracks: Seek Time



Three Tracks Plus A Head (Right: With Seek)
(e.g., read to sector 11)

- **Seek:** Move the disk arm to the correct track
 - **Seek time:** Time to move head to the track contain the desired sector.
 - One of the most costly disk operations.

Phases of Seek

■ Acceleration → Coasting → Deceleration → Settling

- **Acceleration:** The disk arm gets moving.
- **Coasting:** The arm is moving at full speed.
- **Deceleration:** The arm slows down.
- **Settling:** The head is *carefully positioned* over the correct track.
 - ▶ The settling time is often quite significant, e.g., 0.5 to 2ms.

Transfer

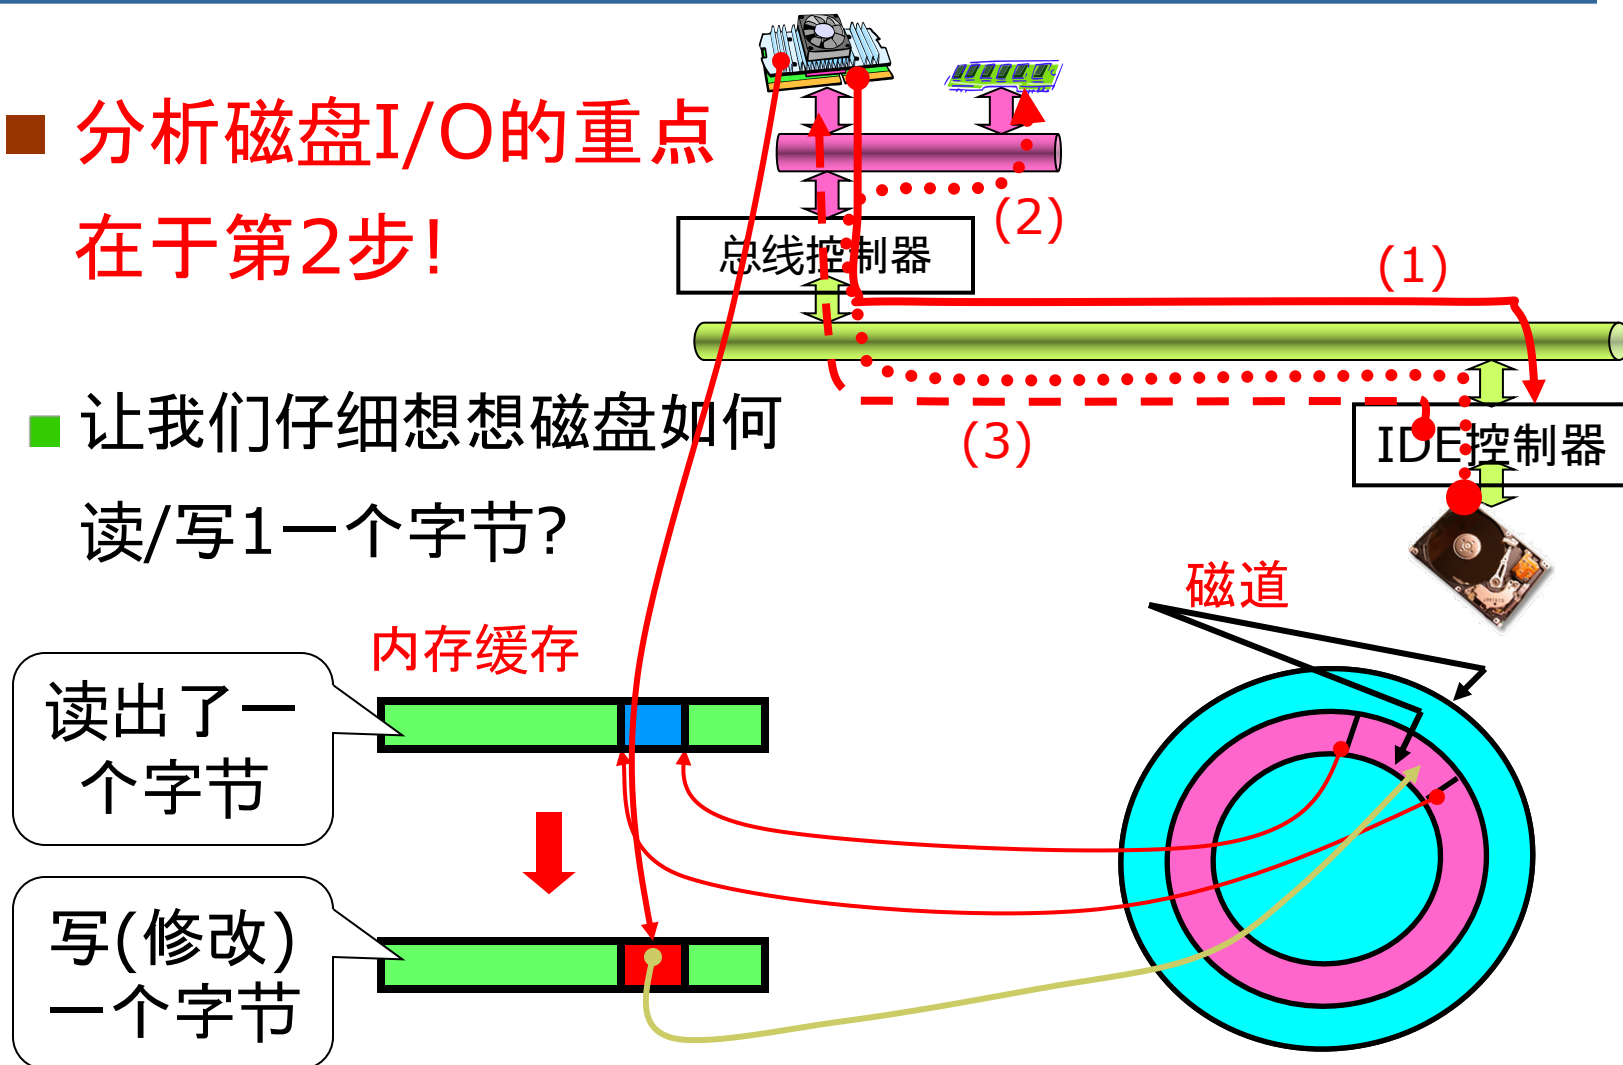
- The final phase of I/O
 - Data is either *read from* or *written* to the surface.
- Complete I/O time:
 - **Seek**
 - Waiting for the **rotational delay**
 - **Transfer**



磁盘的I/O

- 分析磁盘I/O的重点在于第2步!

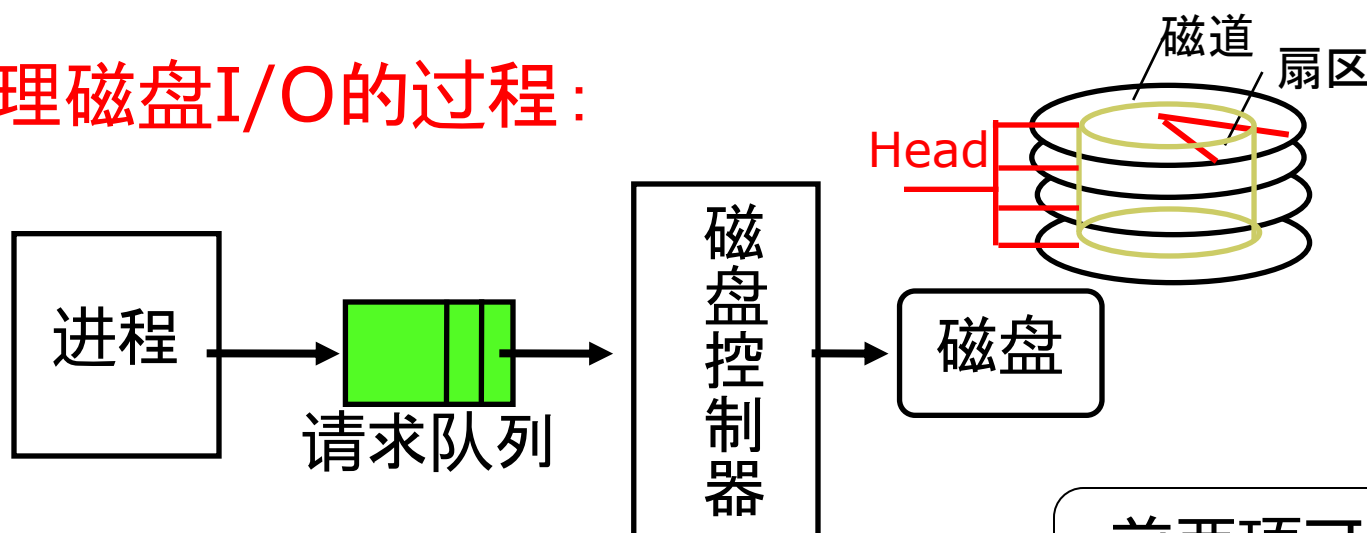
- 让我们仔细想想磁盘如何读/写1一个字节?



- 磁盘I/O: 缓存队列 → 控制器 → 寻道 → 旋转 → 传输!

磁盘I/O的分析

■ 整理磁盘I/O的过程：



■ 我们最关心的磁盘什么时候读/写完？

前两项可以忽略！

磁盘访问延迟 = 队列时间 + 控制器时间 +
寻道时间 + 旋转时间 + 传输时间

3ms to
12 ms

(半周): 4
ms to 8 ms

约
0.25ms

■ 关键所在：最小化寻道时间和旋转延迟！

I/O过程是解开许多磁盘问题的钥匙

■ 磁盘调度：

磁盘访问延迟 = 队列时间 + 控制器时间 +
寻道时间 + 旋转时间 + 传输时间

前两项可以忽略！

12 ms to 8 ms

8 ms to 4 ms

约0.25ms

■ 多个磁盘访问请求出现在请求队列怎么办？ **调度**

■ 调度的目标是什么？调度时主要考察什么？

目标当然是平均
访问延迟小！

寻道时间是主要
矛盾！

■ **磁盘调度：输入多个磁道请求，给出服务顺序！**



FCFS磁盤調度


$$130+146+85+108+110+59+2=640$$

■ 最直觀、最公平的調度：

■ 一個實例：磁頭開始磁道位置=53，

請求隊列=98, 183, 37, 122, 14, 124, 65,

FCFS：磁頭共
移動640磁道！



在移動過程中把經過的請求處理了？！

183-53=130

122-37=85

124-14=110

183-37=146

122-14=108

124-65=59

67-65=2



磁頭在長途奔襲！

SSTF磁盘调度

■ Shortest-seek-time First最短寻道时间优先:

■ 继续该实例: 磁头开始位置=53;

请求队列=98, 183, 37, 122, 14, 124, 65, 67



SSTF: 磁头共移动
236(14+53+169)
磁道, 要少很多!

如果在处理183之前
又来一些中间磁道
的请求, 则...


■ SSTF存在饥饿问题

SCAN磁盘调度(扫描/电梯算法)

■ SSTF+中途不回折: 每个请求都有处理机会

■ 继续该实例: 磁头开始位置=53;

请求队列=98, 183, 37, 122, 14, 124, 65, 67



SCAN: 磁头共移动
 $53 + 183 = 236$ 磁道,
和SSTF一样!

这些请求的等待时间较长, 只因所在方向不够幸运!

根据其特征,
SCAN也被称为
电梯算法!


■ SCAN导致延迟不均

C-SCAN磁盘调度

■ SCAN+直接移到另一端: 两端请求都能很快处理

■ 继续该实例: 磁头开始位置=53;

请求队列=98, 183, 37, 122, 14, 124, 65, 67



CSCAN中的Circular
是环的意思!

CSCAN: 磁头共移
动 $53+199+134$ 磁
道!

其中199会较快!


■ $14 \rightarrow 0$ ($183 \rightarrow 199$)
没有必要

C-LOOK磁盤調度

■ CSCAN+看一看：前面沒有請求就回移

■ 繼續該實例：磁頭開始位置=53；

請求隊列=98, 183, 37, 122, 14, 124, 65, 67






■ LOOK和C-LOOK是比较合理的缺省算法

操作系统中所有的算法都要因地制宜！



C-LOOK磁盤調度

- 繼續該實例：磁頭開始位置=53；
請求隊列=98, 183, 37, 122, 14, 124, 65, 67



1) 磁道請求隊列的形式

2) 新磁道請求如何入隊列

$C[i+1] < X < c[i]$ 或者
 $X > C[i+1] > c[i]$

53-37-14- 183-124-122-98-67-
65

磁盘编址

如何管理磁盘， 首先对磁盘的扇区进行编号！

- 出厂的磁盘需要低级格式化(物理格式化):
将连续的磁性记录材料分成物理扇区
- 扇区 = 头 + 数据区 + 尾
- 头、尾中包含只有磁盘控制器能识别的扇区号码和纠错信息

什么是磁盘的逻辑格式化？
第12章 文件系统！

I/O过程是解开许多磁盘问题的钥匙

- **磁盘寻址**: 对于内存, 我们往往更关心存放内容的地址

- 实际上就是扇区怎么编址?


- 显然这个地址是(盘面(H) + 磁道(C) + 扇区(S))

- 寻道和旋转费时多 \Rightarrow 花最少时间访问最多扇区的方案: **磁臂不动、磁盘旋转一周, 访问磁头遇到的所有扇区。**


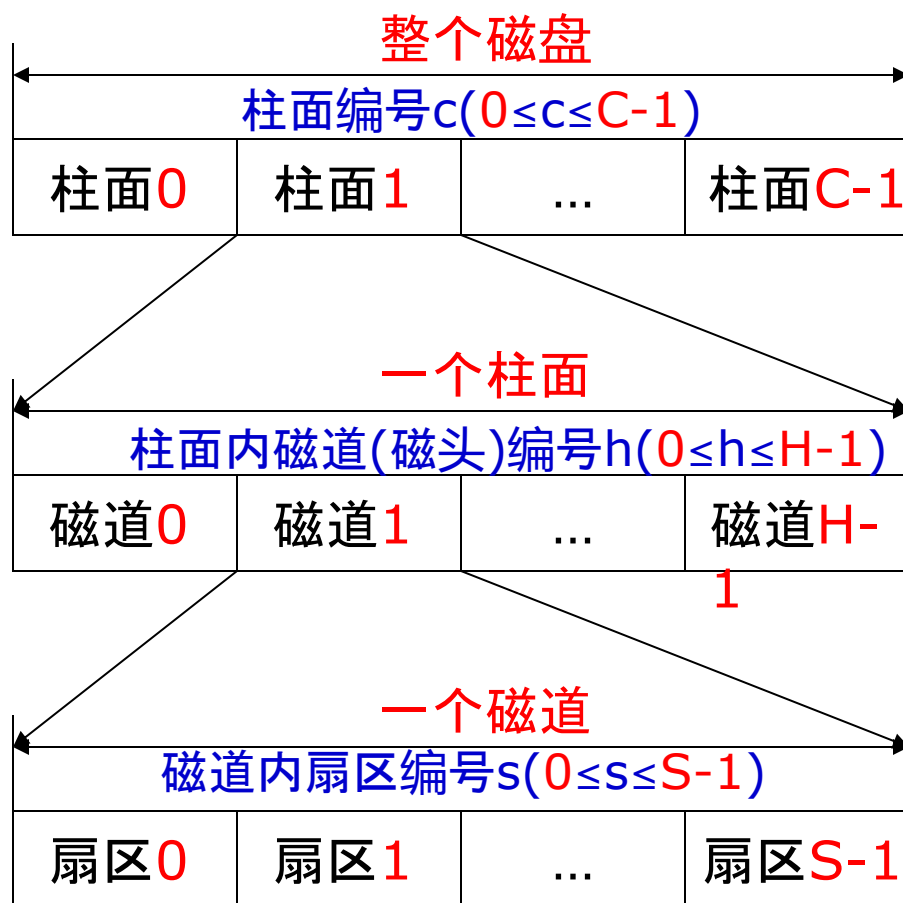
让这些扇区的编址邻近:
因为局部性!

- 扇区编址(1): **CHS(Cylinder/Head/Sector)**

- 扇区编址(2): **扇区编号(Logical Block Addressing LBA)**



扇区编号—现代磁盘的常见寻址方式



体现了局部性!

- **扇区编号**, 按照 (C, H, S) 将扇区形成一维扇区数组, 数组索引就是扇区编号

某扇区 (c, h, s) 编号 $A = c * H * S + h * S + s$ 扇区总数 = $C * H * S$
 已知 A, S , 则 $s = A \% S; h = [A / S] \% H; c = [A / (H * S)]$

扇区编号—现代磁盘的常见寻址方式

- chs(Cylinder/Head/Sector)模式
- 以前, 硬盘的容量还非常小, 采用与软盘类似的结构生产硬盘.
- 也就是**硬盘盘片的每一条磁道都具有相同的扇区数**
- 由此产生了所谓的3D参数 (Disk Geometry).:
- 磁柱面数(Cylinders), 头数(Heads), 扇区数(Sectors per track), 以及相应的寻址方式.

扇区编号—现代磁盘的常见寻址方式

- chs(Cylinder/Head/Sector)模式
- 磁头数(Heads) 表示硬盘总共有几个磁头,也就是有几面盘片, 最大为 256 (用 8 个二进制位存储);
- 柱面数(Cylinders) 表示硬盘每一面盘片上有几条磁道, 最大为 1024(用 10 个二进制位存储);
- 扇区数(Sectors per track) 表示每一条磁道上有几个扇区, 最大为 63 (用 6 个二进制位存储).
- 每个扇区一般是 512个字节;
- 所以磁盘最大容量为:
- $256 * 1024 * 63 * 512 / 1048576 = 8064 \text{ MB}$

扇区编号—现代磁盘的常见寻址方式

- chs(Cylinder/Head/Sector)模式
- 这种方式会浪费很多磁盘空间 (与软盘一样)
- 为了进一步提高硬盘容量, 产生了等密度结构硬盘, 外圈磁道的扇区比内圈磁道多.
- 采用这种结构后, 硬盘不再具有实际的3D参数, 寻址方式也改为线性寻址, 即以扇区为单位进行寻址
- 为了与使用chs寻址的兼容 (如使用BIOS Int13H接口的软件), 在硬盘控制器内部安装了一个地址翻译器, 由它负责将老式3D参数翻译成新的线性参数.

IDE硬盘控制器的寄存器

- 有一组命令寄存器组(Task File Registers), I/O的端口地址为1F0H~1F7H
 - 1F2H 扇区计数寄存器
 - 1F3H 扇区号, 或LBA块地址0~7
 - 1F4H 柱面数低8位, 或LBA块地址8~15
 - 1F5H 柱面数高8位, 或LBA块地址16~23
 - 1F6H 驱动器/磁头, 或LBA块地址24~27
 - 1F7H 状态寄存器 命令寄存器
-
- CHS或LBA在磁头寄存器中指定



想一想.....磁盘驱动应如何实现？

```
do_hd = intr_addr; // do_hd 函数会在中断程序中被调用。
outb_p(hd_info[drive].ctl, HD_CMD); // 向控制寄存器输出控制字节。
port=HD_DATA; // 置 dx 为数据寄存器端口 (0x1f0)。
outb_p(hd_info[drive].wpcom>>2, ++port); // 参数：写预补偿柱面号 (需除 4)。
outb_p(nsect, ++port); // 参数：读/写扇区总数。
outb_p(sect, ++port); // 参数：起始扇区。
outb_p(cyl, ++port); // 参数：柱面号低 8 位。
outb_p(cyl>>8, ++port); // 参数：柱面号高 8 位。
outb_p(0xA0 | (drive<<4) | head, ++port); // 参数：驱动器号+磁头号。
outb(cmd, ++port); // 命令：硬盘控制命令。
```

Linux 0.11 下实现磁盘读写驱动片段



磁盘速度与内存速度的差异

- 1) 磁盘往往不是严格按需读取, 而是每次都会预读, 即使只需要一个字节, 磁盘也会从这个位置开始, 顺序向后读取**一定扇区长度**的数据放入内存。
- 2) 这样做的理论依据是计算机科学中著名的**局部性原理**: 当一个数据被用到时, 其附近的数据也通常会马上被使用。



Module 6: I/O与存储

1. I/O devices
2. Hard Disk Drives
3. **RAID**
4. 其他I/O设备



RAID (Redundant Array of Inexpensive Disks)

- **Use multiple disks** in concert to build a **faster**, **bigger**, and more **reliable** disk system.
 - RAID just looks like a big disk to the host system.
- Advantage
 - **Performance & Capacity**: Using multiple disks in parallel
 - **Reliability**: RAID can tolerate the loss of a disk.

RAIDs provide these advantages **transparently** to systems that use them.

RAID Interface

- When a RAID receives I/O request,
 1. The RAID **calculates** which disk to access.
 2. The RAID **issue** one or more **physical I/Os** to do so.

- RAID example: A mirrored RAID system
 - Keep two copies of each block (each one on a separate disk)
 - Perform two physical I/Os for every one logical I/O it is issued.

RAID Internals

- A microcontroller
 - Run firmware to direct the operation of the RAID
- Volatile memory (such as DRAM)
 - Buffer data blocks
- Non-volatile memory
 - Buffer writes safely
- Specialized logic to perform parity calculation

Fault Model

- RAIDs are designed to **detect** and **recover** from certain kinds of disk faults.
- **Fail-stop** fault model
 - A disk can be in one of two states: *Working* or *Failed*.
 - ▶ Working: all blocks can be read or written.
 - ▶ Failed: the disk is permanently lost.
 - RAID controller can immediately observe when a disk has failed.

How to evaluate a RAID

■ Capacity

- How much useful capacity is available to systems?

■ Reliability

- How many disk faults can the given design tolerate?

■ Performance



RAID Level 0: Striping

- RAID Level 0 is the simplest form as **striping** blocks.
 - **Spread the blocks** across the disks in a round-robin fashion.
 - No redundancy
 - Excellent performance and capacity

Disk 0	Disk 1	Disk 2	Disk 3
0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

-----> Stripe
(The blocks in the same row)

RAID-0: Simple Striping
(Assume here a 4-disk array)

RAID Level 0 (Cont.)

- Example) RAID-0 with a bigger chunk size
 - Chunk size : 2 blocks (8 KB)
 - A Stripe: 4 chunks (32 KB)

Disk 0	Disk 1	Disk 2	Disk 3	chunk size: 2blocks
0	2	4	6	
1	3	5	7	
8	10	12	14	
9	11	13	15	

Striping with a Bigger Chunk Size

Chunk Sizes

- Chunk size mostly affects performance of the array
 - **Small chunk size**
 - ▶ Increasing the parallelism
 - ▶ Increasing positioning time to access blocks
 - **Big chunk size**
 - ▶ Reducing intra-file parallelism
 - ▶ Reducing positioning time

**Determining the “best” chunk size is hard to do.
Most arrays use larger chunk sizes (e.g., 64 KB)**

RAID Level 0 Analysis

N : the number of disks

- **Capacity** → RAID-0 is perfect.
 - Striping delivers N disks worth of useful capacity.
- **Performance** of striping → RAID-0 is excellent.
 - All disks are utilized often in parallel.
- **Reliability** → RAID-0 is bad.
 - Any disk failure will lead to data loss.

Evaluating RAID Performance

- Consider two performance metrics
 - Single request latency
 - Steady-state throughput

- Workload
 - **Sequential**: access 1MB of data (block (B) ~ block (B + 1MB))
 - **Random**: access 4KB at random logical address

- A disk can transfer data at
 - S MB/s under a sequential workload
 - R MB/s under a random workload

Evaluating RAID Performance Example

■ sequential (S) vs random (R)

- **Sequential** : transfer 10 MB on average as continuous data.
- **Random** : transfer 10 KB on average.
- Average seek time: 7 ms
- Average rotational delay: 3 ms
- Transfer rate of disk: 50 MB/s

■ Results:

- $S = \frac{\text{Amount of Data}}{\text{Time to access}} = \frac{10 \text{ MB}}{210 \text{ ms}} = 47.62 \text{ MB /s}$
- $R = \frac{\text{Amount of Data}}{\text{Time to access}} = \frac{10 \text{ KB}}{10.195 \text{ ms}} = 0.981 \text{ MB /s}$

Evaluating RAID-0 Performance

N : the number of disks

- Single request latency
 - Identical to that of a single disk.

- Steady-state throughput
 - **Sequential** workload : $N \cdot S$ MB/s
 - **Random** workload : $N \cdot S$ MB /s



RAID Level 1 : Mirroring

- RAID Level 1 tolerates **disk failures**.
 - **Copy** more than one of **each block** in the system.
 - Copy block places on a separate disk.

Disk 0	Disk 1	Disk 2	Disk 3
0	0	1	1
2	2	3	3
4	4	5	5
6	6	7	7

Simple RAID-1: Mirroring (Keep two physical copies)

- ▶ RAID-10 (RAID 1+0) : mirrored pairs and then stripe
- ▶ RAID-01 (RAID 0+1) : contain two large striping arrays, and then mirrors

RAID-1 Analysis

N : the number of disks

- **Capacity:** RAID-1 is Expensive
 - The useful capacity of RAID-1 is $N/2$.
- **Reliability:** RAID-1 does well.
 - It can tolerate the failure of any one disk (up to $N/2$ failures depending on which disk fail).



Performance of RAID-1

- Two physical writes to complete
 - It suffers the worst-case seek and rotational delay of the two request.
 - Steady-state throughput
 - ▶ **Sequential Write** : $\frac{N}{2} \cdot S$ MB/s
 - Each logical write must result in two physical writes.
 - ▶ **Sequential Read** : $\frac{N}{2} \cdot S$ MB/s
 - Each disk will only deliver half its peak bandwidth.
 - ▶ **Random Write** : $\frac{N}{2} \cdot R$ MB/s
 - Each logical write must turn into two physical writes.
 - ▶ **Random Read** : $N \cdot R$ MB/s
 - Distribute the reads across all the disks.

RAID Level 4 : Saving Space With Parity

- Add a single parity block

- A Parity block stores the *redundant information* for that stripe of blocks.

* P: Parity

Disk 0	Disk 1	Disk 2	Disk 3	Disk 4
0	0	1	1	P0
2	2	3	3	P1
4	4	5	5	P2
6	6	7	7	P3

Five-disk RAID-4 system layout

RAID Level 4 (Cont.)

- **Compute parity** : the XOR of all of bits

C0	C1	C2	C3	P
0	0	1	1	$\text{XOR}(0,0,1,1)=0$
0	1	0	0	$\text{XOR}(0,1,0,0)=1$

- **Recover from parity**

- Imagine the bit of the C2 in the first row is lost.
 1. Reading the other values in that row : 0, 0, 1
 2. The parity bit is 0 → even number of 1's in the row
 3. What the missing data must be: a 1.

RAID-4 Analysis

N : the number of disks

■ Capacity

- The useful capacity is $(N - 1)$.

■ Reliability

- RAID-4 tolerates 1 disk failure and no more.



RAID-4 Analysis (Cont.)

■ Performance

● Steady-state throughput

- ▶ Sequential read: $(N - 1) \cdot S$ MB/s
- ▶ Sequential write: $(N - 1) \cdot S$ MB/s

Disk 0	Disk 1	Disk 2	Disk 3	Disk 4
0	1	2	3	P0
4	5	6	7	P1
8	9	10	11	P2
12	13	14	15	P3

Full-stripe Writes In RAID-4

- ▶ Random read: $(N - 1) \cdot R$ MB/s

Random write performance for RAID-4

- Overwrite a block + update the parity
- **Method 1:** *additive parity*
 - Read in all of the other data blocks in the stripe
 - XOR those blocks with the new block (1)
 - **Problem:** the performance scales with the number of disks



Random write performance for RAID-4 (Cont.)



■ Method 2: *subtractive parity*

C0	C1	C2	C3	P
0	0	1	1	XOR(0,0,1,1)=0

$$P(new) = (C2(old) \text{ XOR } C2(new)) \text{ XOR } P(old)$$

- Update C2(old) → C2(new)
 1. Read in the old data at C2 (C2(old)=1) and the old parity (P(old)=0)
 2. Calculate P(new):
 - If C2(new)==C2(old) → P(new)==P(old)
 - If C2(new)≠C2(old) → Flip the old parity bit

Small-write problem

- The parity disk can be a **bottleneck**.
 - Example: update blocks 4 and 13 (marked with *)

Disk 0	Disk 1	Disk 2	Disk 3	Disk 4
0	1	2	3	P0
*4	5	6	7	+P1
8	9	10	11	P2
12	*13	14	15	+P3

Writes To 4, 13 And Respective Parity Blocks.

- ▶ Disk 0 and Disk 1 can be accessed in parallel.
- ▶ Disk 4 prevents any parallelism.

RAID-4 throughput under random small writes is $(\frac{R}{2})$ MB/s (*terrible*).

A I/O latency in RAID-4

■ A single read

- Equivalent to the latency of a single disk request.

■ A single write

- Two reads and then two writes
 - ▶ Data block + Parity block
 - ▶ The reads and writes can happen in parallel.
- Total latency *is about twice* that of a single disk.

RAID Level 5: Rotating Parity

- RAID-5 is **solution of** small write problem.
 - Rotate the parity blocks across drives.
 - Remove the parity-disk bottleneck for RAID-4

Disk 0	Disk 1	Disk 2	Disk 3	Disk 4
0	1	2	3	P0
5	6	7	P1	4
10	11	P2	8	9
15	P3	12	13	14
P4	16	17	18	19

RAID-5 With Rotated Parity

RAID-5 Analysis

N : the number of disks

■ Capacity

- The useful capacity for a RAID group is $(N - 1)$.

■ Reliability

- RAID-5 tolerates 1 disk failure and no more.

RAID-5 Analysis (Cont.)

N : the number of disks

■ Performance

- Sequential read and write
 - A single read and write request
- } Same as RAID-4
- Random read : a little better than RAID-4
 - ▶ RAID-5 can utilize all of the disks.
 - Random write : $\frac{N}{4} \cdot R$ MB/s
 - ▶ The factor of four loss is cost of using parity-based RAID.

RAID Comparison: A Summary

N : the number of disks

D : the time that a request to a single disk take

	RAID-0	RAID-1	RAID-4	RAID-5
Capacity	N	$N/1$	$N-1$	$N-1$
Reliability	0	1 (for sure) $\frac{N}{2}$ (if lucky)	1	1
Throughput				
Sequential Read	$N \cdot S$	$(N/2) \cdot S$	$(N-1) \cdot S$	$(N-1) \cdot S$
Sequential Write	$N \cdot S$	$(N/2) \cdot S$	$(N-1) \cdot S$	$(N-1) \cdot S$
Random Read	$N \cdot R$	$N \cdot R$	$(N-1) \cdot R$	$N \cdot R$
Random Write	$N \cdot R$	$(N/2) \cdot R$	$\frac{1}{2} R$	$\frac{N}{4} R$
Latency				
Read	D	D	D	D
Write	D	D	$2D$	$2D$

RAID Capacity, Reliability, and Performance

RAID Comparison: A Summary

- **Performance** and do not care about reliability → RAID-0 (Striping)
- **Random I/O** performance and **Reliability** → RAID-1 (Mirroring)
- **Capacity** and **Reliability** → RAID-5
- **Sequential I/O** and Maximize **Capacity** → RAID-5





谢谢！

